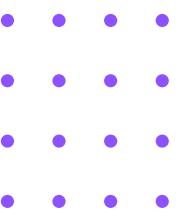


# Classez des images à l'aide d'algorithmes de Deep Learning

## Objectif

- Réaliser un algorithme de détection de la race de chiens



# Présentation du projet

**Situation : bénévole pour une association de protection des animaux**

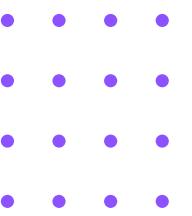
**Problématique : difficulté d'indexer les images des animaux de l'association**

**Objectif : développer un algorithme de détection de la race d'un chien**

**Avantage : amélioration de l'efficacité de l'indexation des images**

**Utilisation des données du Stanford Dogs Dataset pour entraîner le modèle**

- <http://vision.stanford.edu/aditya86/ImageNetDogs/>
- **20580 images de chiens réparties en 120 races**



# Modèle de base

## Sélection de 10 races

```
# Création d'une partition de 80% des images pour l'entraînement
train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset='training',
    seed=420,
    image_size=img_size,
    batch_size=batch_size,
    label_mode='categorical')

# Création d'une partition temporaire de 20% des images
temp_ds = tf.keras.preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset='validation',
    seed=420,
    image_size=img_size,
    batch_size=batch_size,
    label_mode='categorical')

# Séparation en deux partitions égales (10%) de la partition temporaire pour la validation et les tests
val_ds = temp_ds.take(len(temp_ds)//2)
test_ds = temp_ds.skip(len(temp_ds)//2)

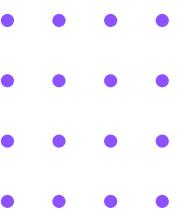
# Nombre de classes possible de classification
num_classes = len(train_ds.class_names)
```

### Générateur

- traitement par répertoire
- séparation train / temp
- redimensionne les images
- chargement par batch
- categorical\_crossentropy loss

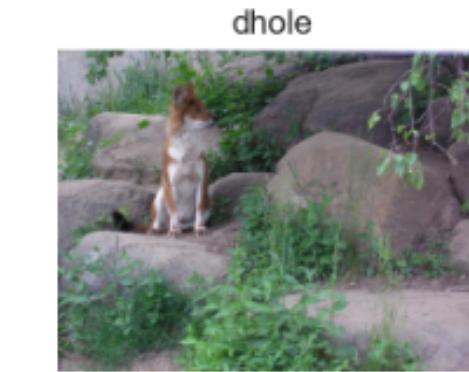
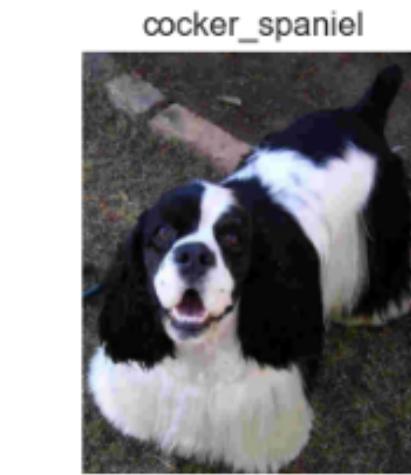
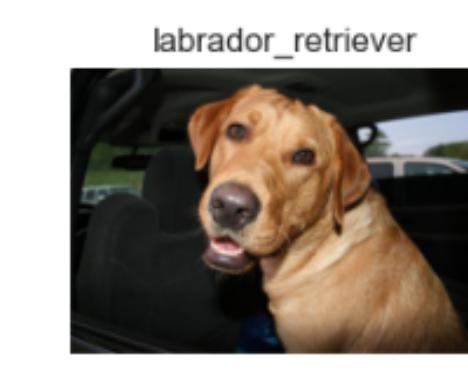
### Séparation val / test

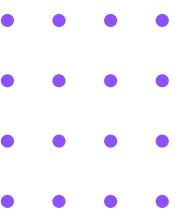
### Nombre de classes



# Modèle de base

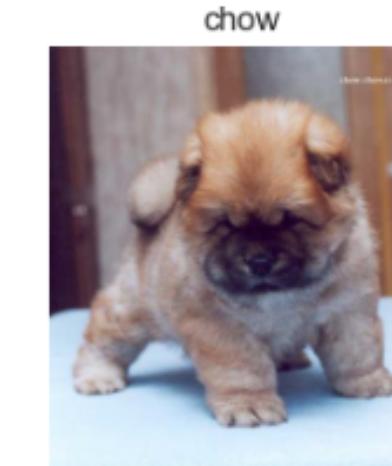
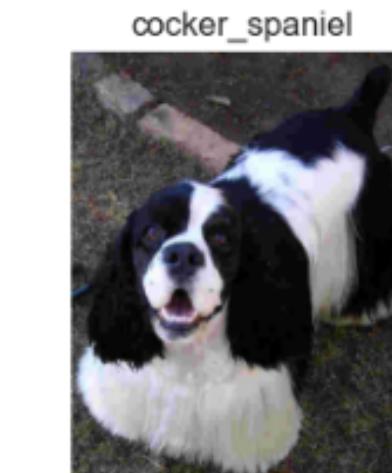
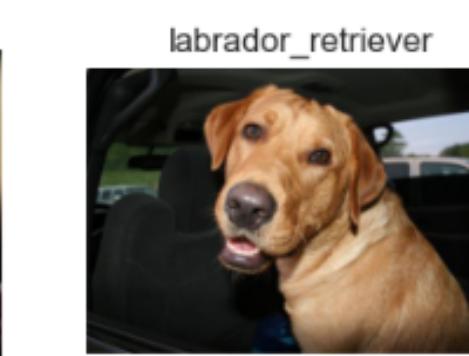
## Sélection de 10 races





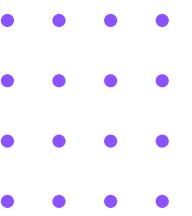
# Modèle de base

## Sélection de 10 races



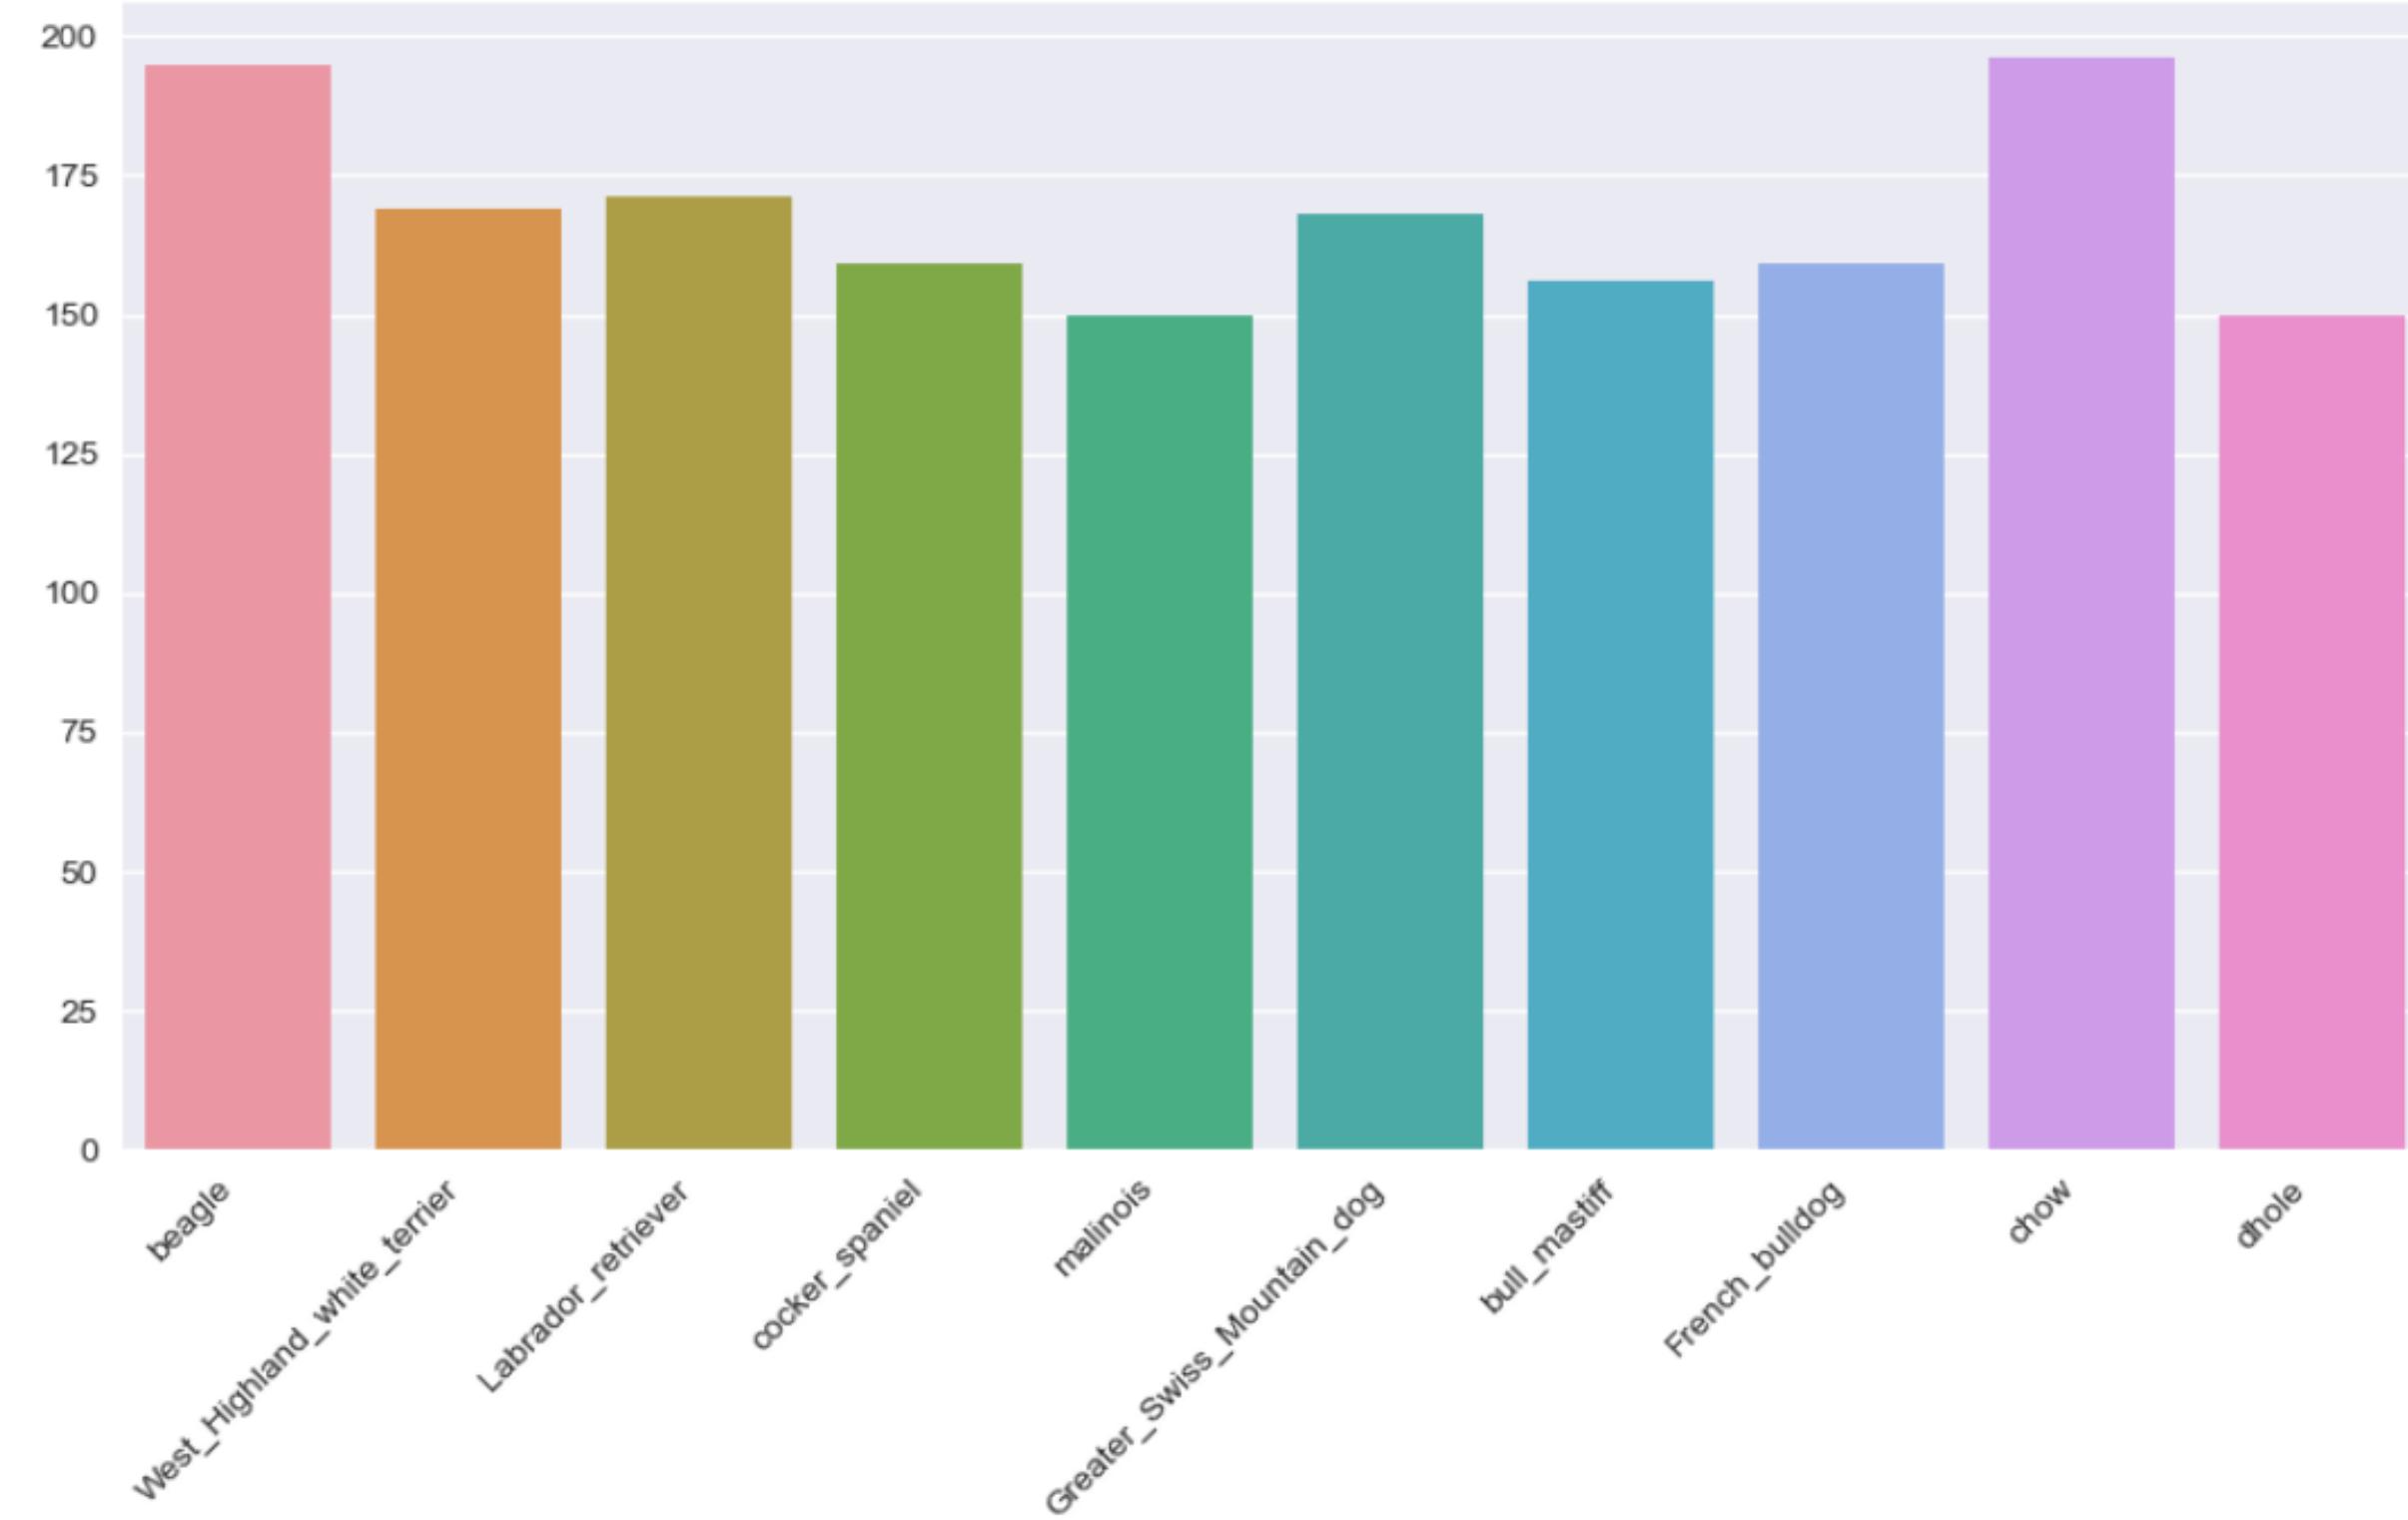
Grande variabilité en termes de taille, de résolution et de qualité

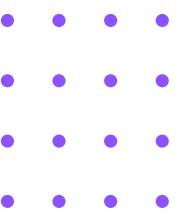




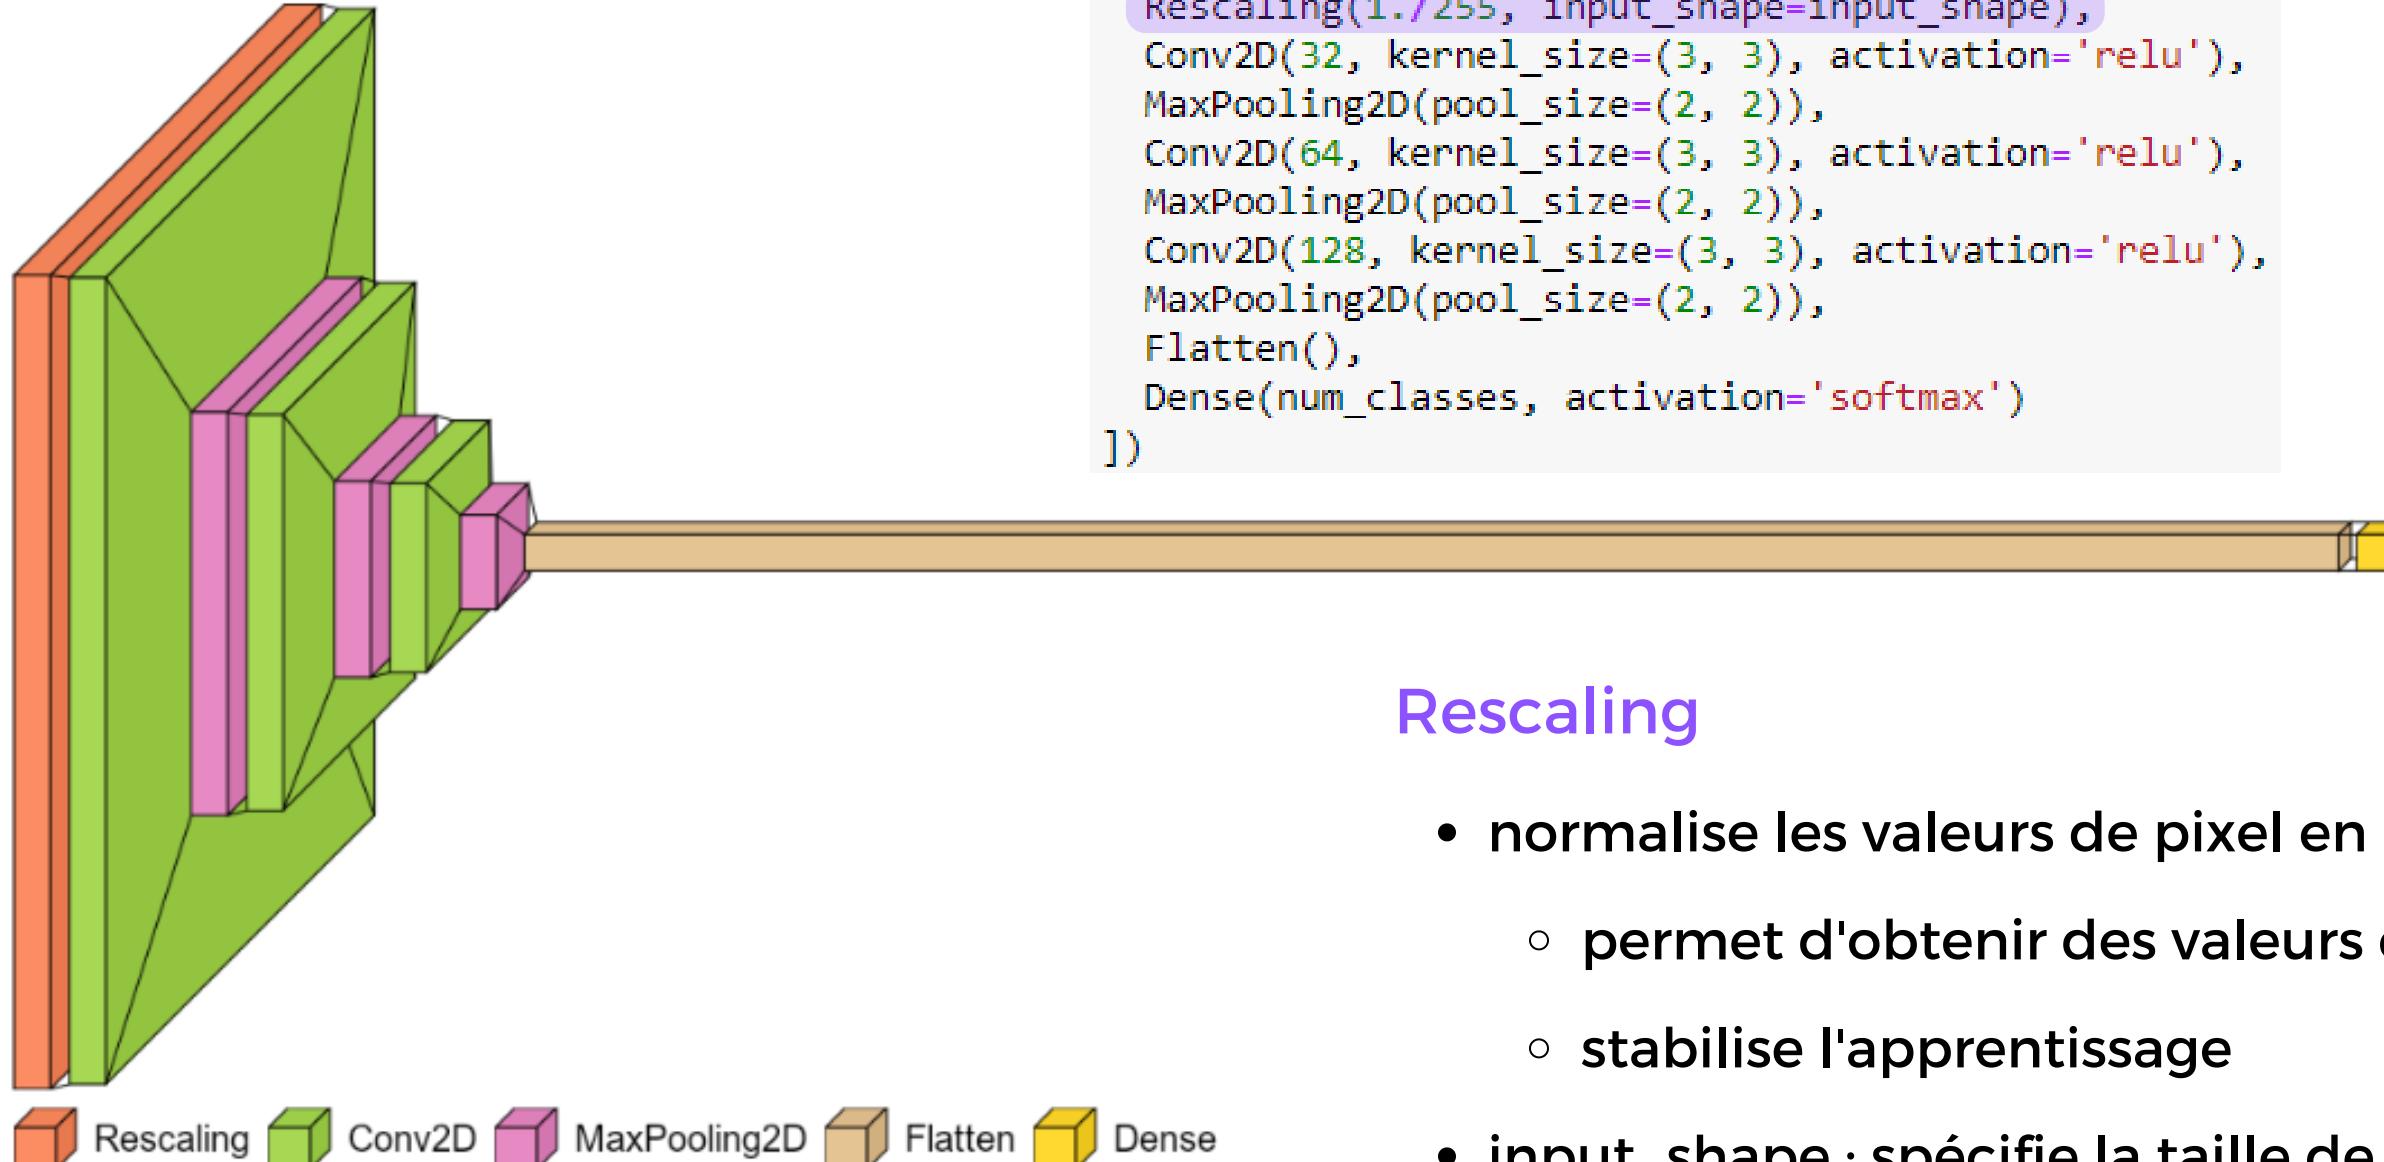
# Modèle de base

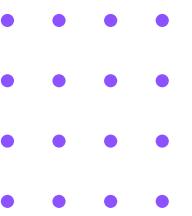
Nombre d'images par race de chiens



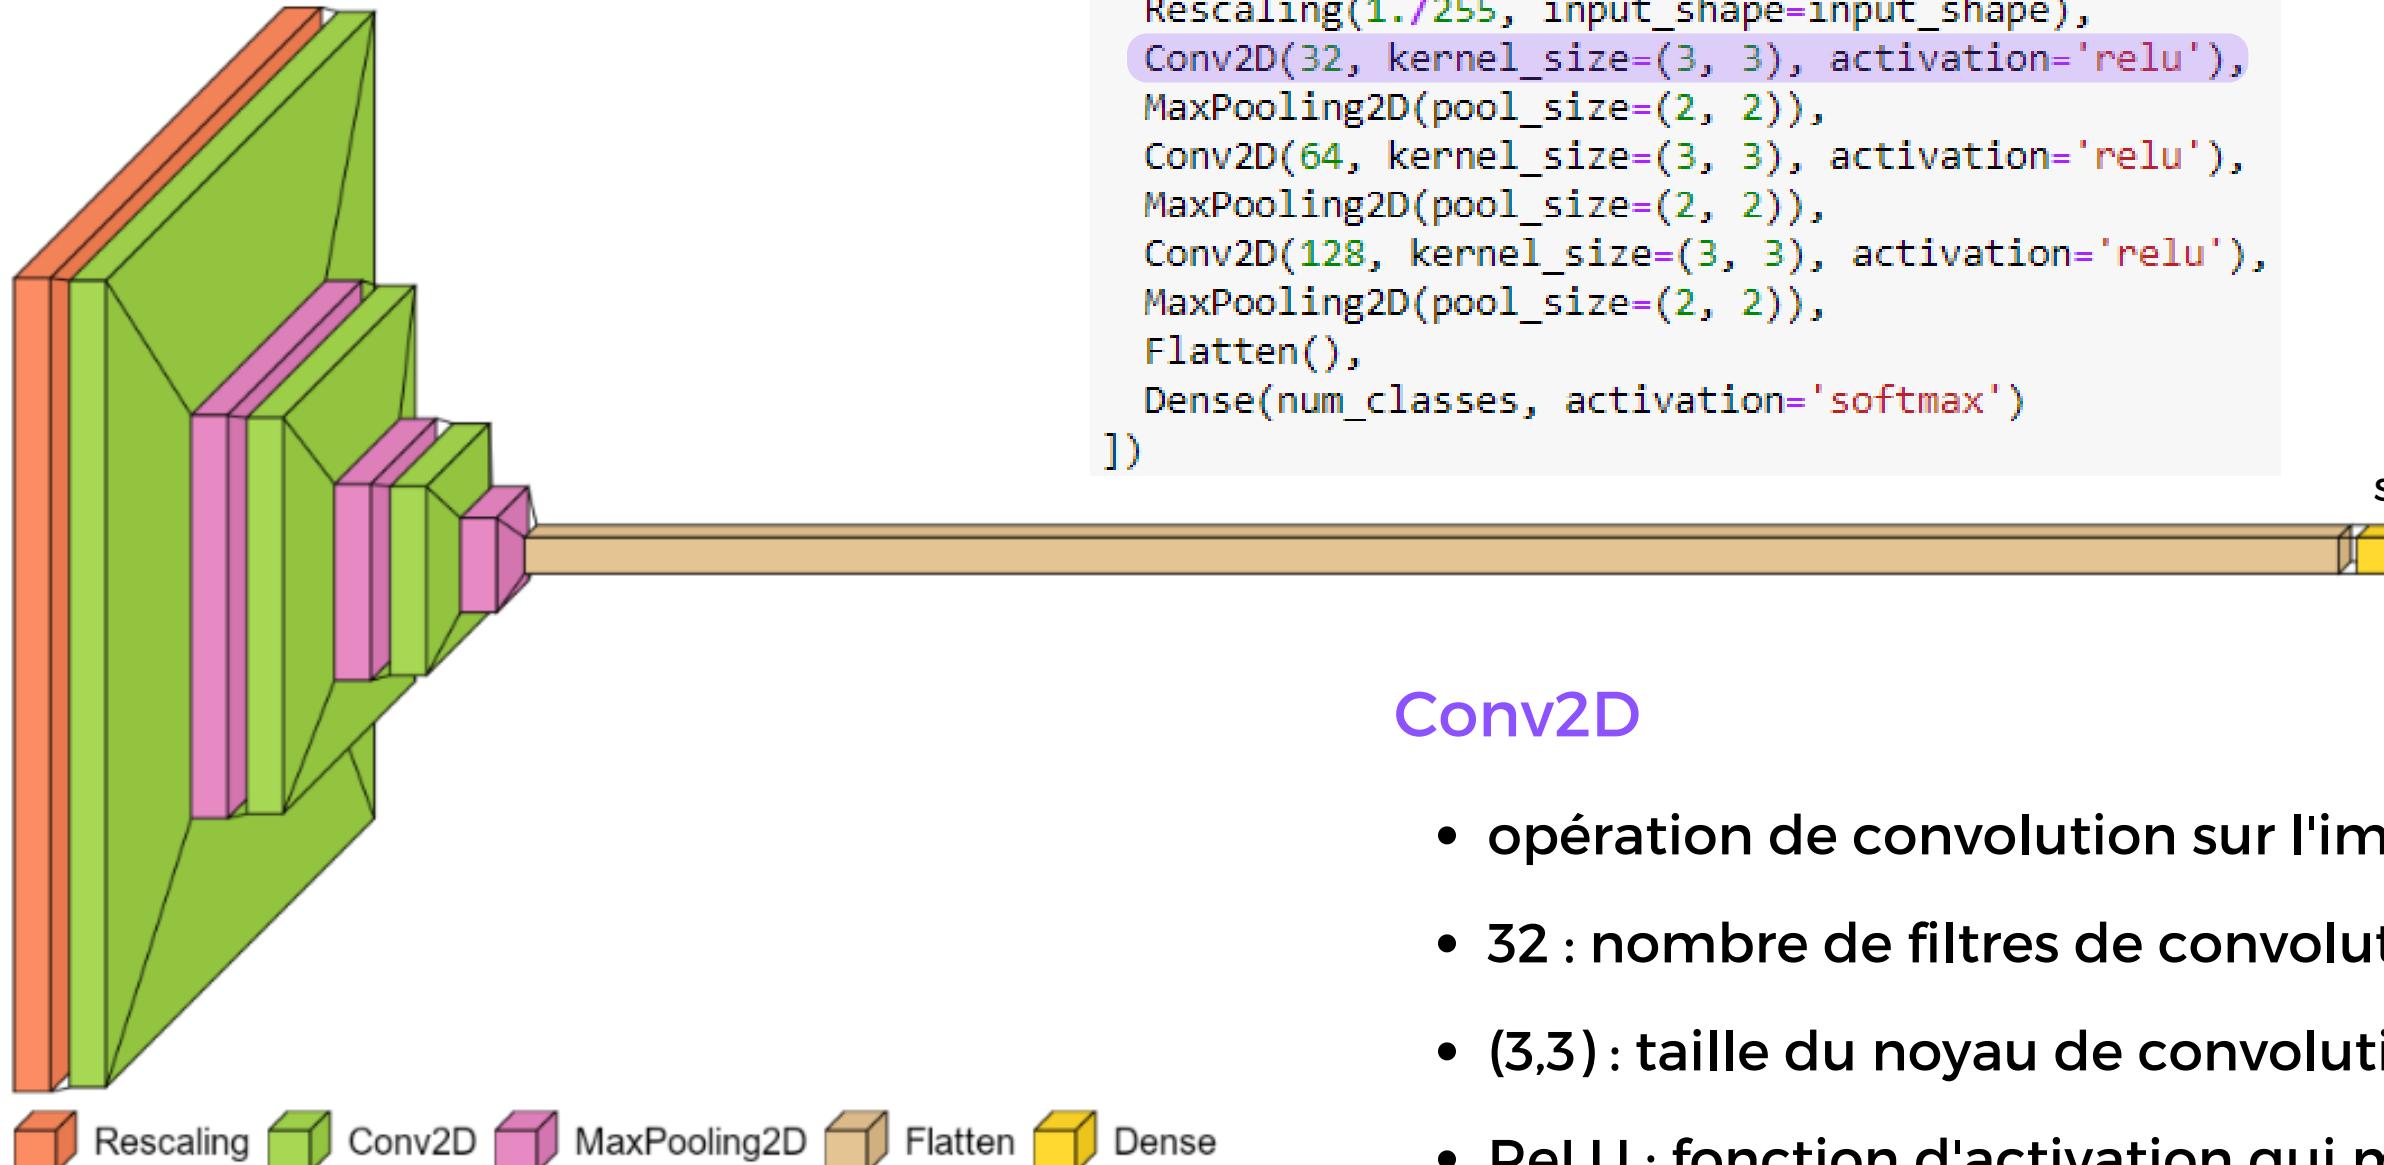


# Modèle de base





# Modèle de base



```
model = Sequential([
    Rescaling(1./255, input_shape=input_shape),
    Conv2D(32, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(128, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(num_classes, activation='softmax')
])
```

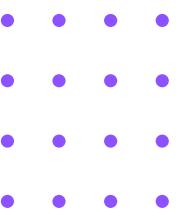
Source : [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

## Conv2D

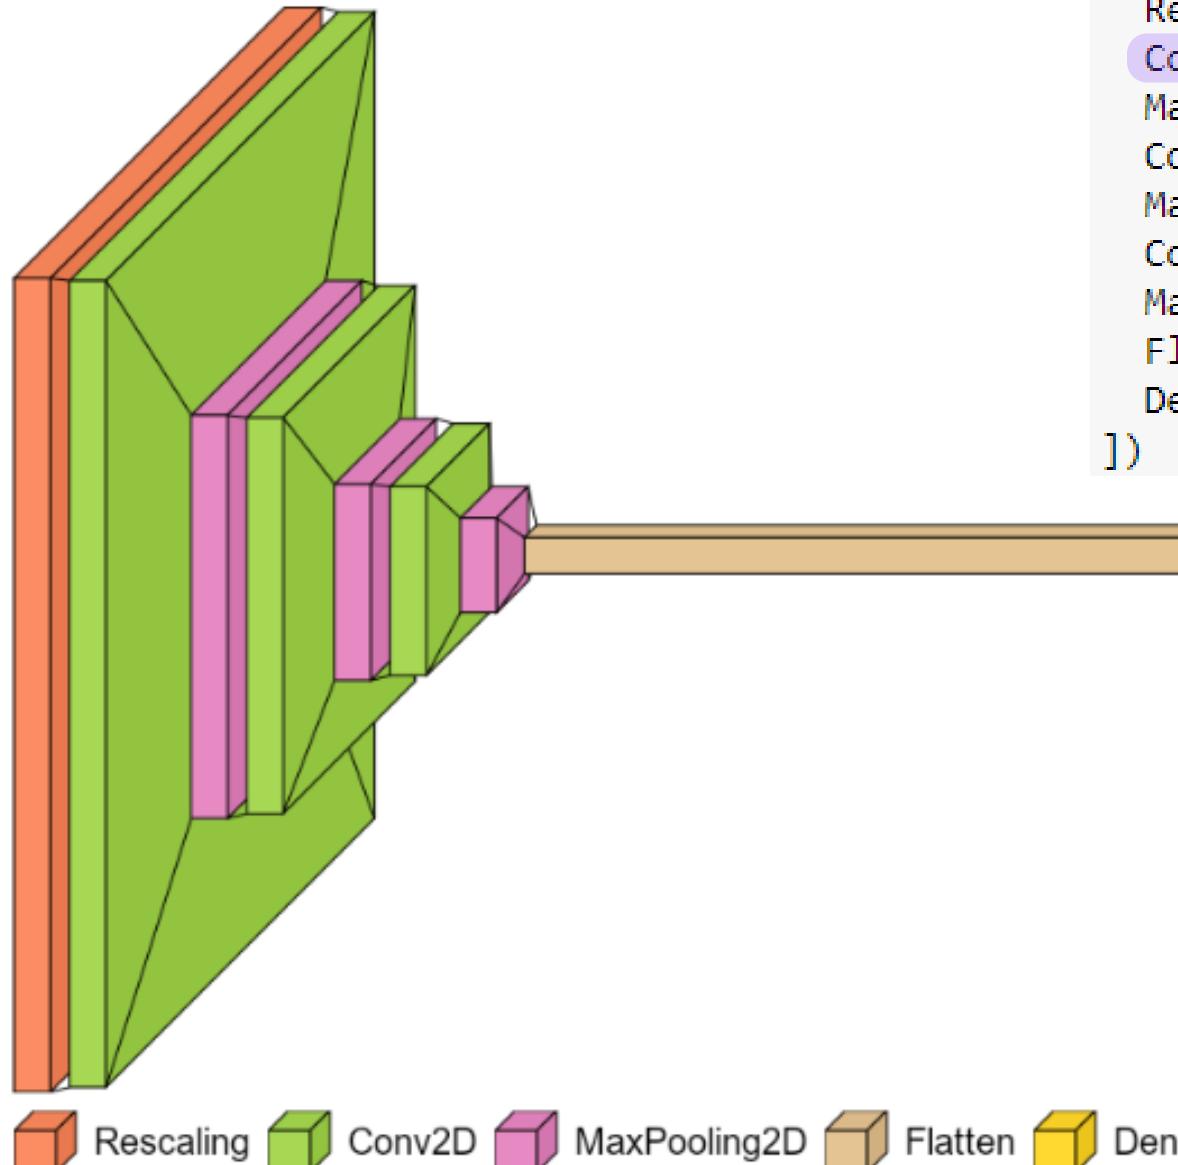
- opération de convolution sur l'image d'entrée
- 32 : nombre de filtres de convolution appliqués
- (3,3) : taille du noyau de convolution
- ReLU : fonction d'activation qui mappe les valeurs négatives sur 0 et conserve les valeurs positives
- produit terme à terme (produit de Hadamard)

$$\underbrace{\begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix}}_{\text{Une partie de l'image}} \odot \underbrace{\begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{pmatrix}}_{\text{Filtre de convolution}} = \underbrace{\begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}}_{\text{Partie de l'image convoluee}}$$

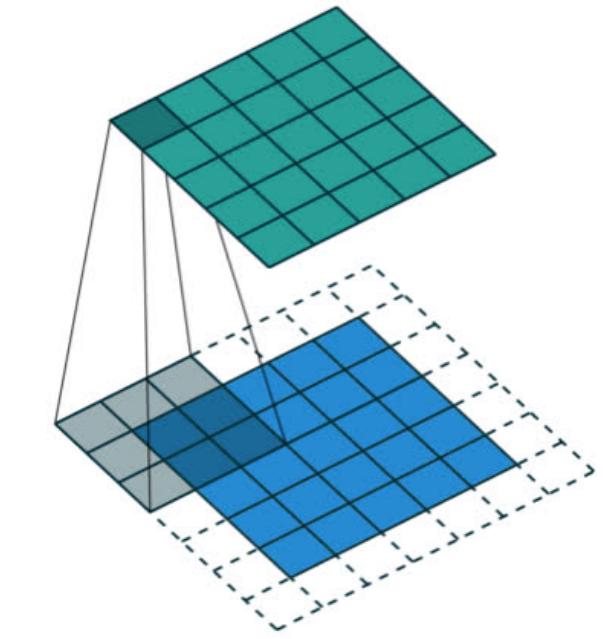




# Modèle de base



```
model = Sequential([
    Rescaling(1./255, input_shape=input_shape),
    Conv2D(32, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(128, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(num_classes, activation='softmax')
])
```



Source : [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)

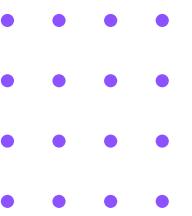
1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

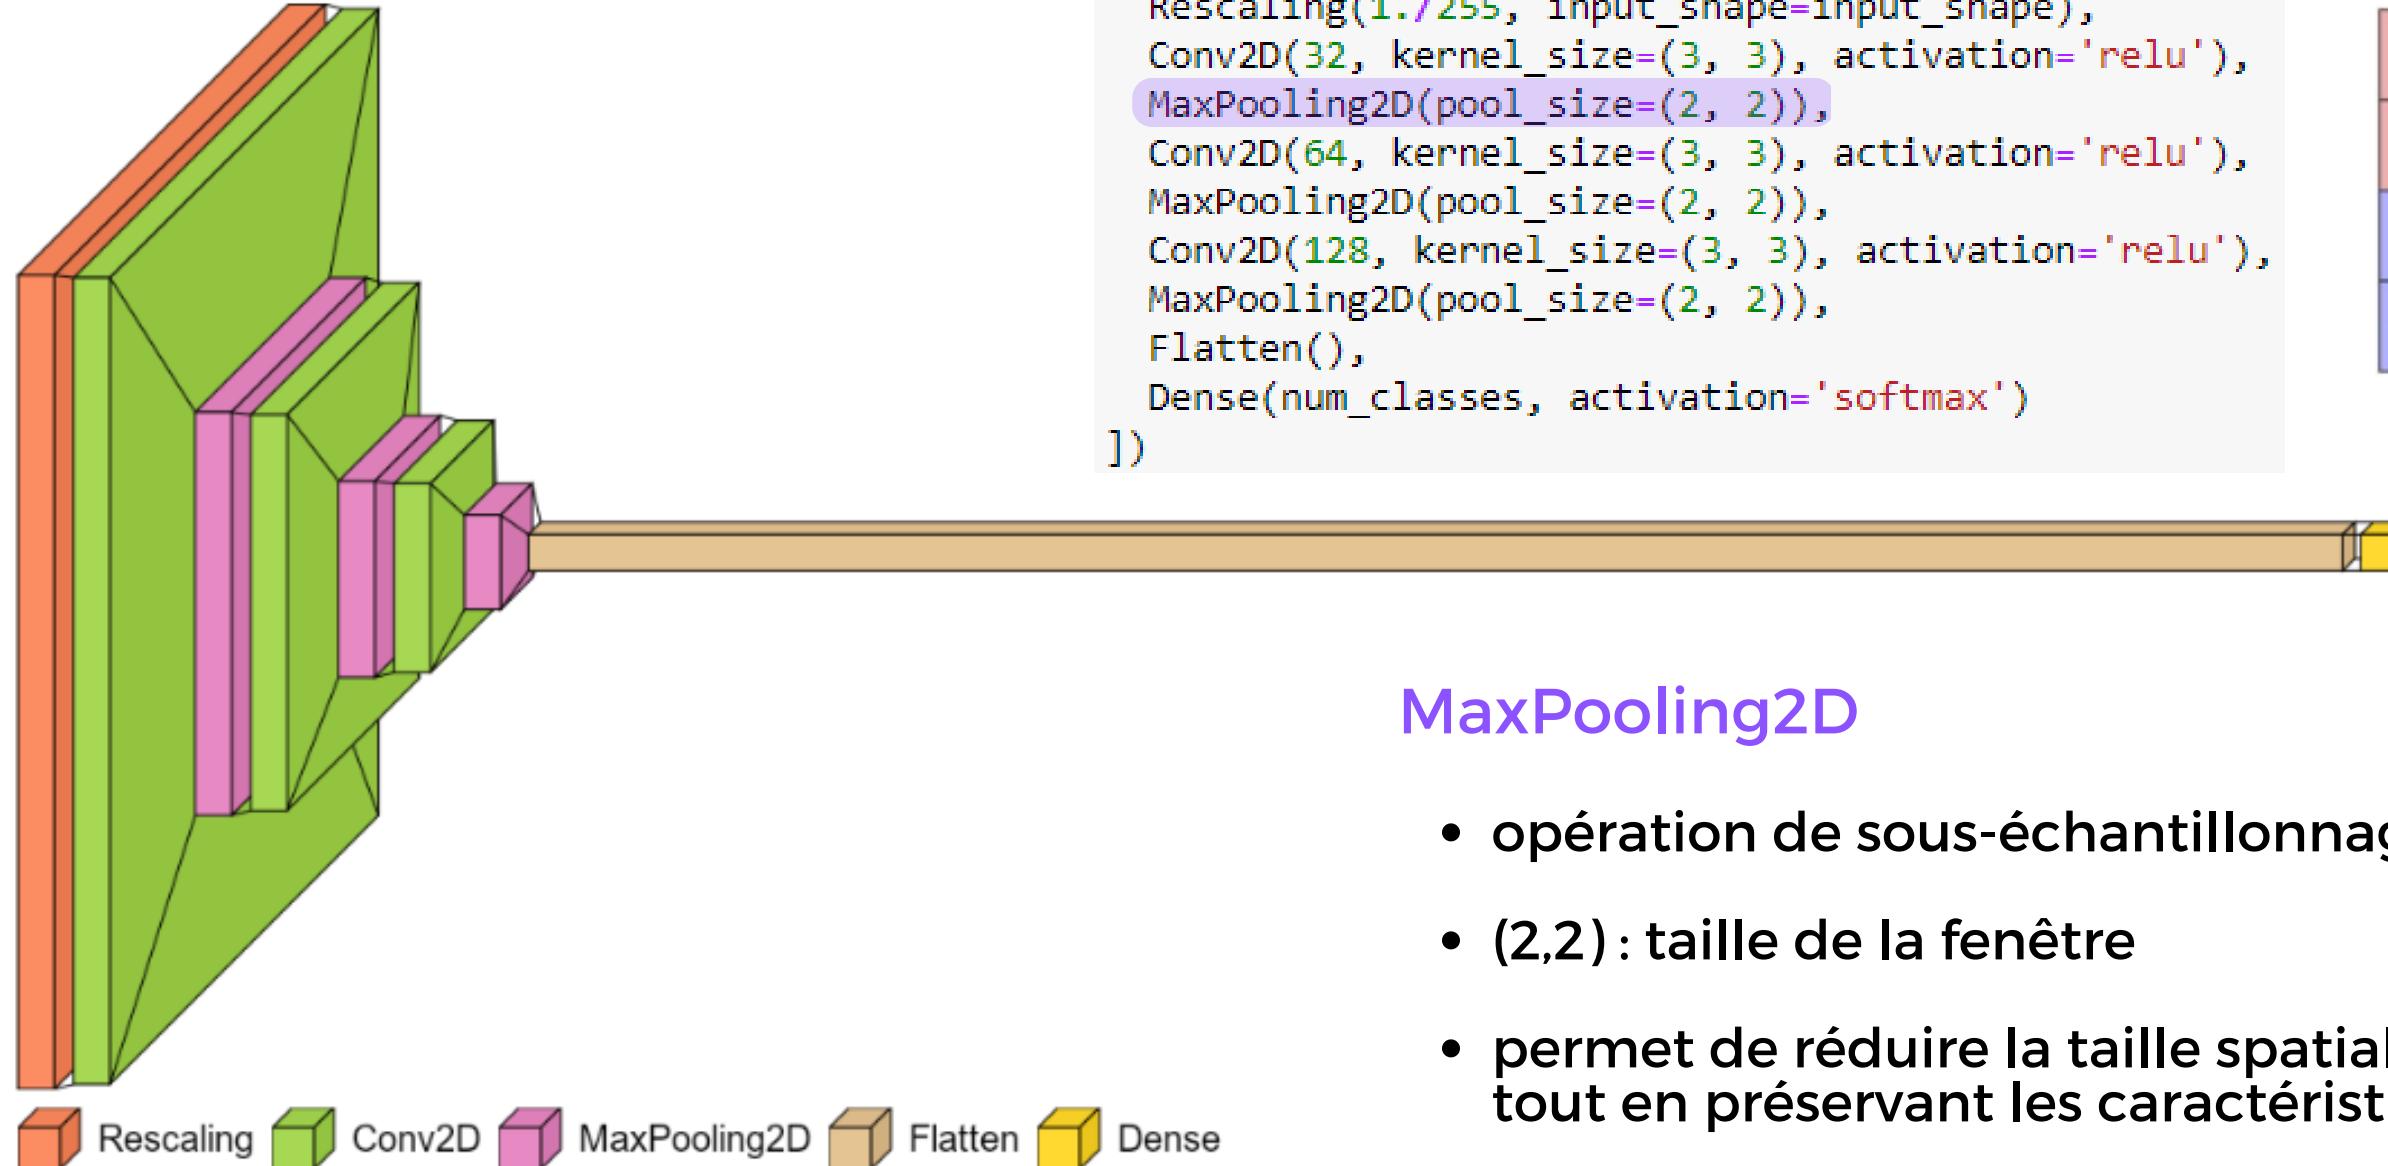
4		

Convolved  
Feature

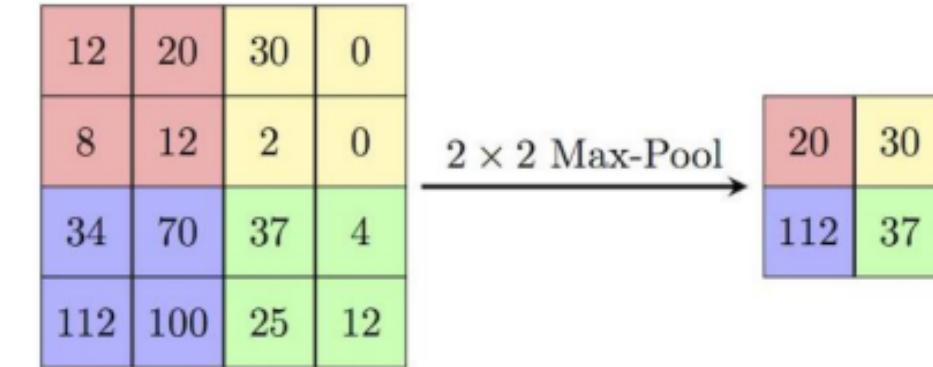
Source : IceCream Labs



# Modèle de base

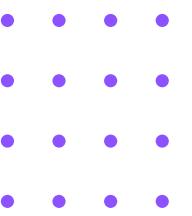


```
model = Sequential([
    Rescaling(1./255, input_shape=input_shape),
    Conv2D(32, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(128, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(num_classes, activation='softmax')
])
```

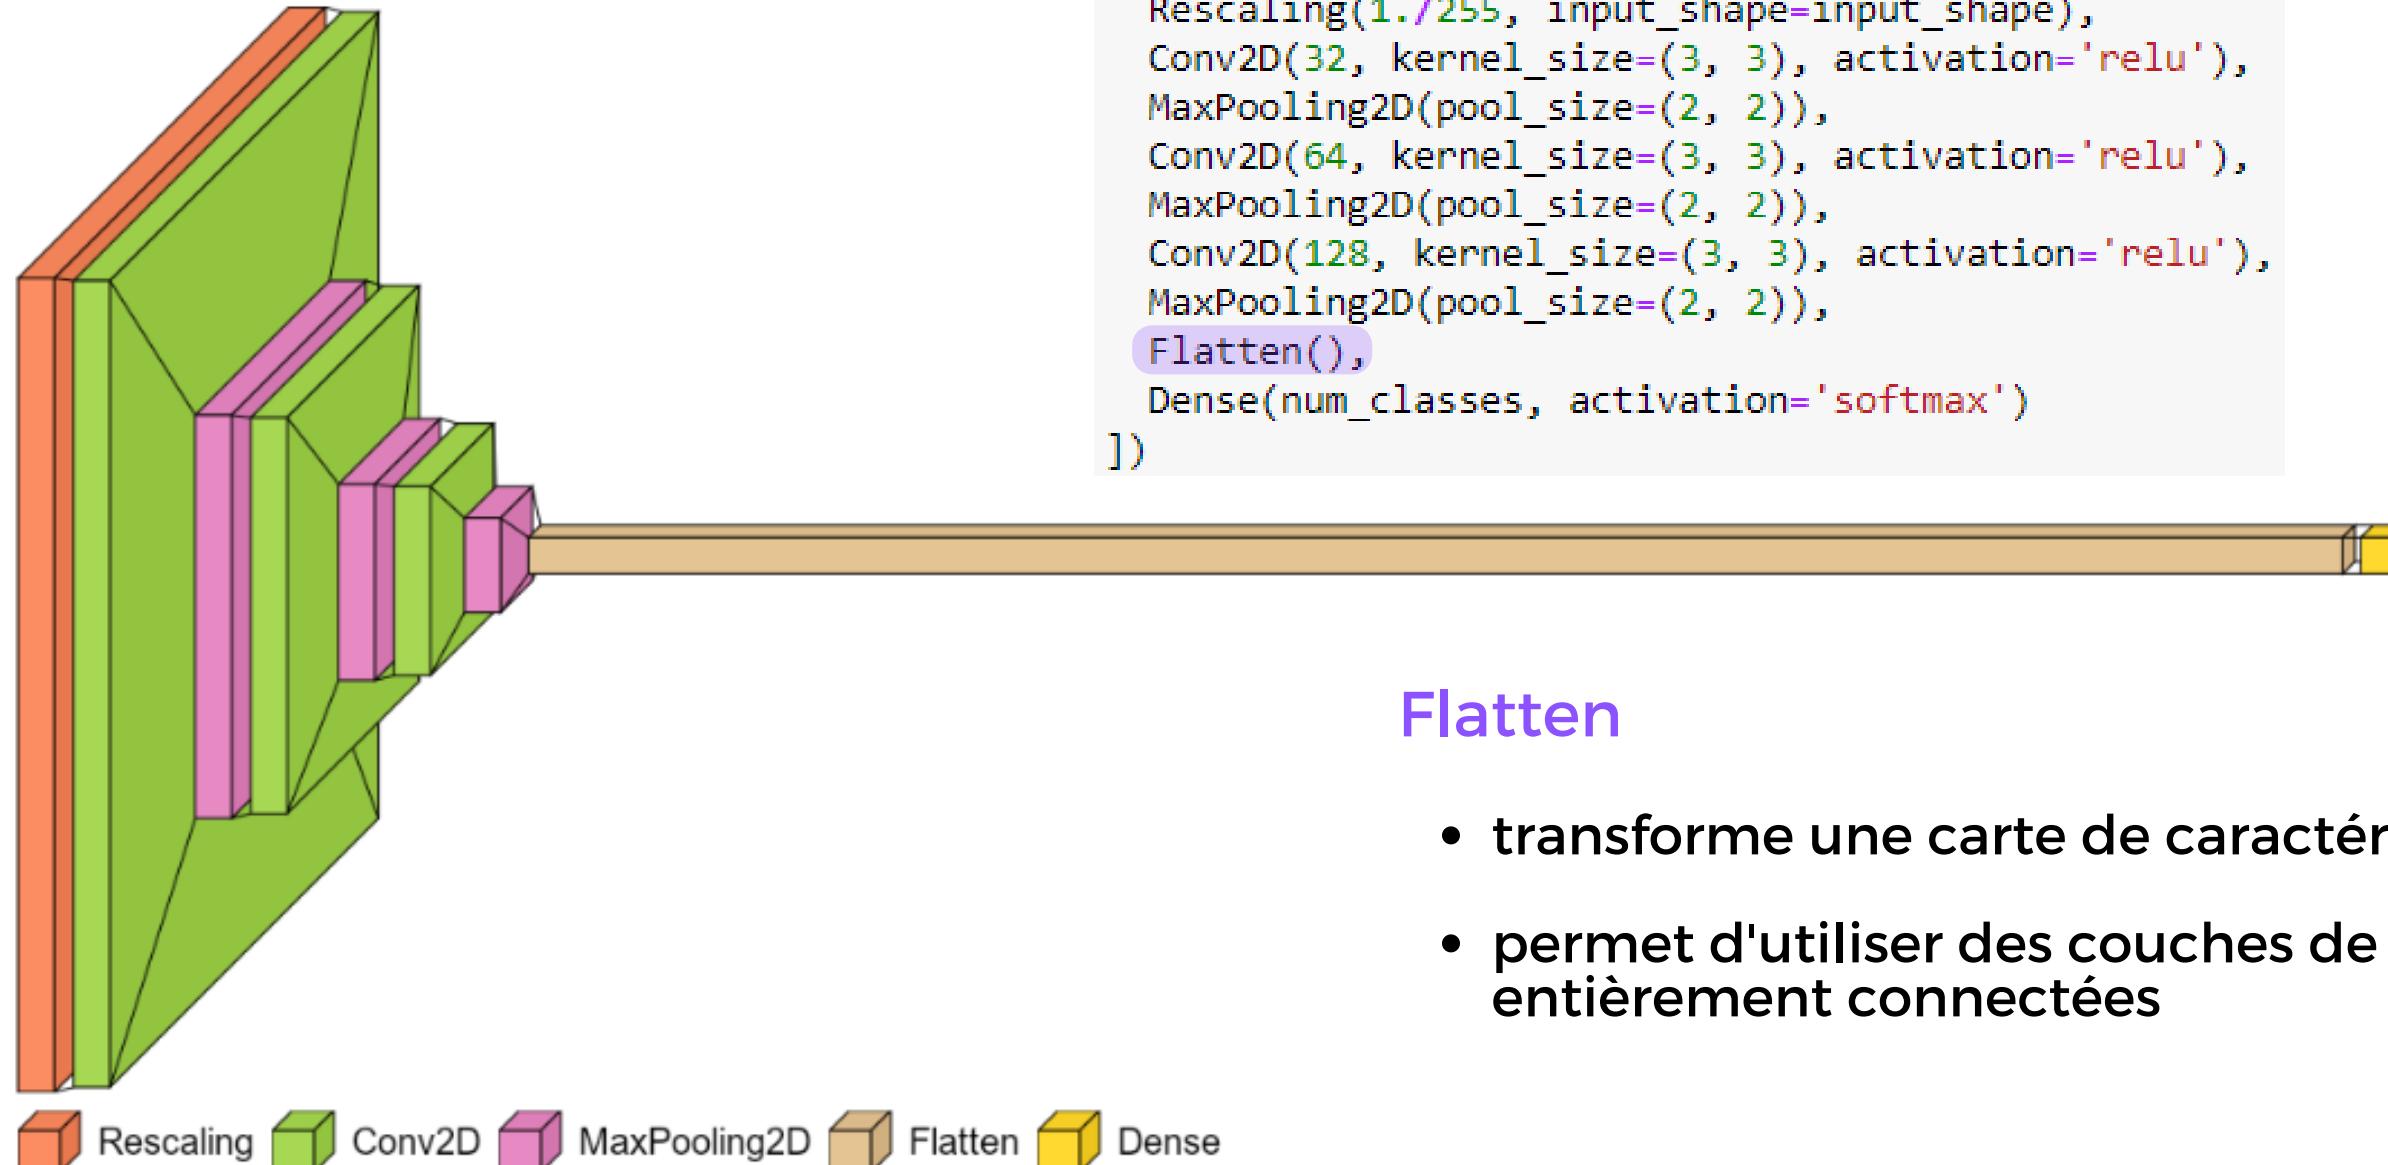


## MaxPooling2D

- opération de sous-échantillonnage
- (2,2) : taille de la fenêtre
- permet de réduire la taille spatiale de la carte de caractéristiques tout en préservant les caractéristiques les plus importantes
  - améliore l'efficacité du réseau
  - évite le sur-apprentissage

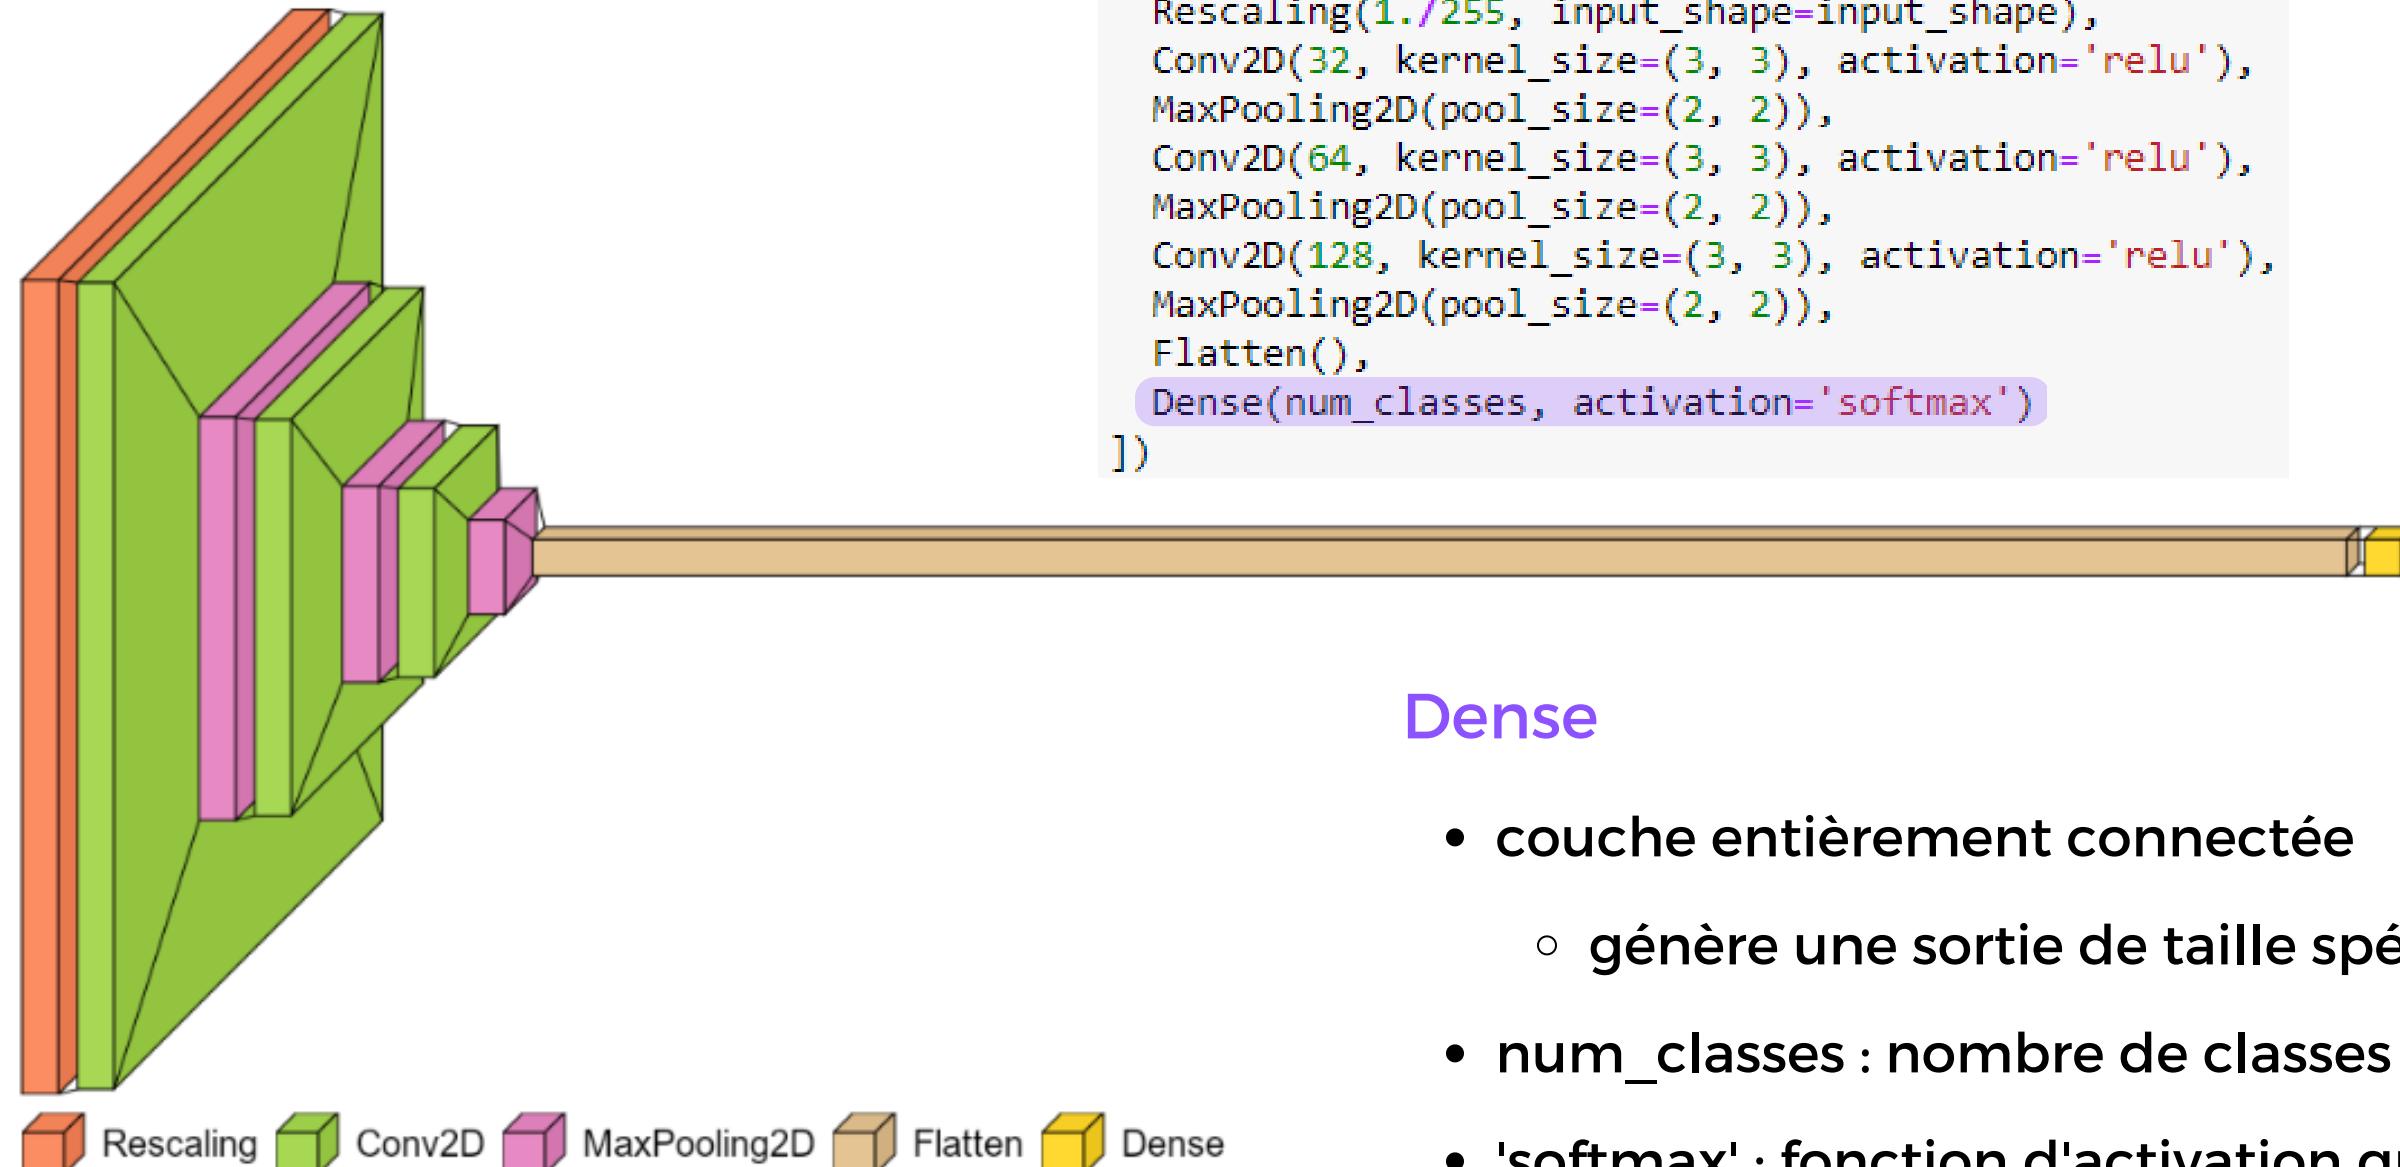


# Modèle de base



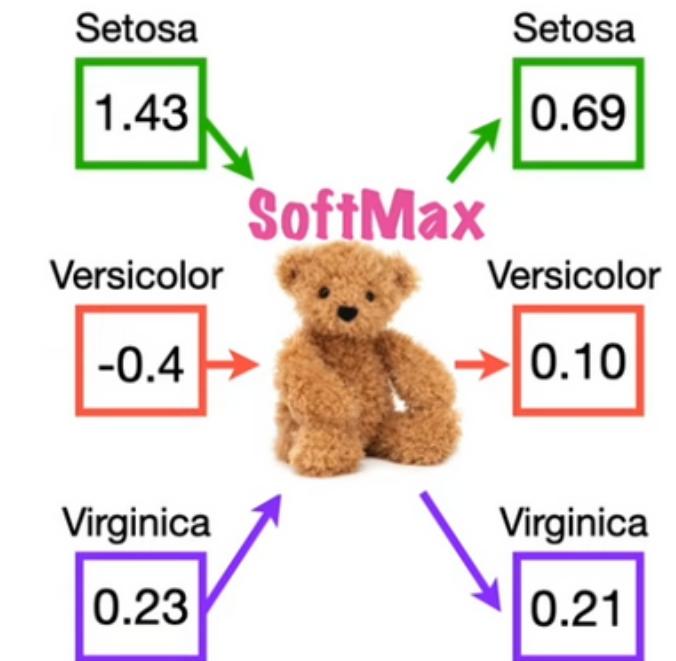


# Modèle de base



## Dense

- couche entièrement connectée
  - génère une sortie de taille spécifiée
- num\_classes : nombre de classes de sortie
- 'softmax' : fonction d'activation qui normalise les valeurs de sortie en une distribution de probabilité
  - idéal pour la classification multi-classes
  - $$\text{SoftMax}_i(\text{Output Values}) = \frac{e^{\text{Output Value}_i}}{\sum_{j=1}^k e^{\text{Output Value}_j}}$$



Source : StatQuestwithJoshStarmer (YouTube)

# Modèle de base

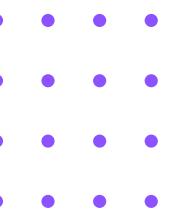


## Optimiseur Adam

- ajuste automatiquement le taux d'apprentissage pour chaque poids du réseau de neurones
- efficace : nécessite moins de mémoire
- optimisation plus rapide et plus efficace pour les ensembles de données de grande taille

# Modèle de base

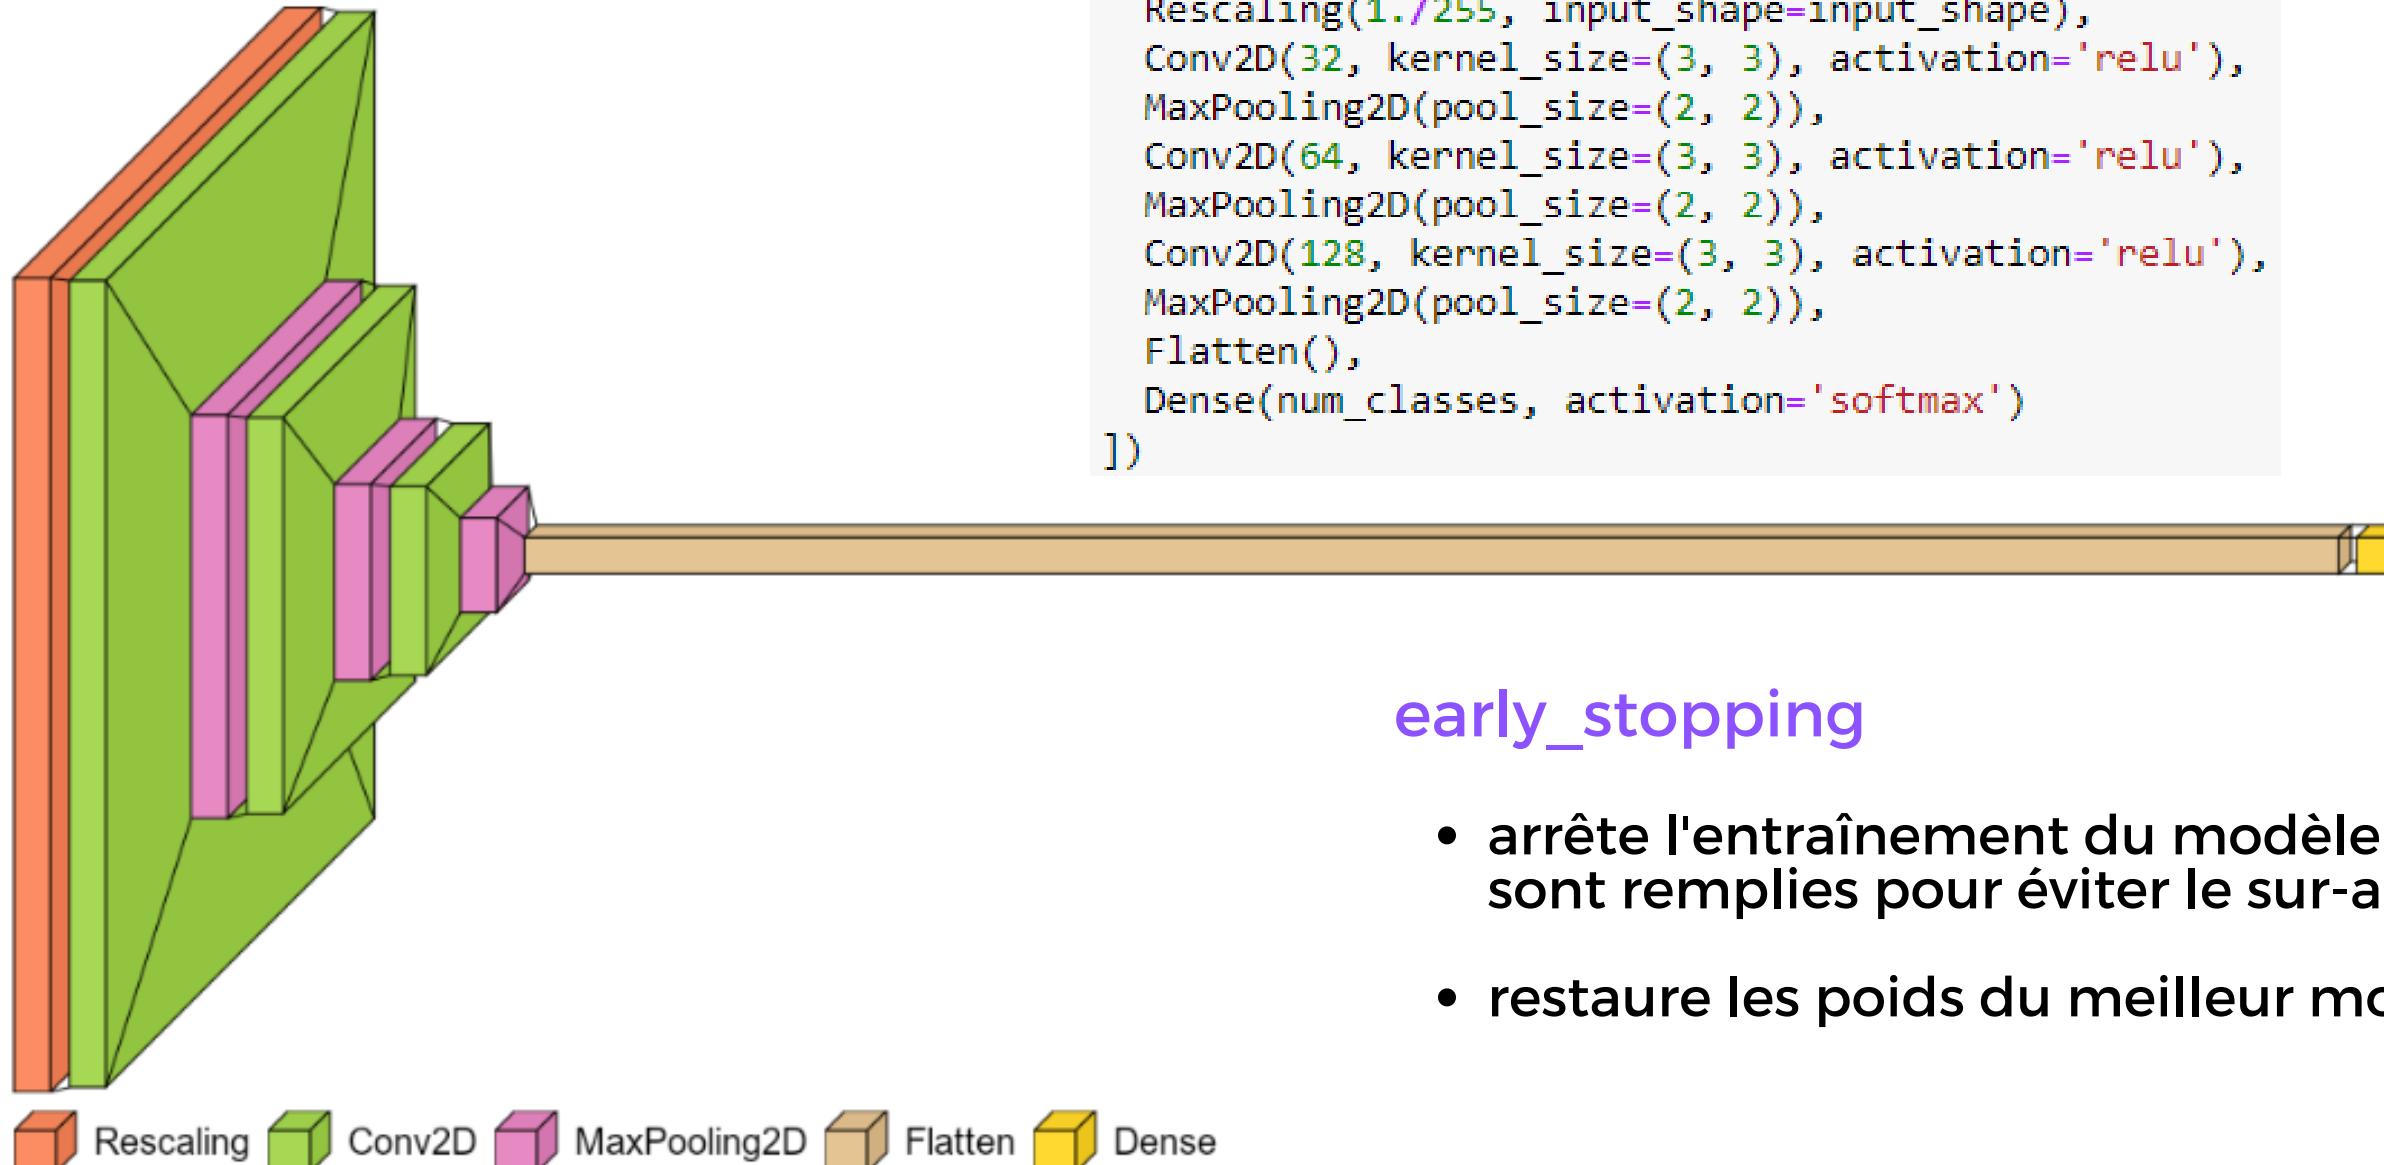




# Modèle de base



# Modèle de base



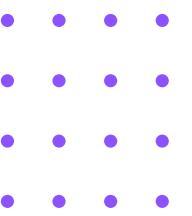
```
model = Sequential([
    Rescaling(1./255, input_shape=input_shape),
    Conv2D(32, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(64, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(128, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(num_classes, activation='softmax')
])
```

```
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['categorical_accuracy']
)

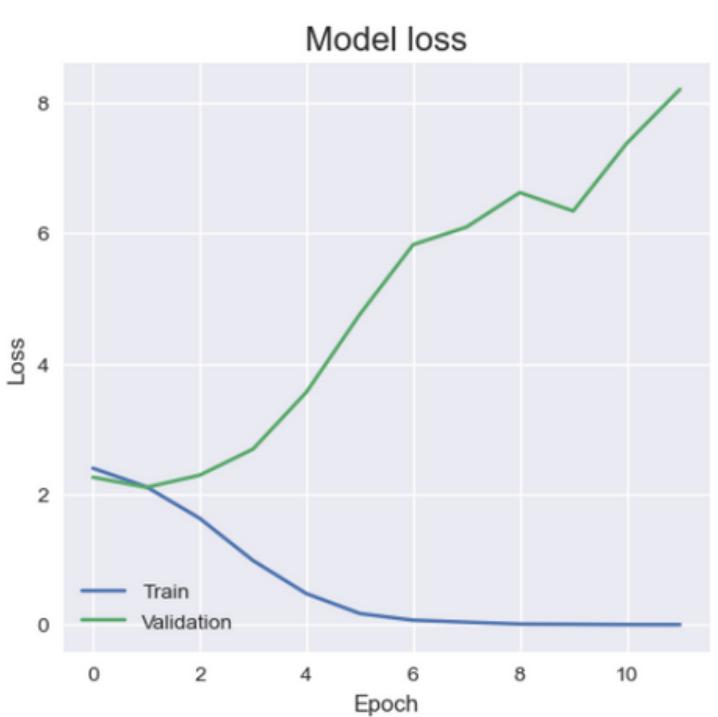
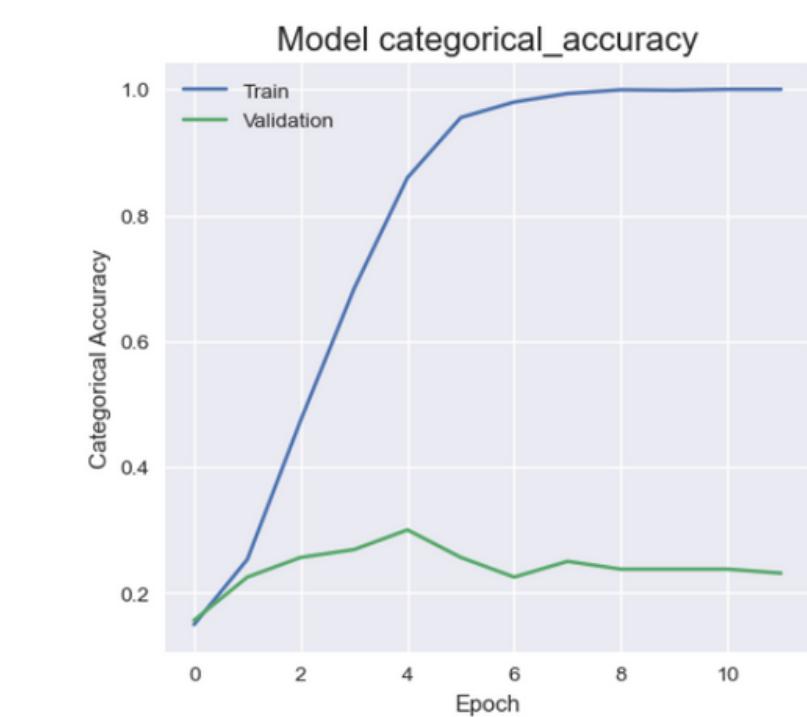
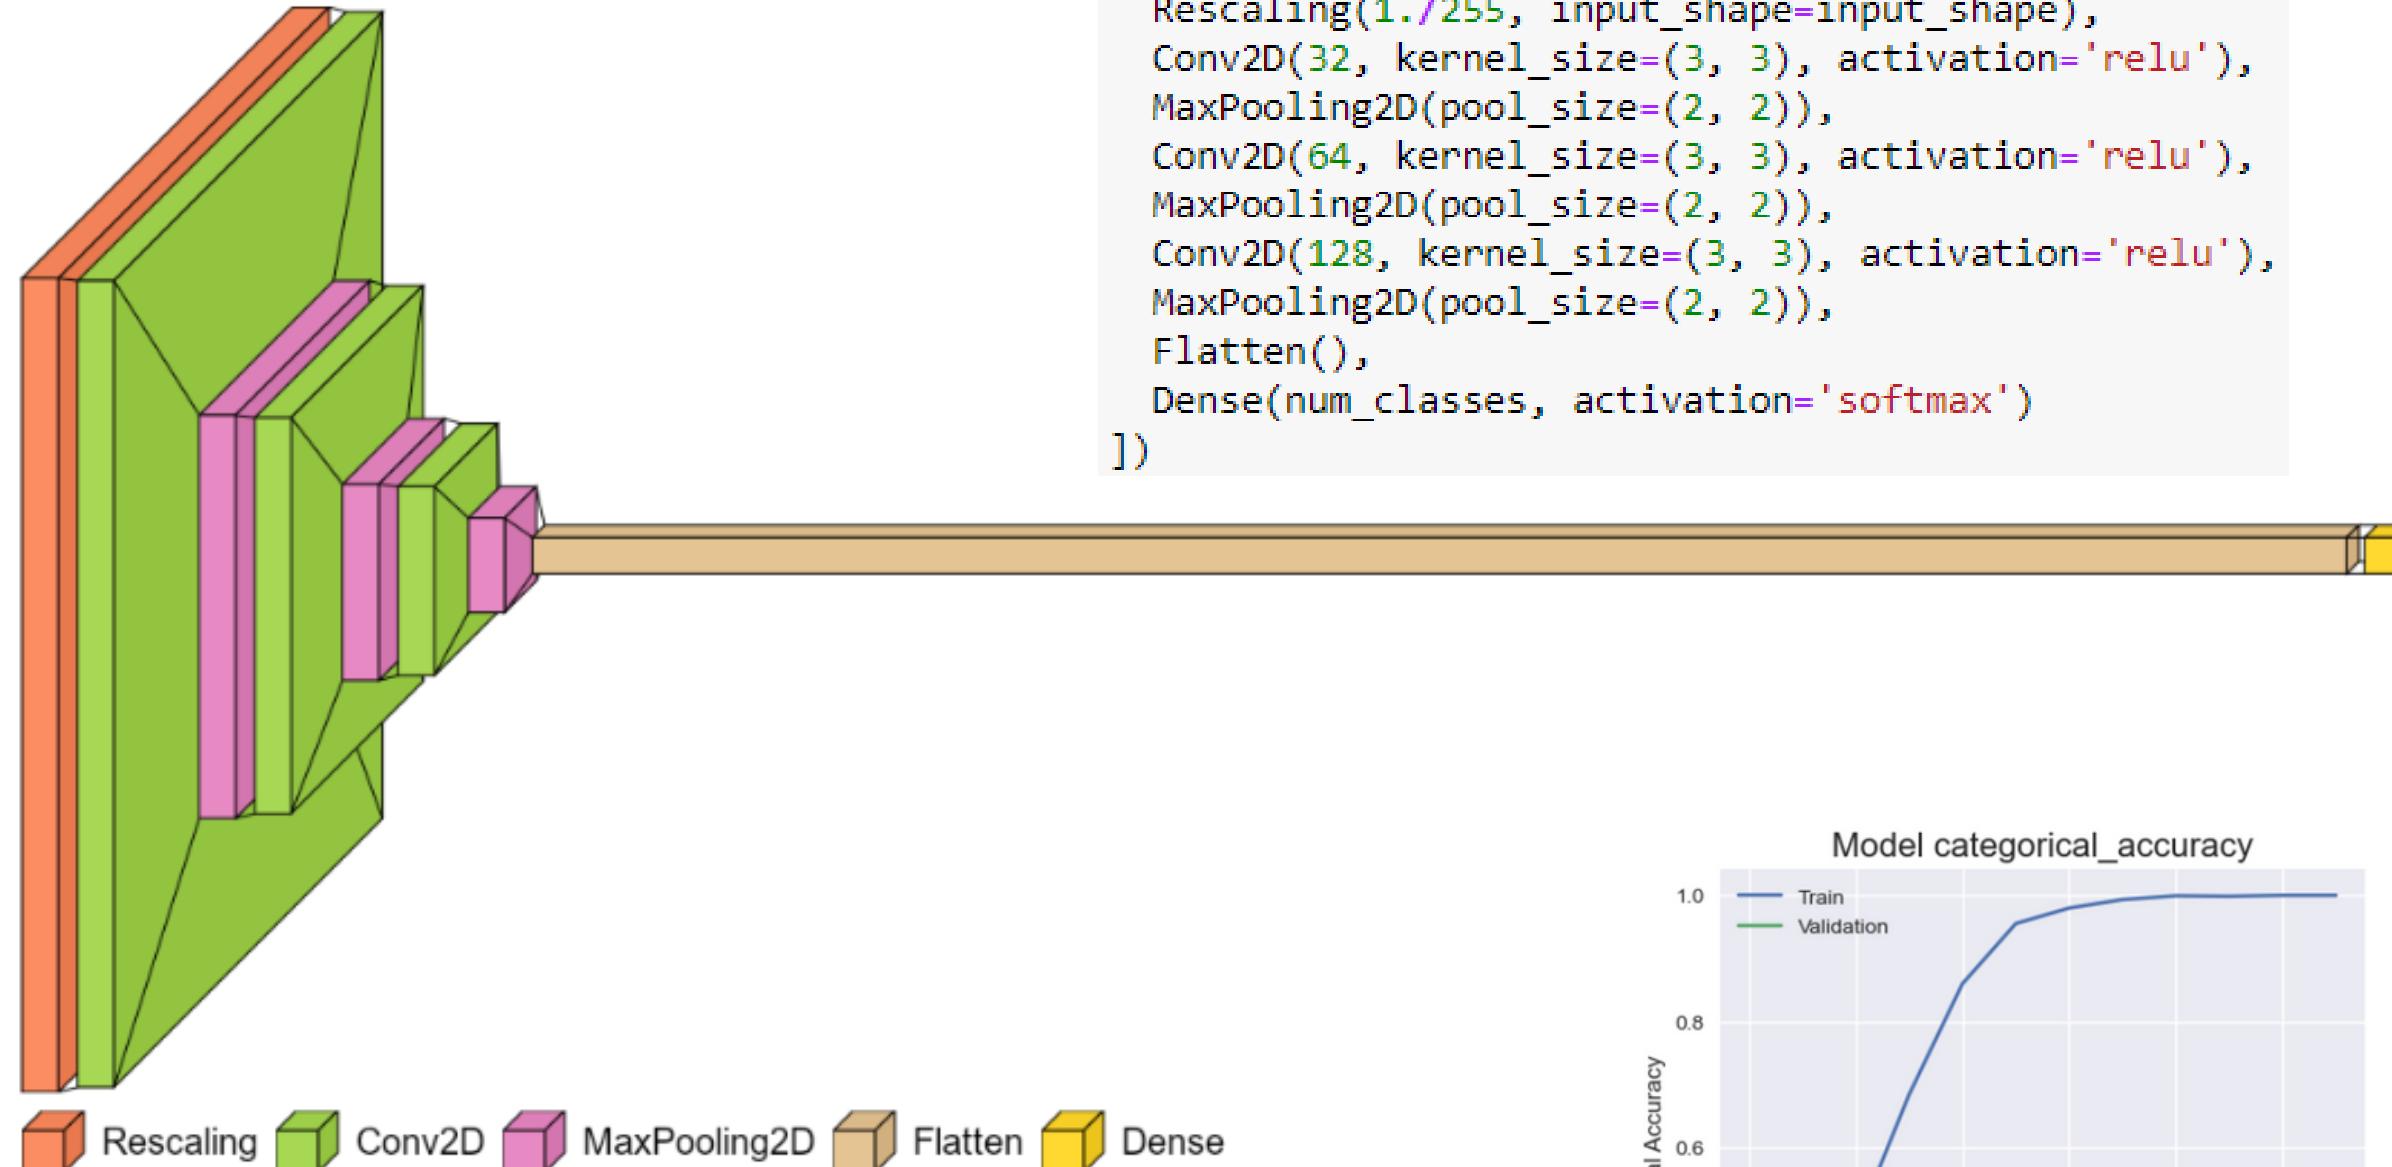
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)
```

## early\_stopping

- arrête l'entraînement du modèle lorsque certaines conditions sont remplies pour éviter le sur-apprentissage.
- restaure les poids du meilleur modèle

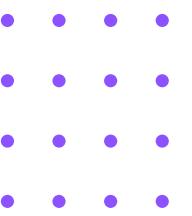


# Modèle de base



Test accuracy: 0.24137930572032928

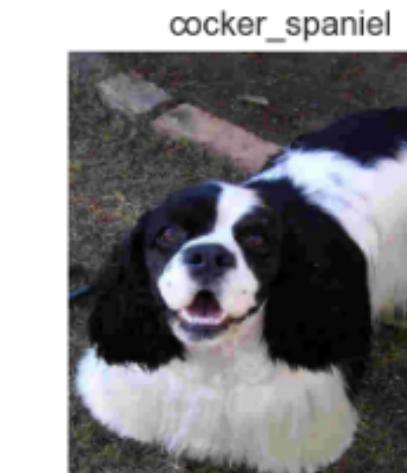
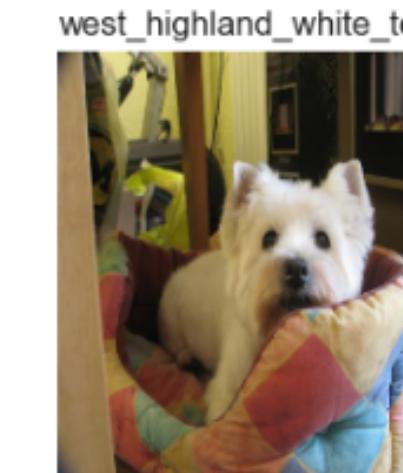
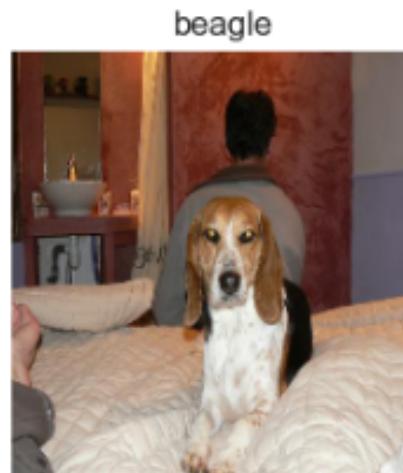
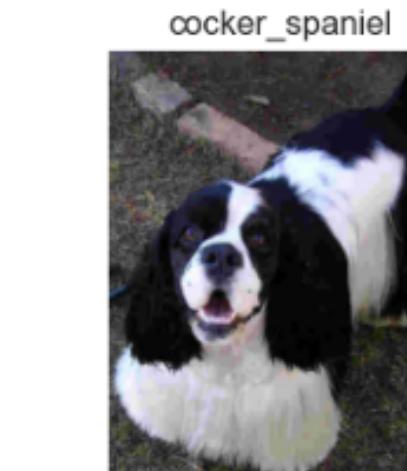
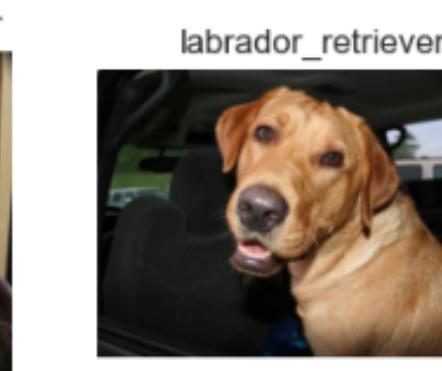




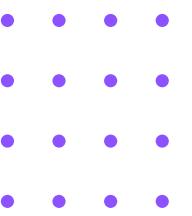
# Modèle de base

## Preprocessing des images

- Redimensionnement et normalisation des images



Objectif: améliore la qualité et l'efficacité de l'entraînement des algos

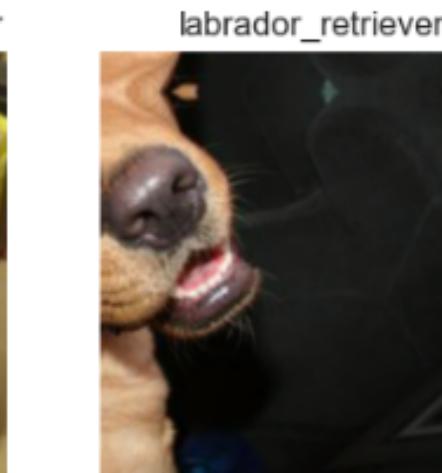
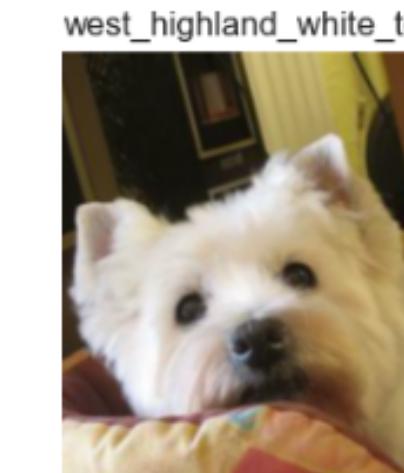
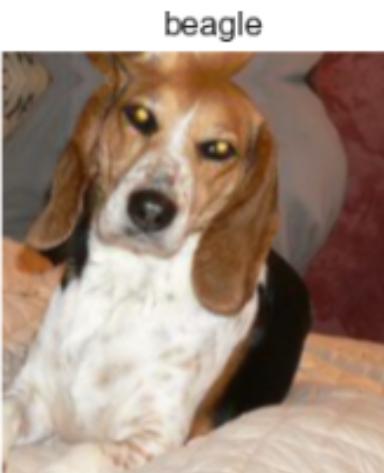
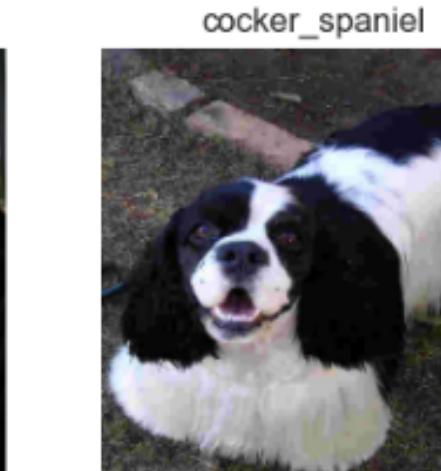
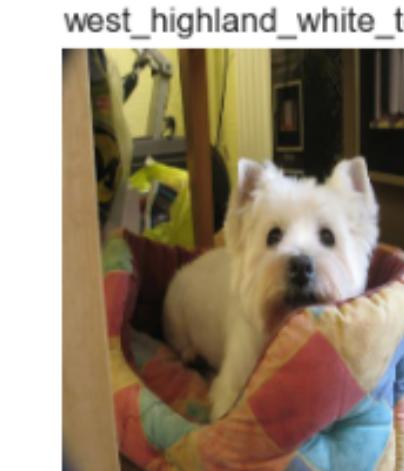


# Modèle de base

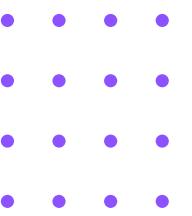
## Preprocessing des images

- **Data Augmentation**

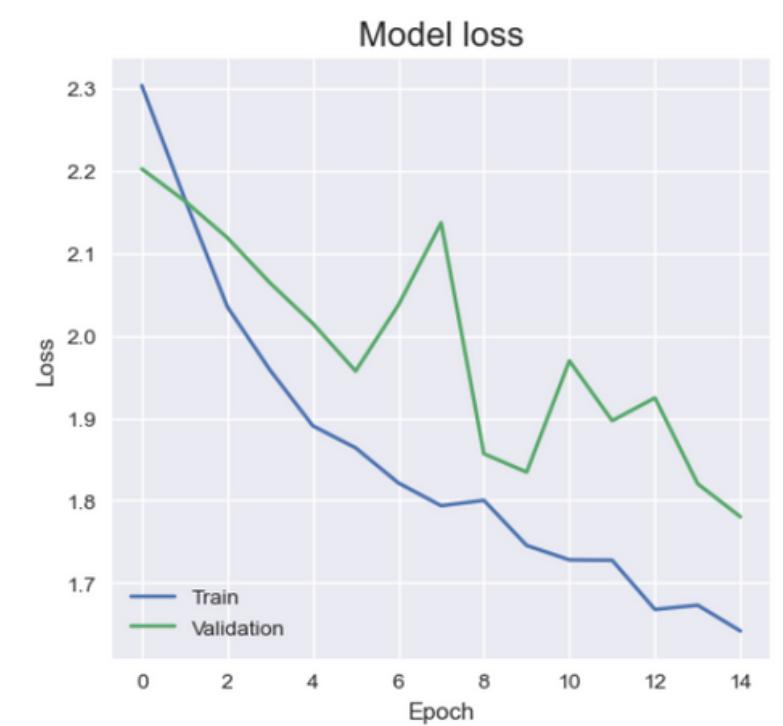
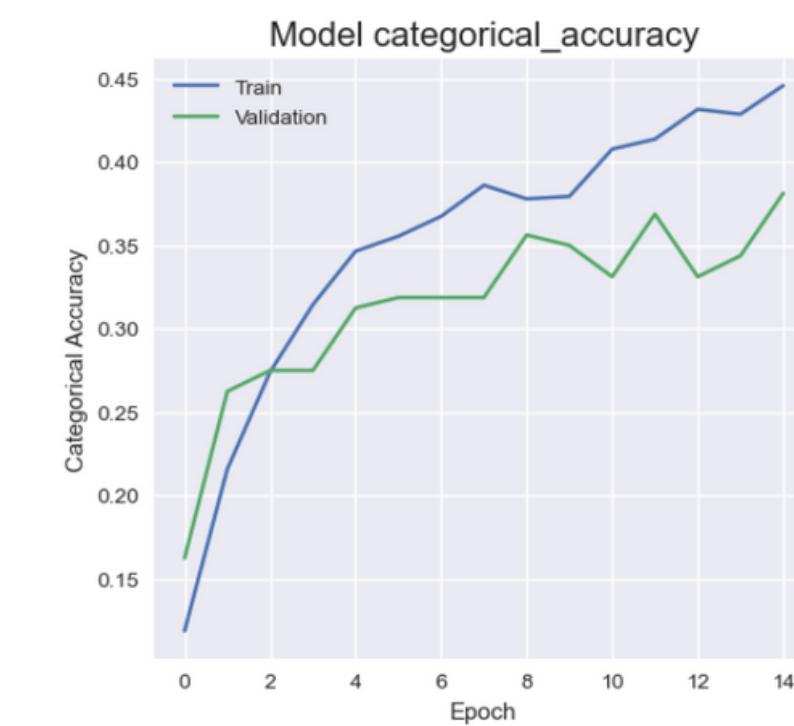
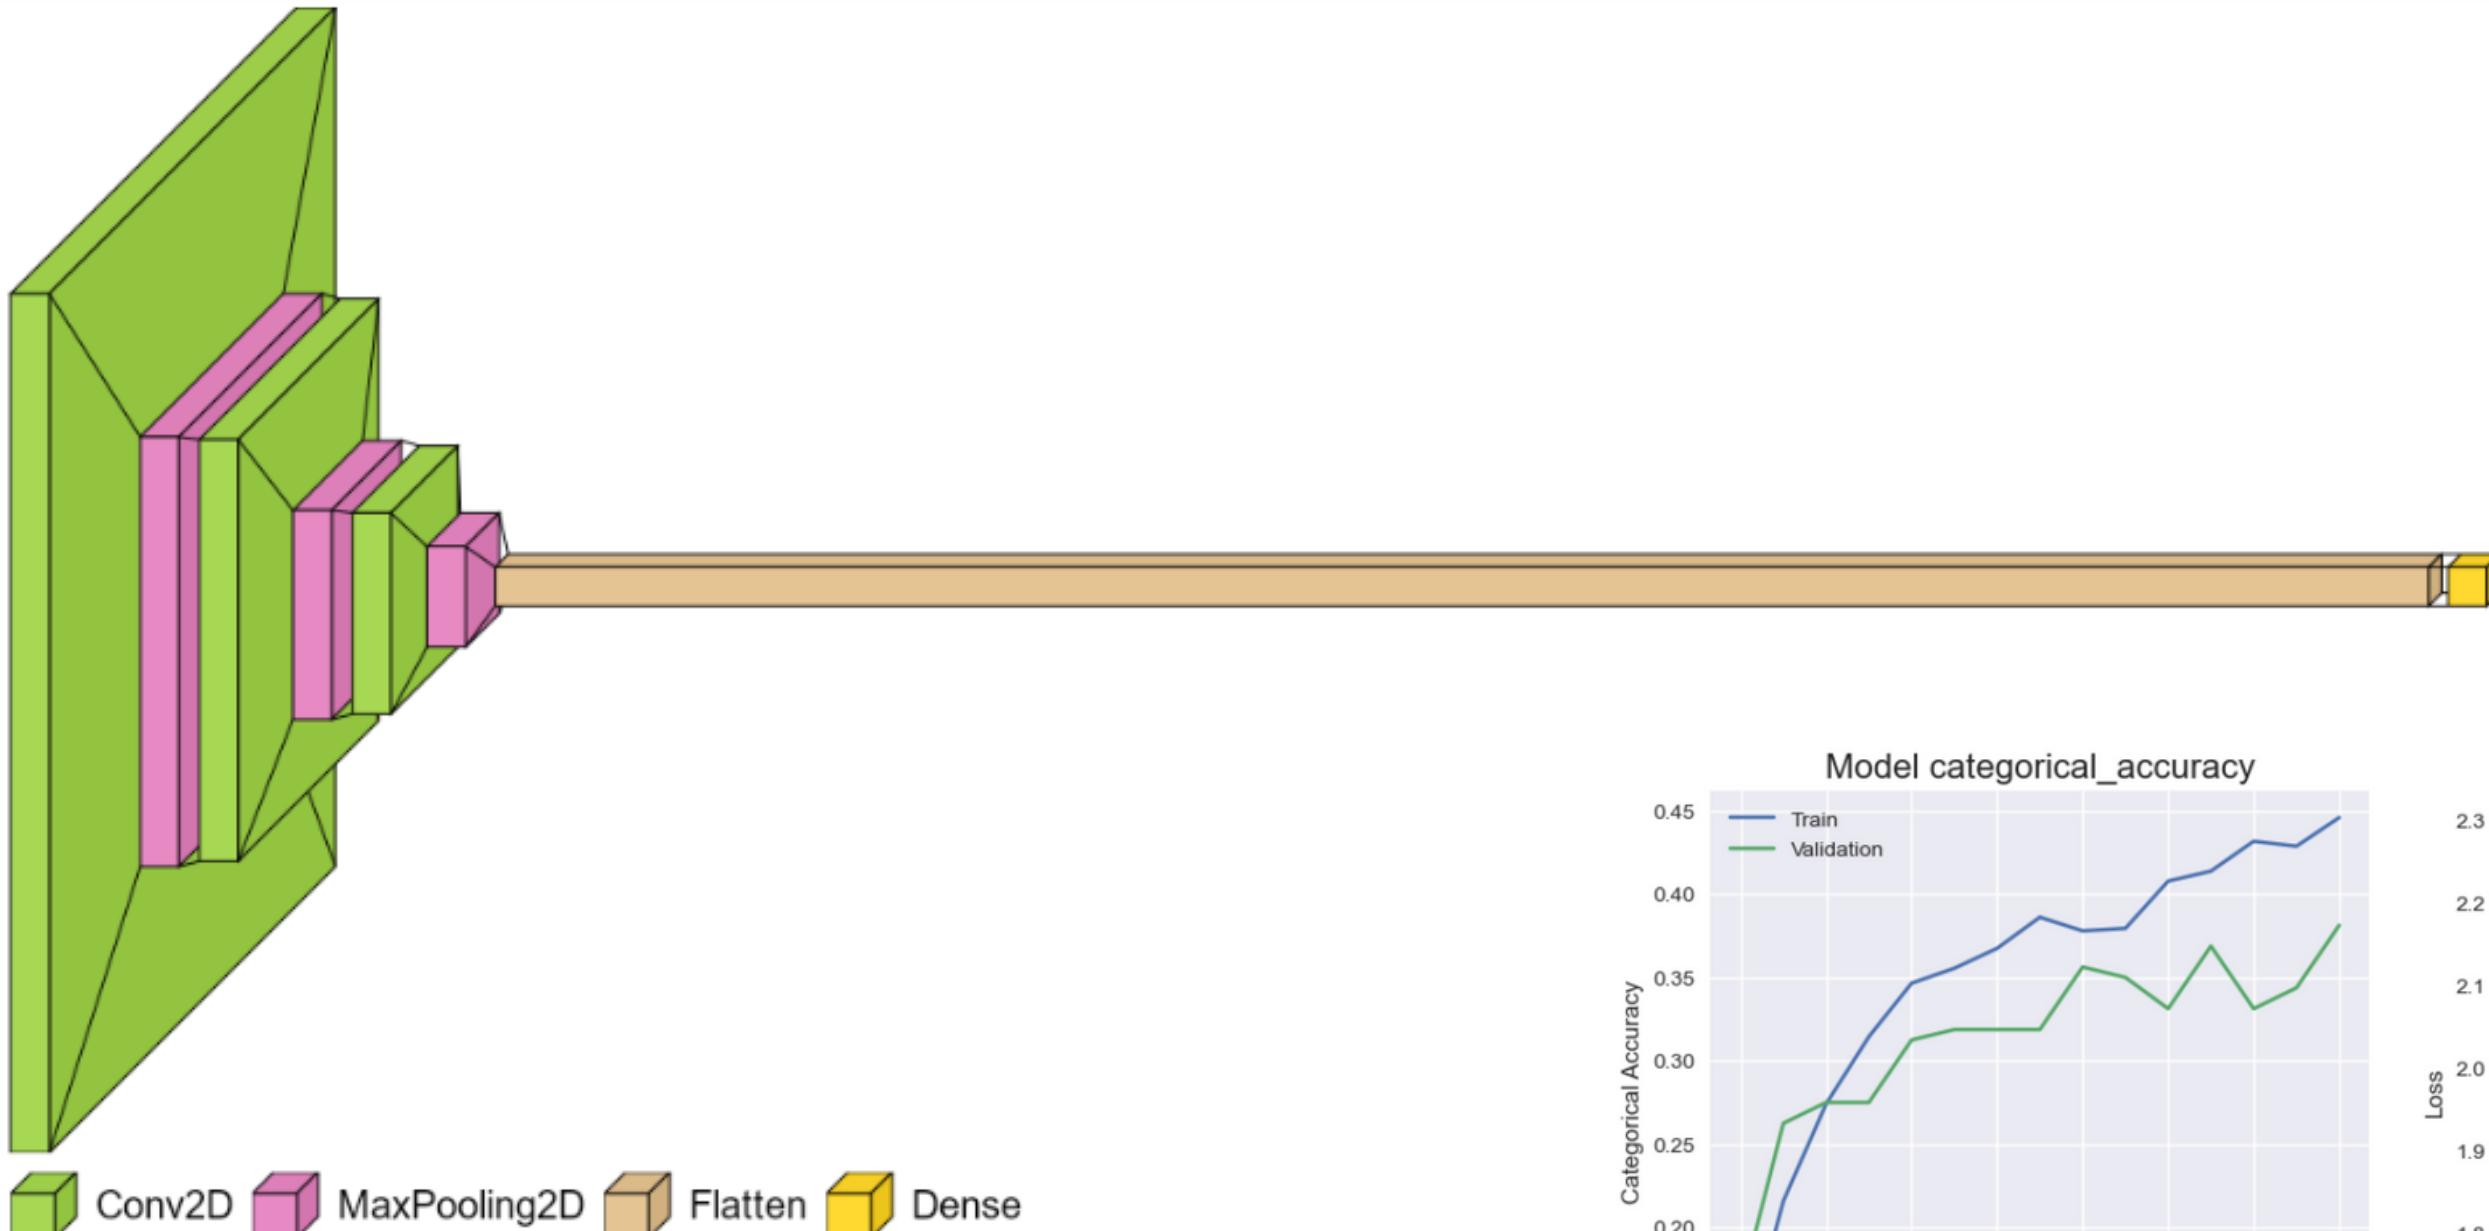
```
RandomFlip("horizontal"),
RandomRotation(0.1),
RandomZoom(0.1),
RandomContrast(0.1),
RandomCrop(img_size[0], img_size[1]),
RandomTranslation(height_factor=0.1, width_factor=0.1),
```



Objectif: apporte davantage de diversité au dataset

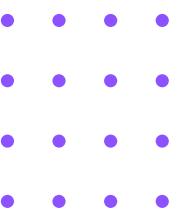


# Modèle de base



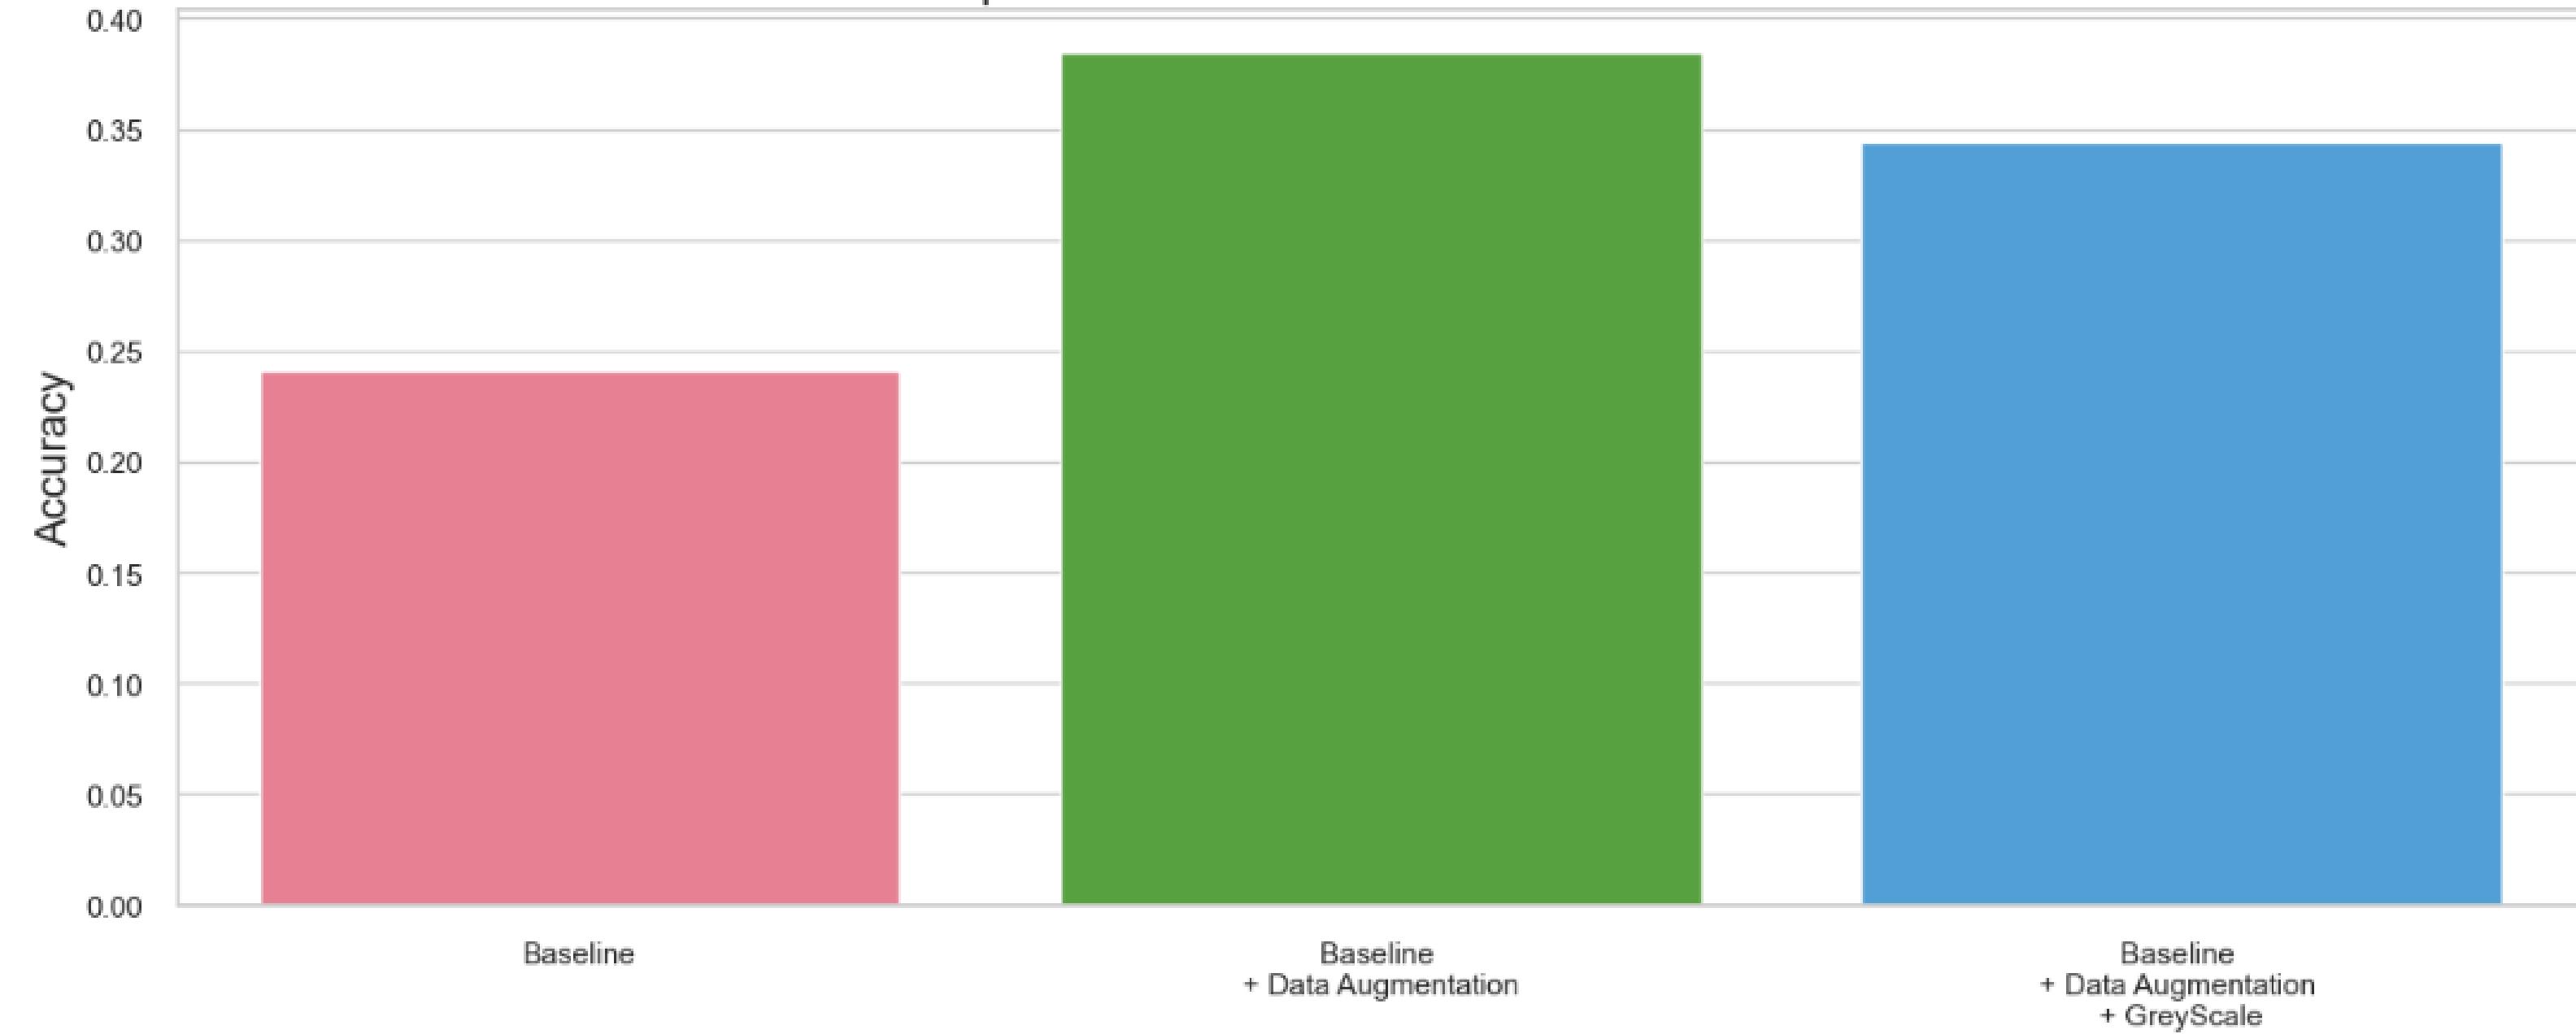
Test accuracy: 0.3850574791431427

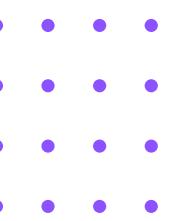




# Modèle de base

Comparaison des exactitudes des modèles

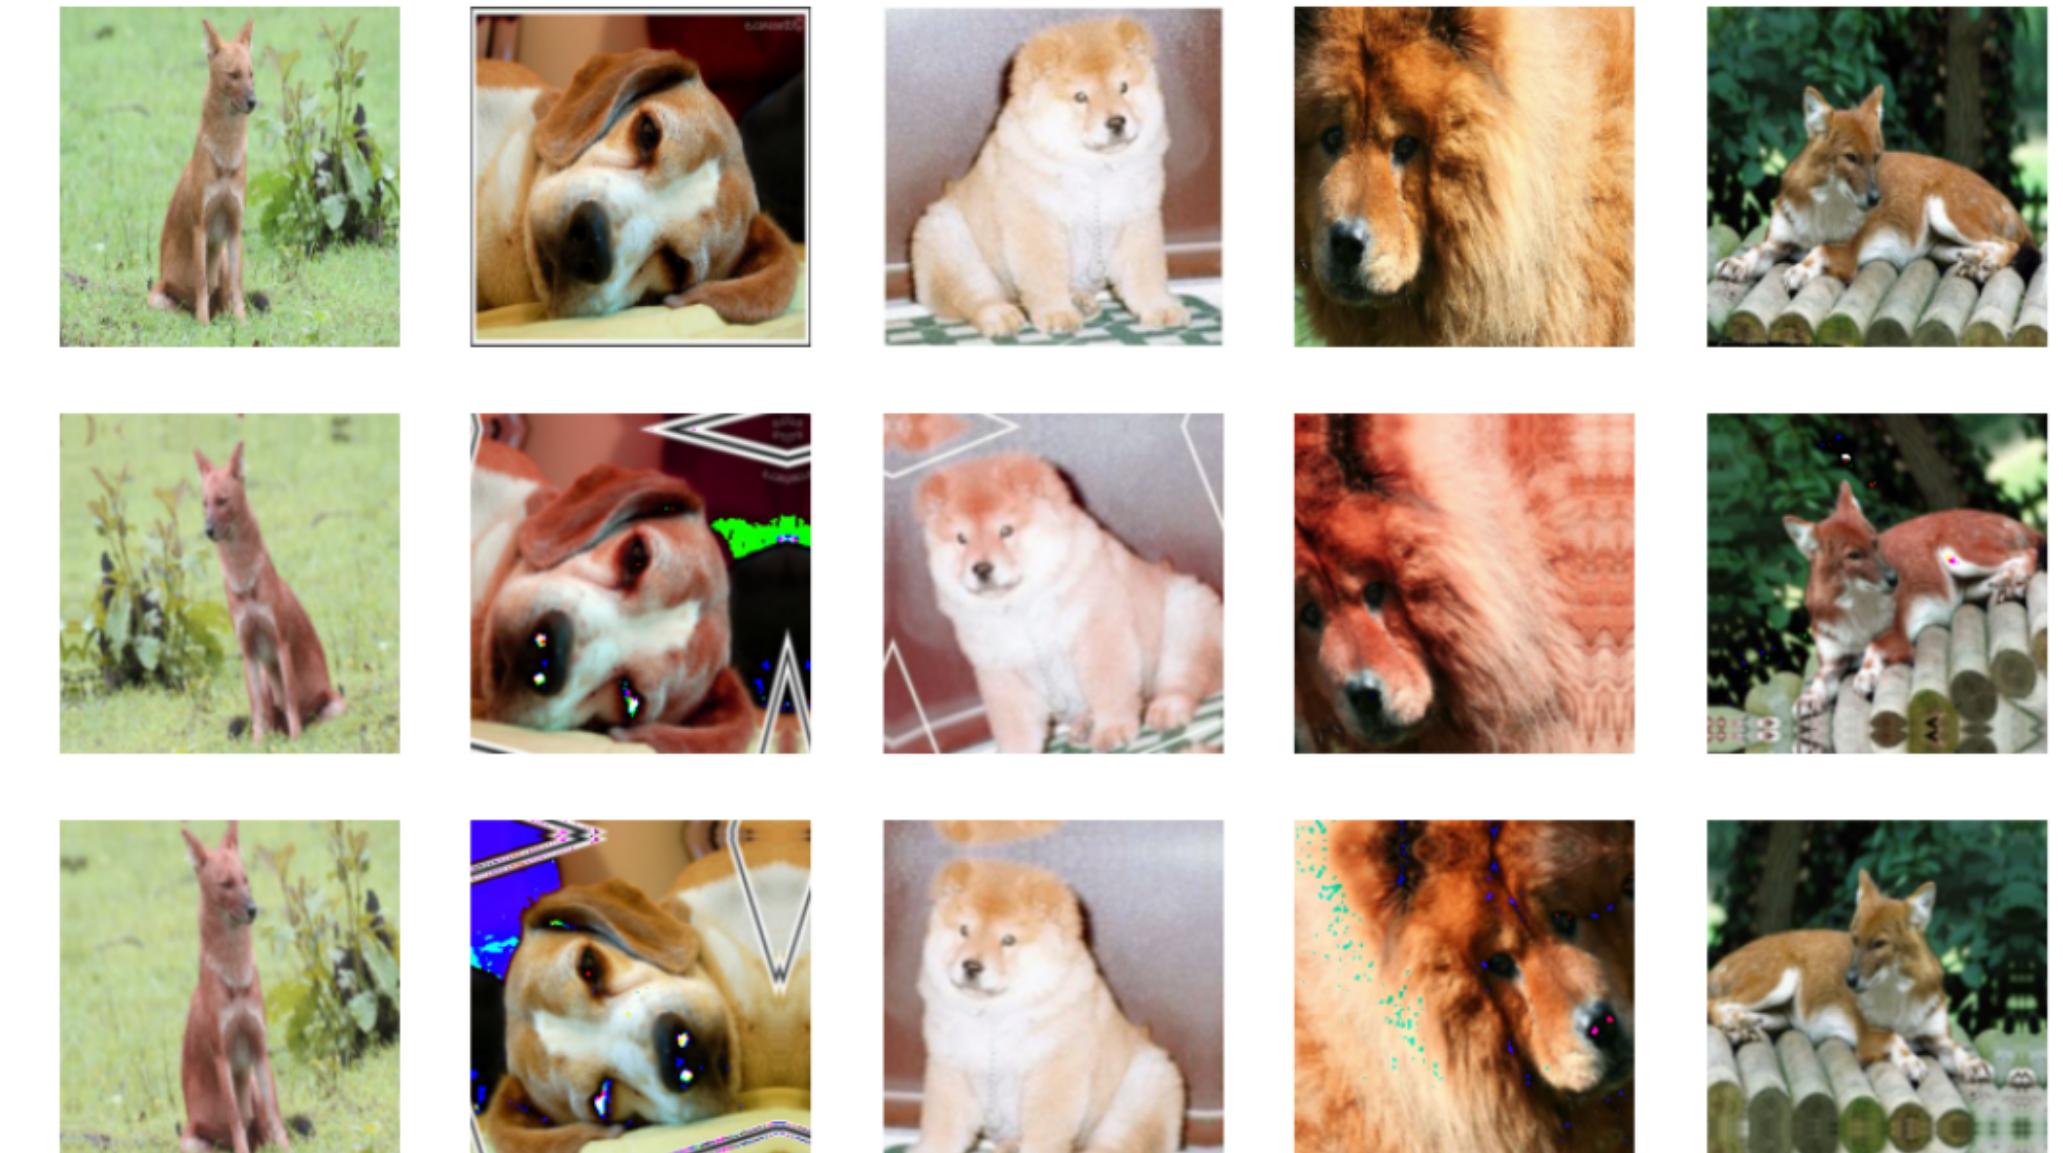




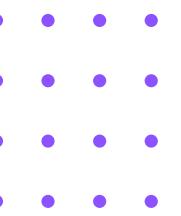
# Création d'un CNN "from scratch"

## Data Augmentation modifiée

```
data_augmentation = Sequential(
    [
        RandomFlip("horizontal"),
        RandomRotation(0.1),
        RandomZoom(0.1),
        RandomContrast(0.05),
        RandomCrop(img_size[0], img_size[1]),
        RandomTranslation(height_factor=0.1, width_factor=0.1),
    ]
)
```

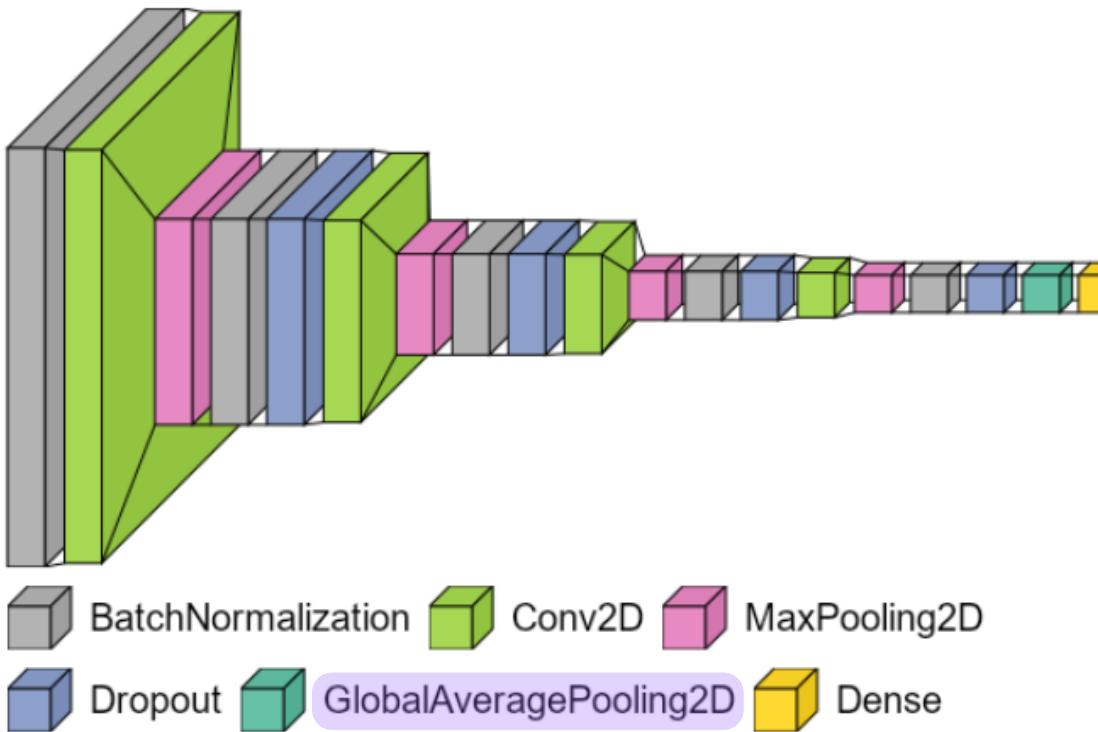


```
def preprocess_image_with_augmentation(image, label):
    image = tf.image.resize(image, img_size)
    image = tf.image.random_brightness(image, 1.0)
    image = tf.image.random_hue(image, 0.05)
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_saturation(image, 0.95, 1.05)
    image = data_augmentation(image)
    image = tf.cast(image, tf.float32) / 255.0
    return image, label
```



# Création d'un CNN "from scratch"

## Hyperparamétrage à la main à partir du modèle AlexNet



### BatchNormalization

- normalise les activations de chaque couche
- aider à stabiliser le processus d'apprentissage

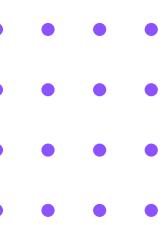
### Dropout

- désactive aléatoirement certains neurones
- réduire le sur-apprentissage en réduisant les interdépendances entre les neurones.

### GlobalAveragePooling2D

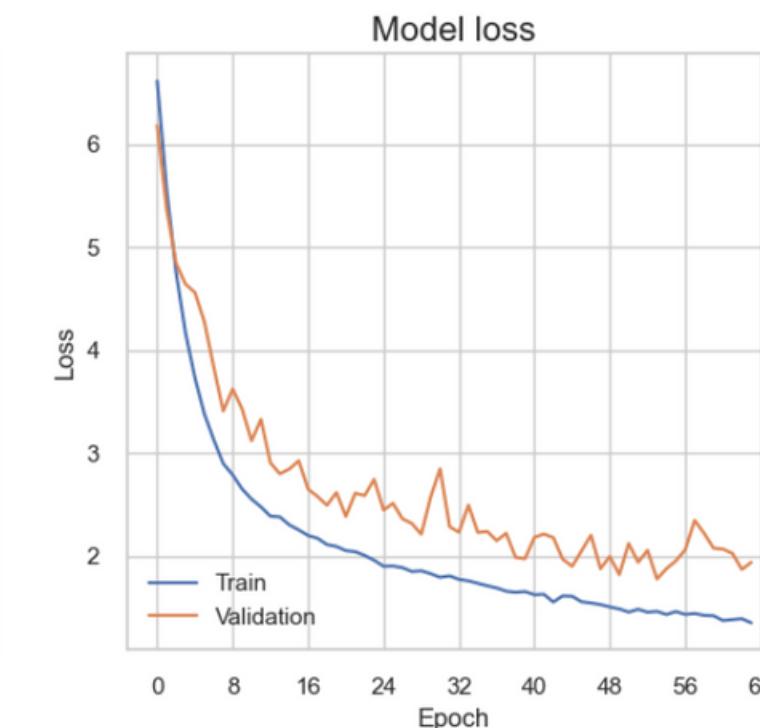
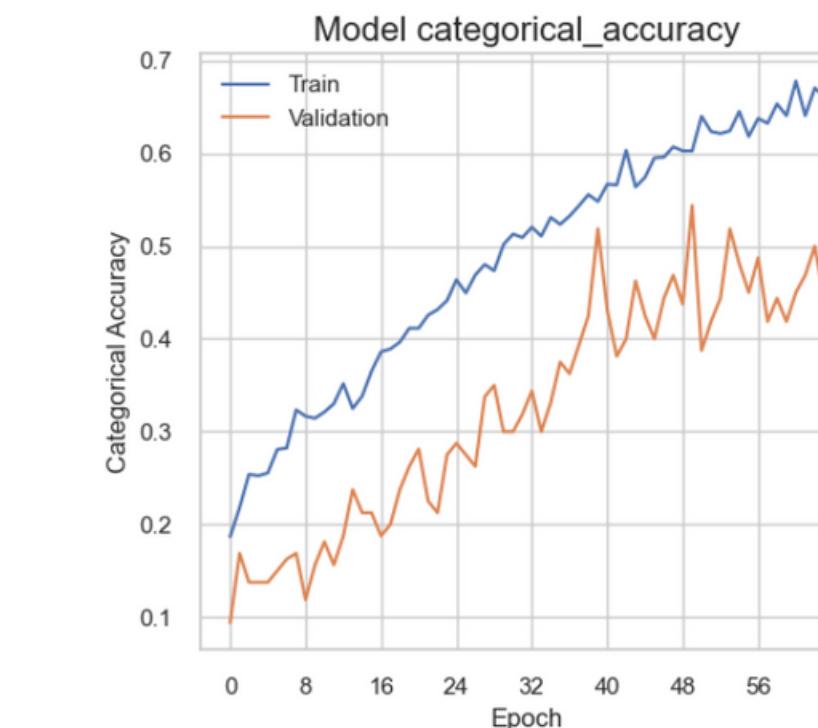
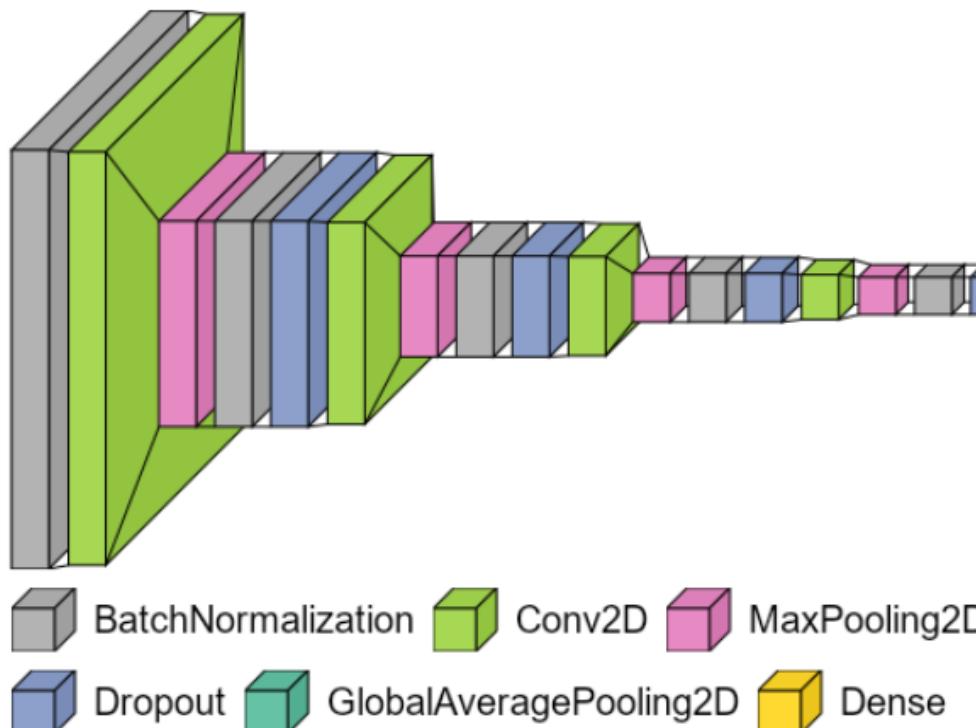
- calcule la moyenne des activations de chaque carte de caractéristiques
- permet de réduire la nb de paramètres (contre l'overfitting)





# Création d'un CNN "from scratch"

## Hyperparamétrage à la main à partir du modèle AlexNet

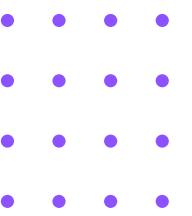


## Definition des hyperparamètres à rechercher

```
param_grid = {
    'num_filters': [[16, 32, 64, 128], [32, 64, 128, 256]],
    'k_sizes': [[5, 3, 3, 3], [4, 3, 2, 3], [5, 3, 3, 2], [3, 3, 3, 3]],
    'k_init': ['he_normal', 'glorot_uniform', 'random_normal'],
    'l2_value': [0.01, 0.02, 0.03],
    'pool_sizes': [[(2, 2), (2, 2), (2, 2), (2, 2)], [(3, 3), (2, 2), (2, 2), (2, 2)]],
    'dropout': [0.20, 0.25, 0.5, 0.75]
}
```

Nombre de combinaisons : 576

Utilisation GPU ~40 heures

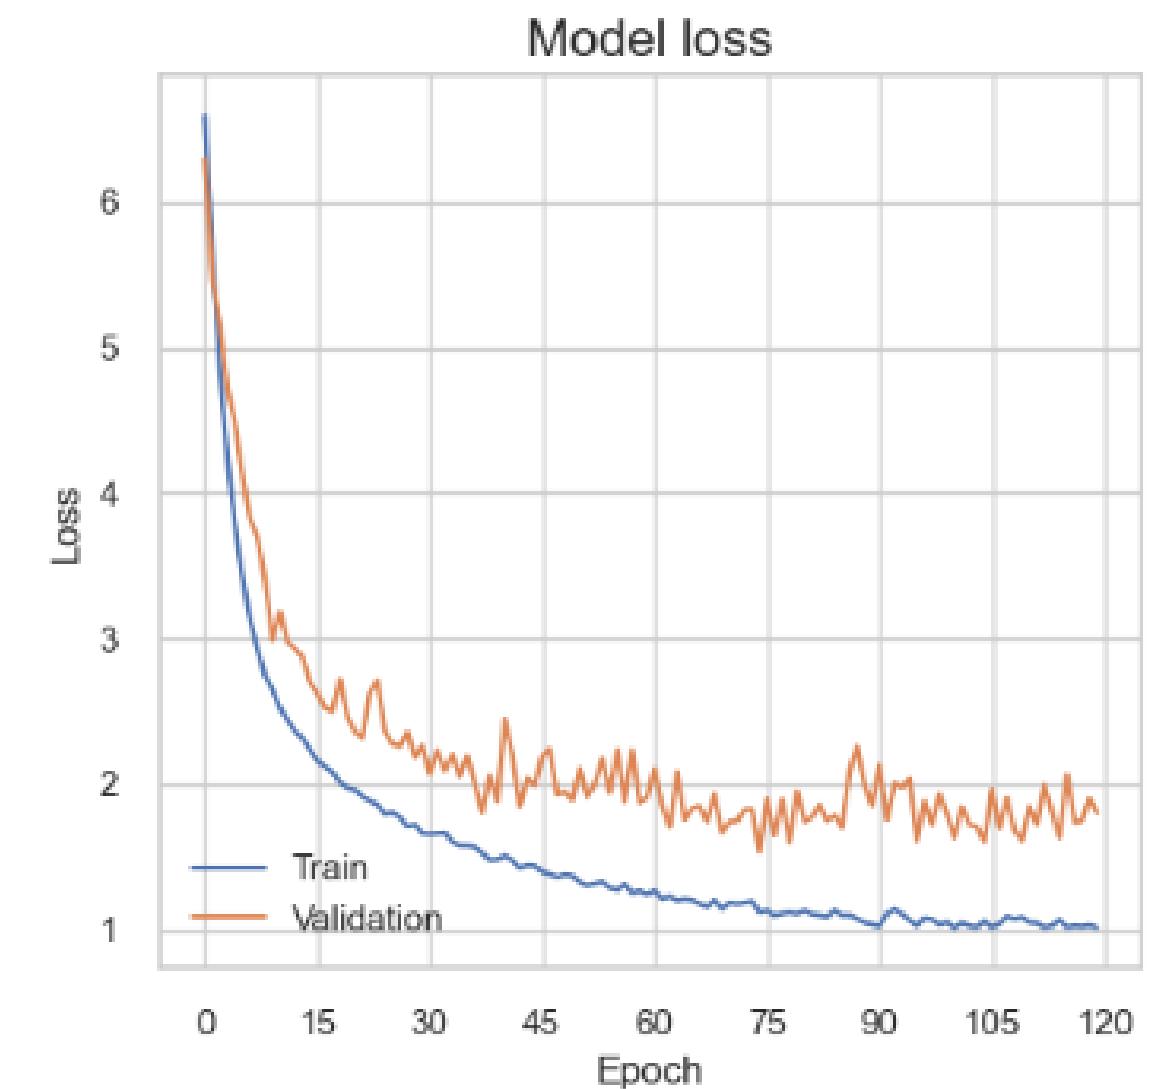
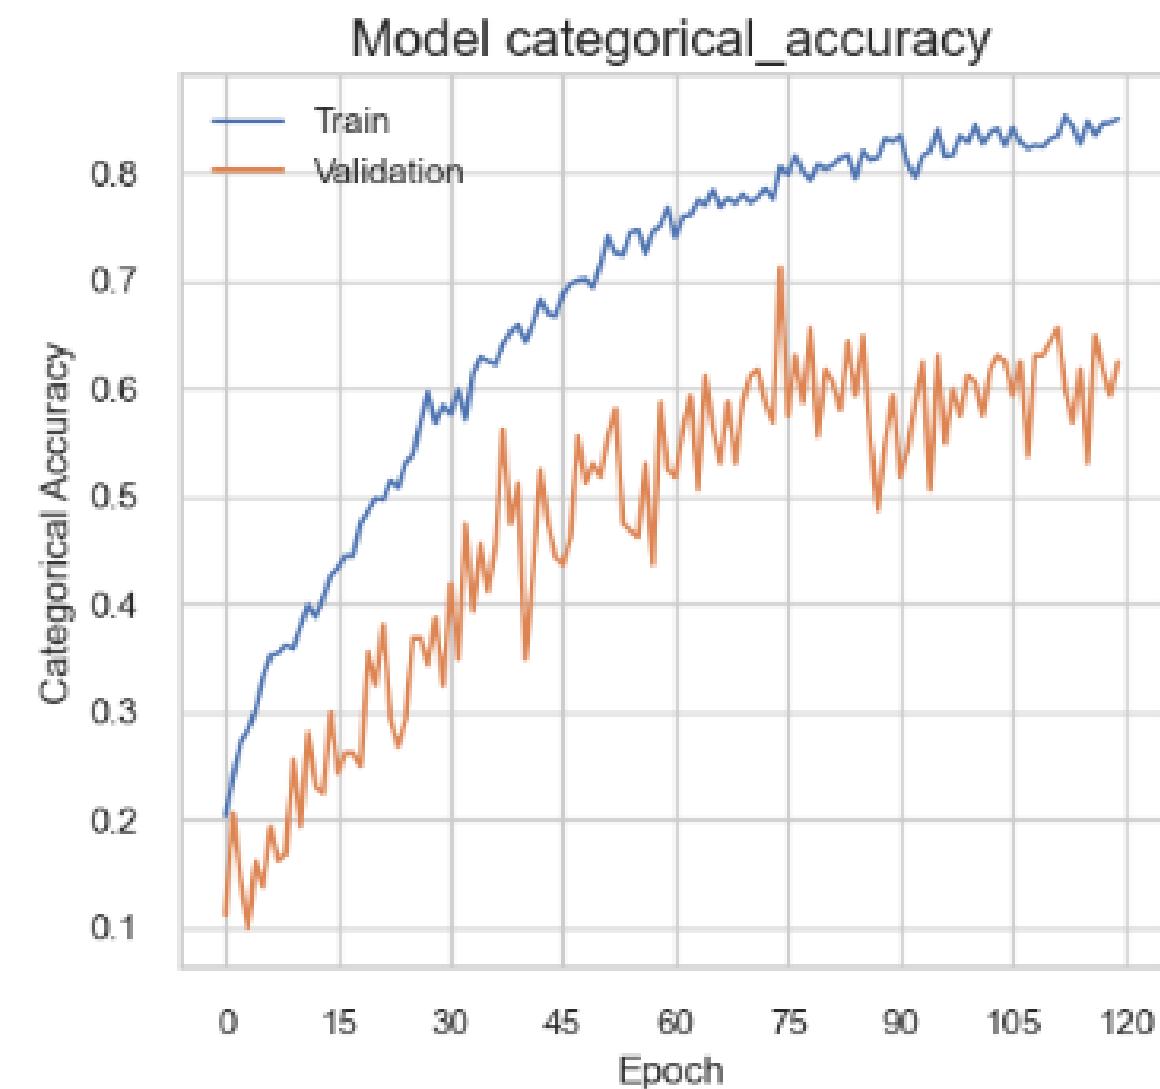


# Création d'un CNN "from scratch"

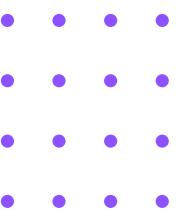
## Meilleurs hyperparamètres

Best hyperparameters: ([16, 32, 64, 128], [5, 3, 3, 3], 'he\_normal', 0.01, [(3, 3), (2, 2), (2, 2), (2, 2)], 0.25)  
Best validation accuracy: 0.59

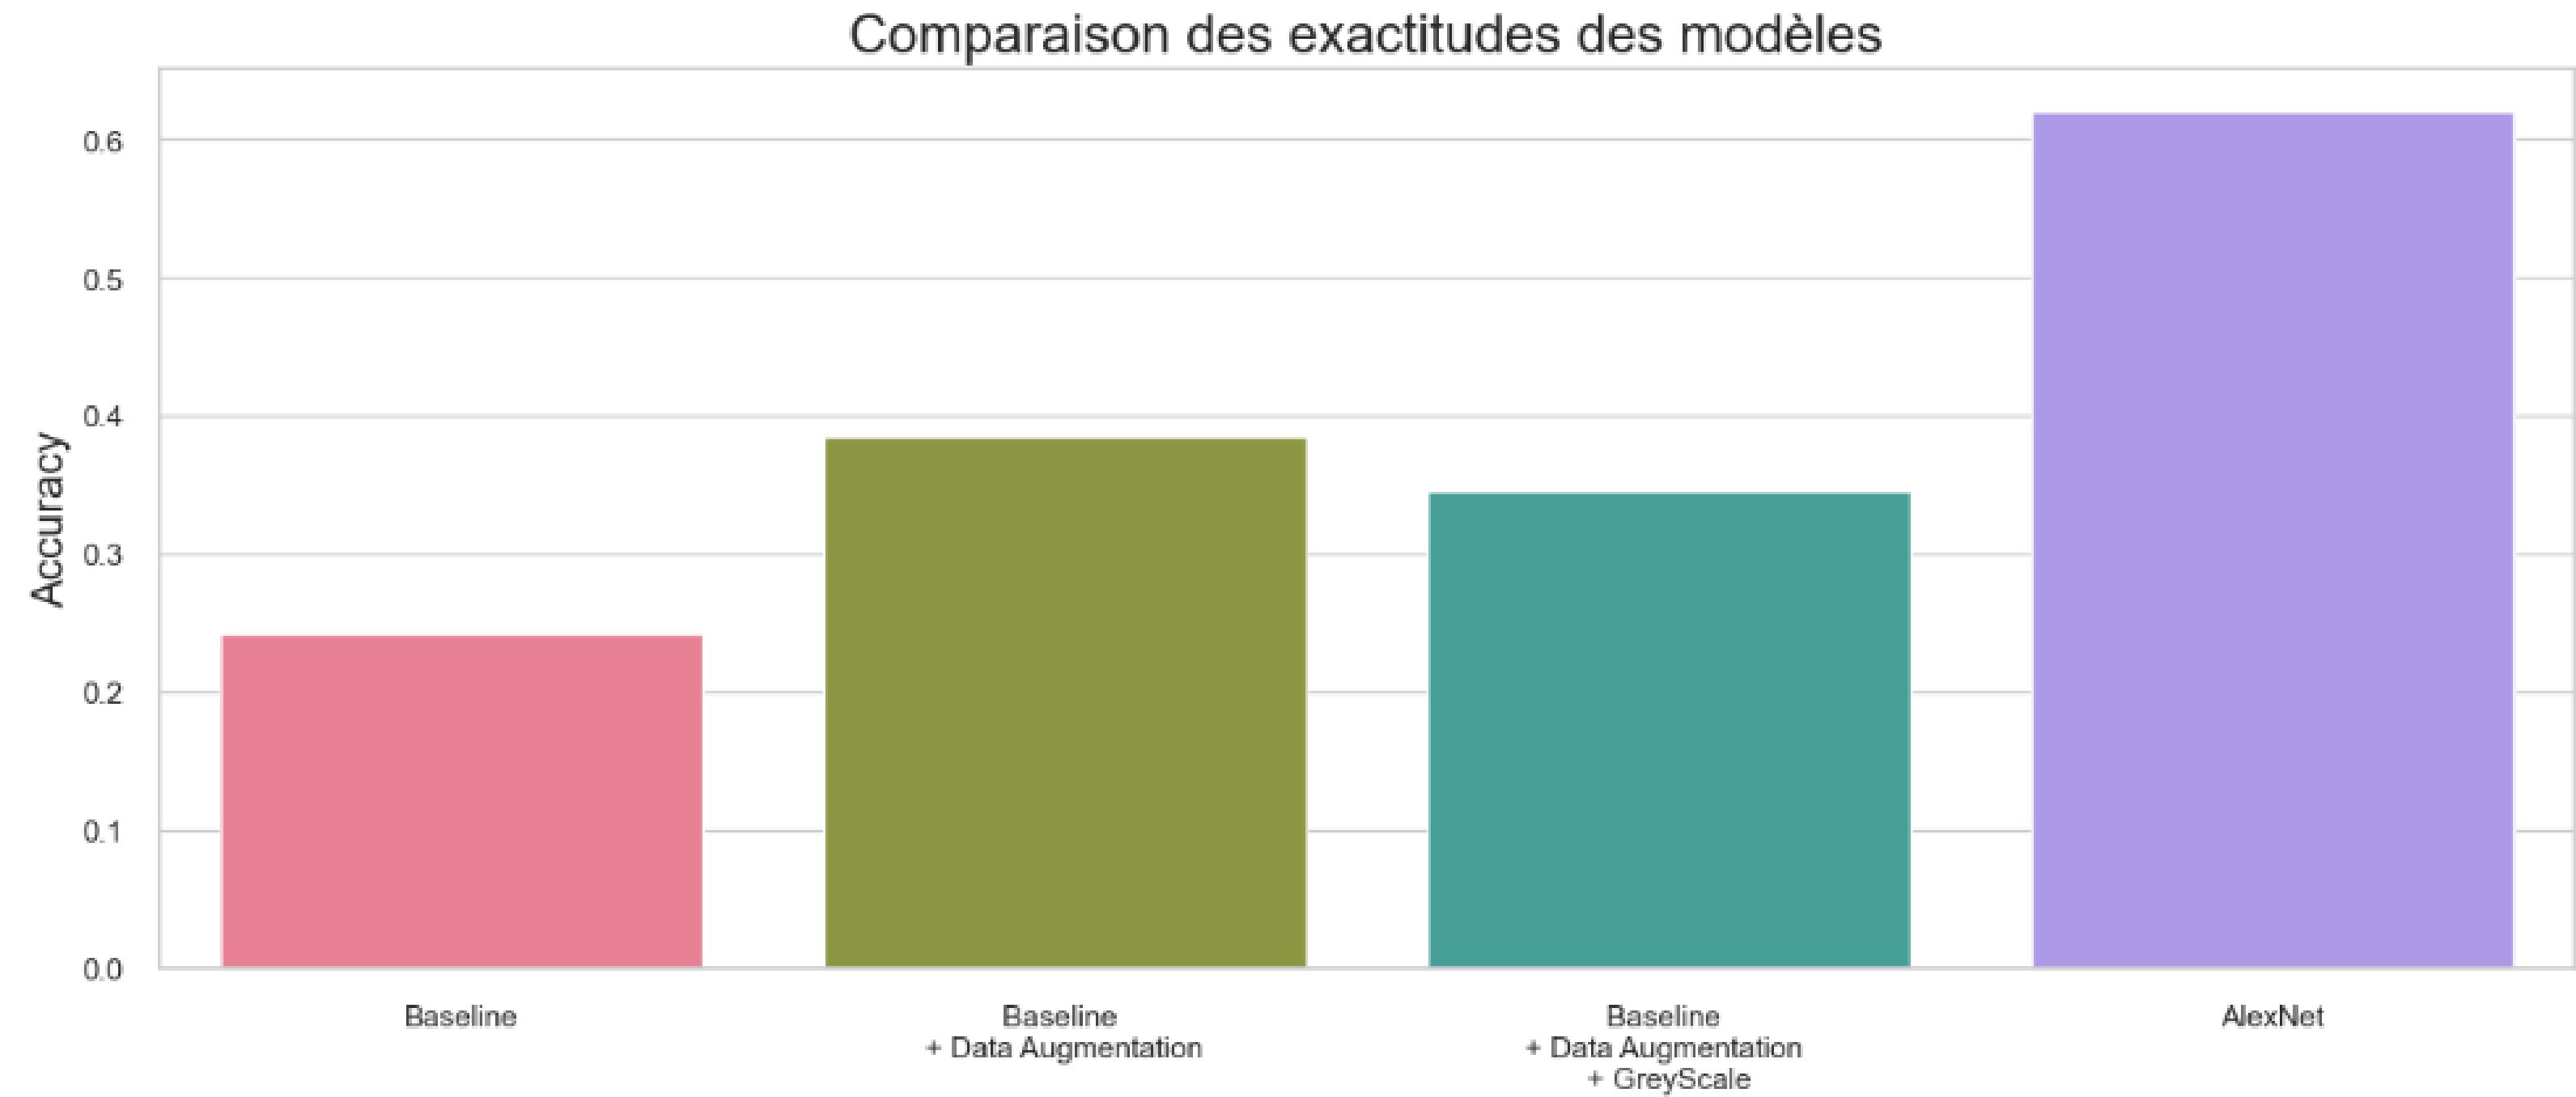
Layer (type)	Output Shape	Param #
<hr/>		
batch_normalization_115 (BatchNormalization)	(None, 224, 224, 3)	12
conv2d_101 (Conv2D)	(None, 220, 220, 16)	1216
max_pooling2d_101 (MaxPooling2D)	(None, 73, 73, 16)	0
batch_normalization_116 (BatchNormalization)	(None, 73, 73, 16)	64
dropout_92 (Dropout)	(None, 73, 73, 16)	0
conv2d_102 (Conv2D)	(None, 71, 71, 32)	4640
max_pooling2d_102 (MaxPooling2D)	(None, 35, 35, 32)	0
batch_normalization_117 (BatchNormalization)	(None, 35, 35, 32)	128
dropout_93 (Dropout)	(None, 35, 35, 32)	0
conv2d_103 (Conv2D)	(None, 33, 33, 64)	18496
max_pooling2d_103 (MaxPooling2D)	(None, 16, 16, 64)	0
batch_normalization_118 (BatchNormalization)	(None, 16, 16, 64)	256
dropout_94 (Dropout)	(None, 16, 16, 64)	0
conv2d_104 (Conv2D)	(None, 14, 14, 128)	73856
max_pooling2d_104 (MaxPooling2D)	(None, 7, 7, 128)	0
batch_normalization_119 (BatchNormalization)	(None, 7, 7, 128)	512
dropout_95 (Dropout)	(None, 7, 7, 128)	0
global_average_pooling2d_23 (GlobalAveragePooling2D)	(None, 128)	0
dense_26 (Dense)	(None, 10)	1290
<hr/>		
Total params: 100,470		
Trainable params: 99,984		
Non-trainable params: 486		

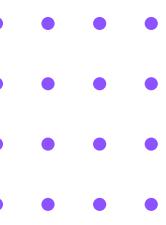


Test accuracy: 0.6206896305084229



# Création d'un CNN "from scratch"

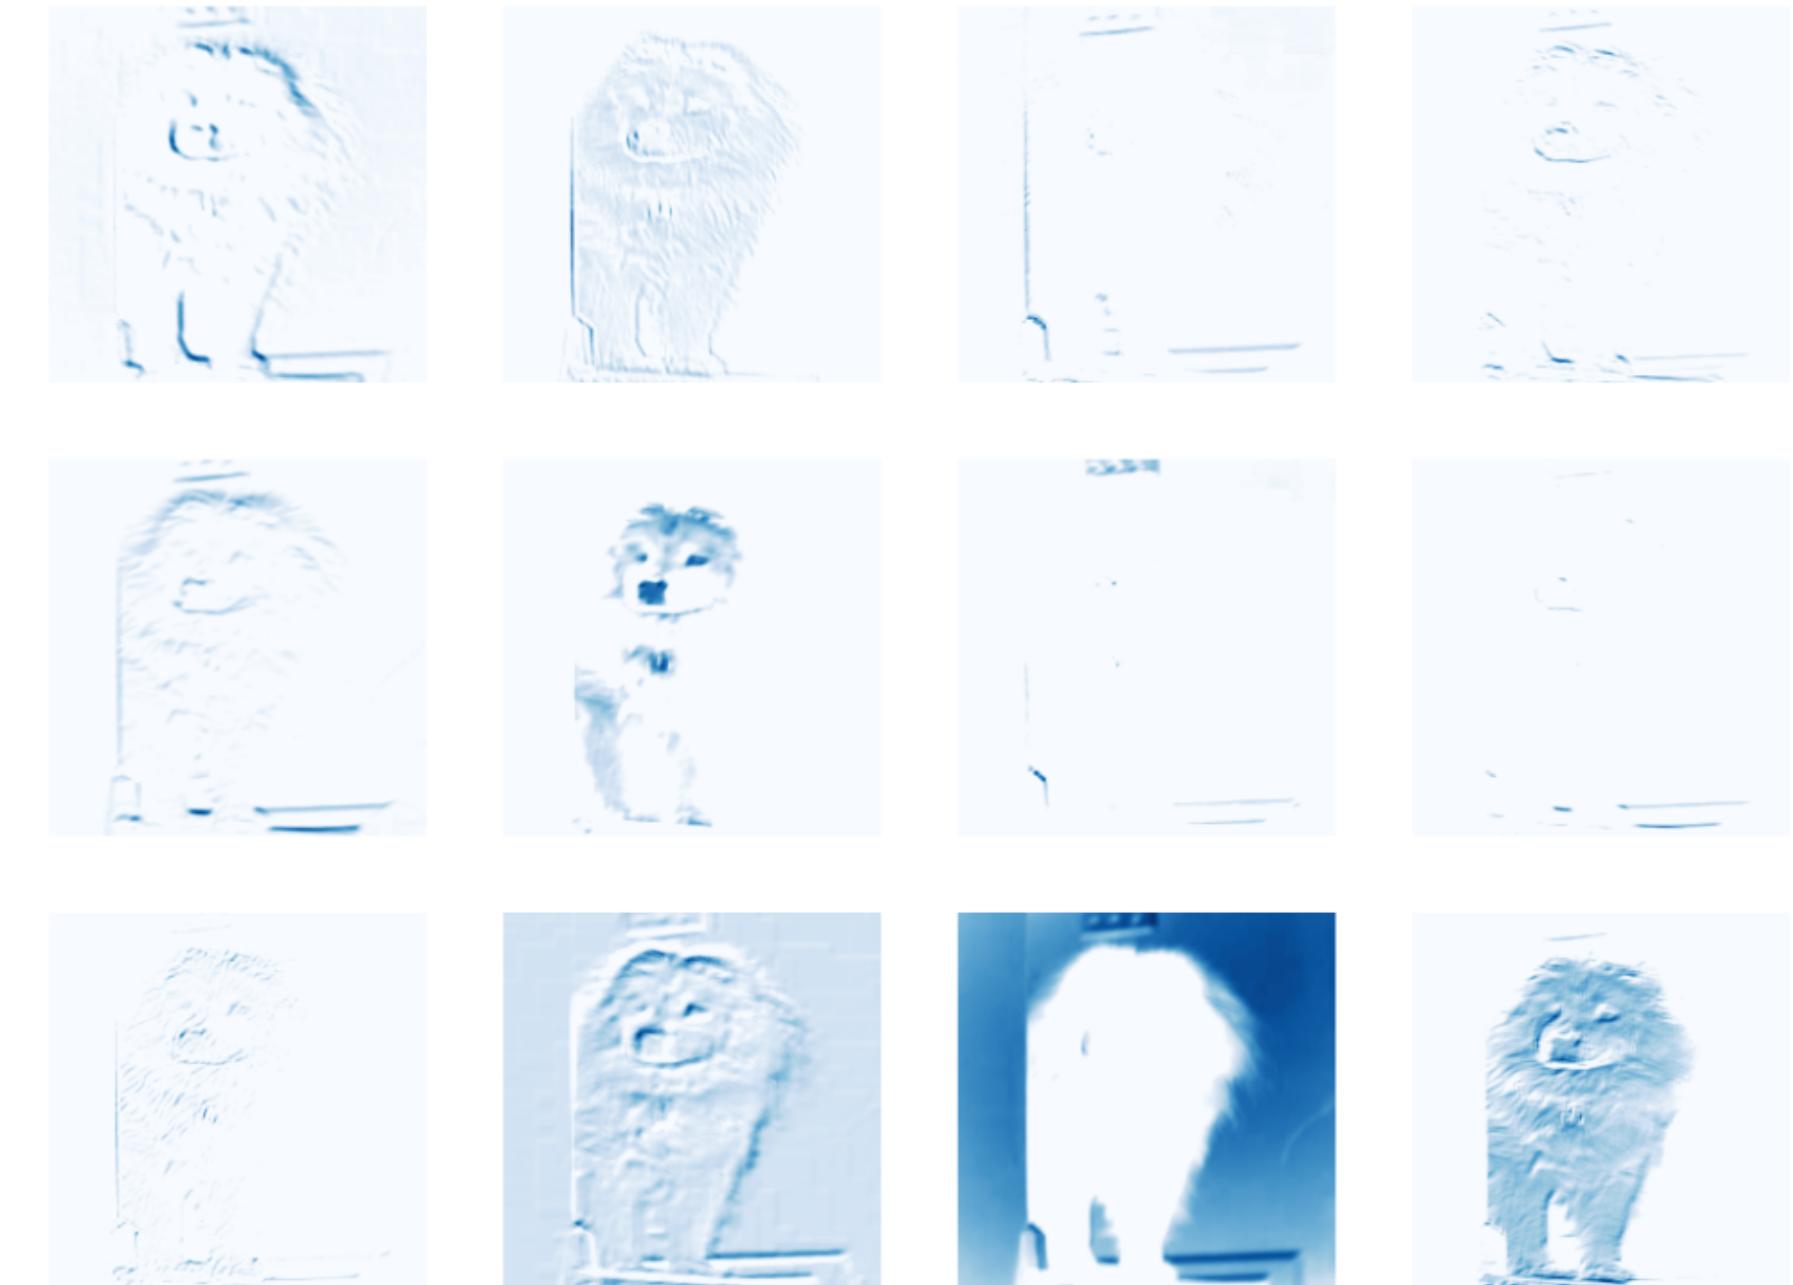


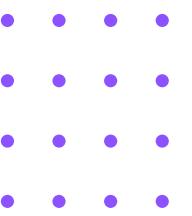


# Création d'un CNN "from scratch"

## Observation de 12 Features Maps (cartes de caractéristiques)

- **1ère convolution  
(contours, bords)**



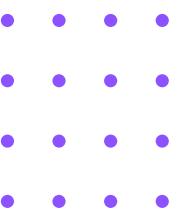


# Création d'un CNN "from scratch"

## Observation de 12 Features Maps (cartes de caractéristiques)

- **2ème convolution**

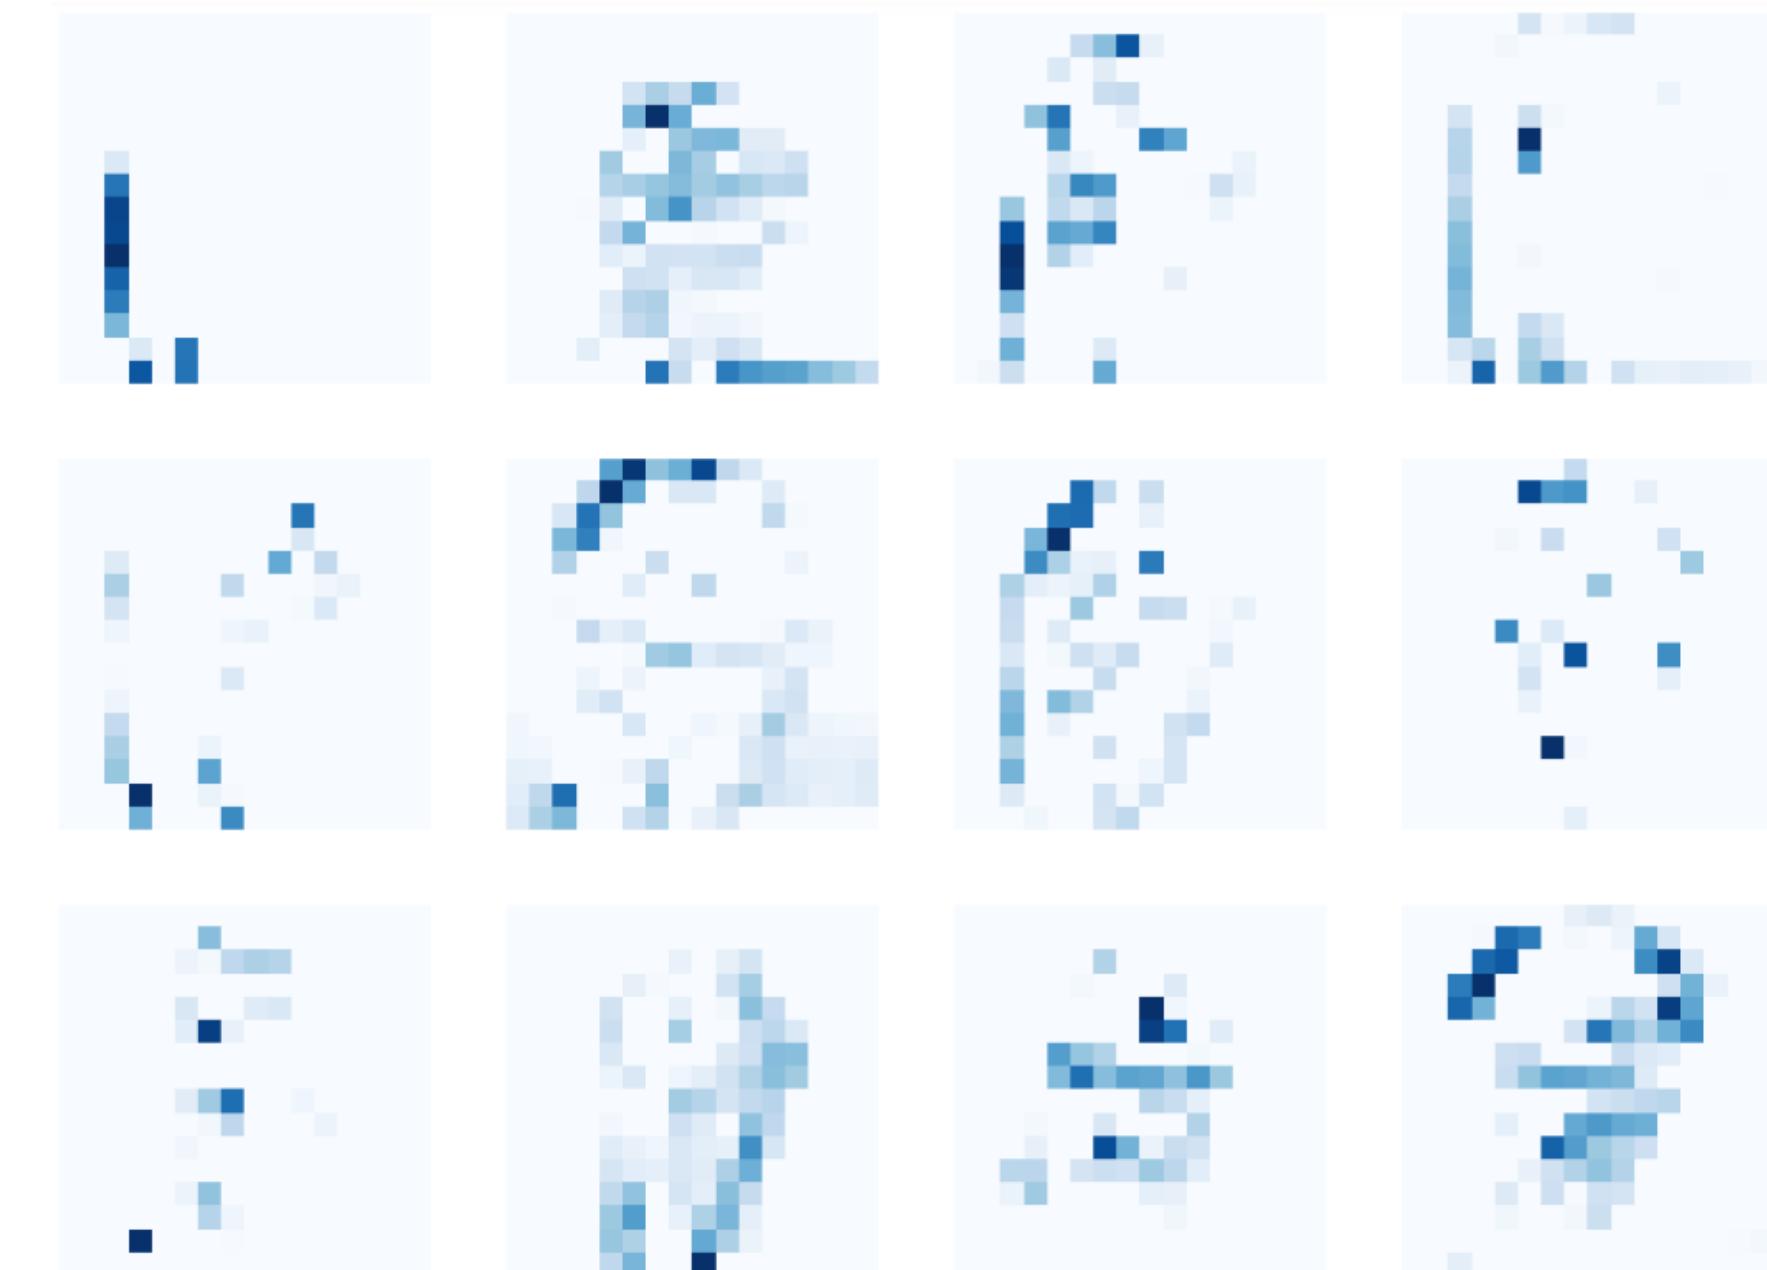


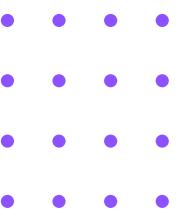


# Création d'un CNN "from scratch"

## Observation de 12 Features Maps (cartes de caractéristiques)

- **3ème convolution  
(plus abstrait)**





# Création d'un CNN "from scratch"

## Analyse des résultats

True: malinois  
Pred: Greater\_Swiss\_Mountain\_dog



True: Labrador\_retriever  
Pred: Greater\_Swiss\_Mountain\_dog



True: Labrador\_retriever  
Pred: Labrador\_retriever



True: chow  
Pred: chow



True: cocker\_spaniel  
Pred: cocker\_spaniel



True: malinois  
Pred: bull\_mastiff



True: chow  
Pred: Labrador\_retriever



True: West\_Highland\_white\_terrier  
Pred: West\_Highland\_white\_terrier

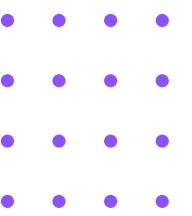


True: dhole  
Pred: dhole



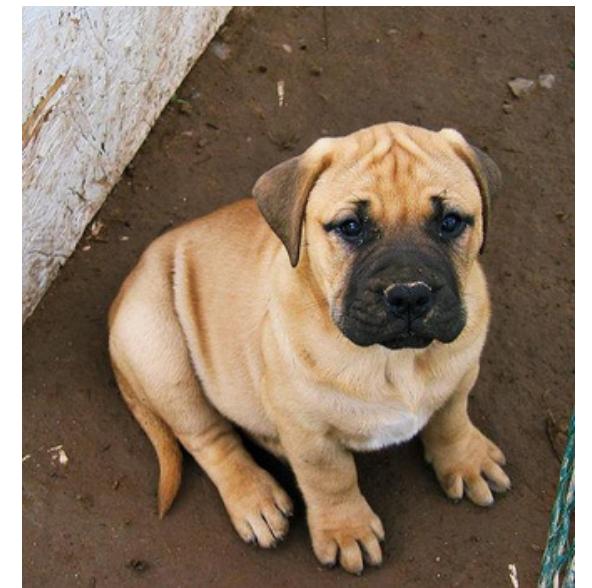
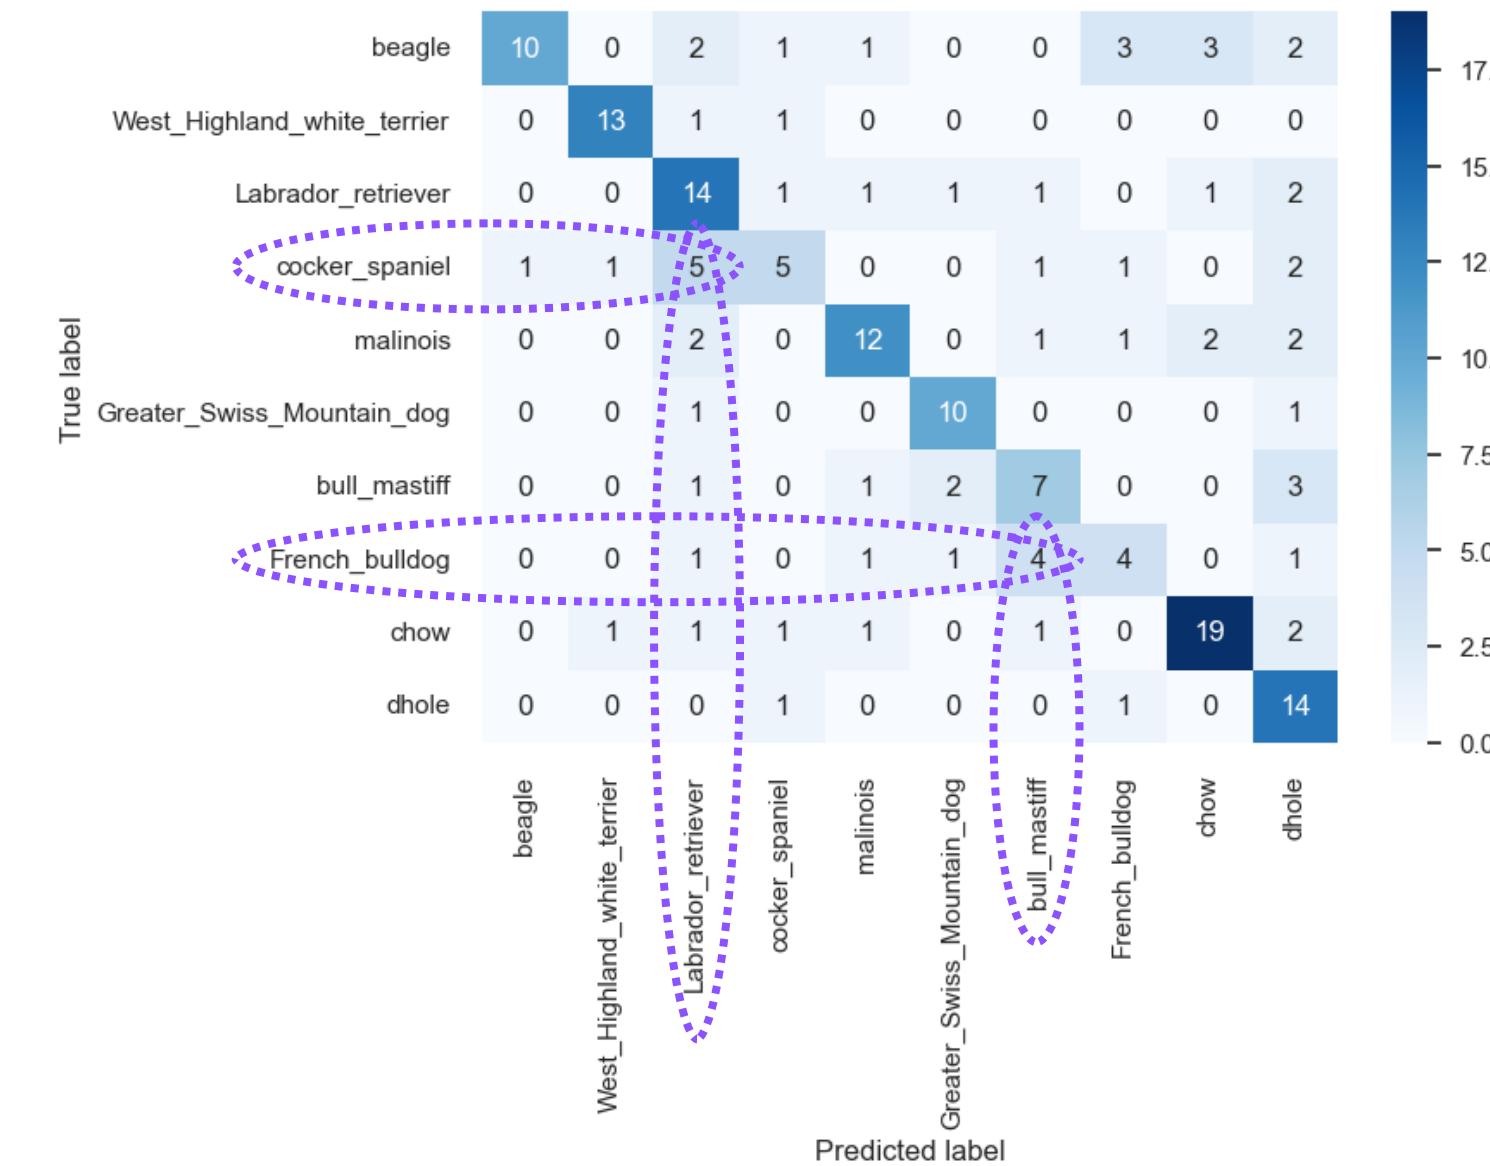
True: French\_bulldog  
Pred: dhole

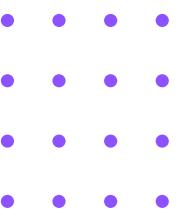




# Création d'un CNN "from scratch"

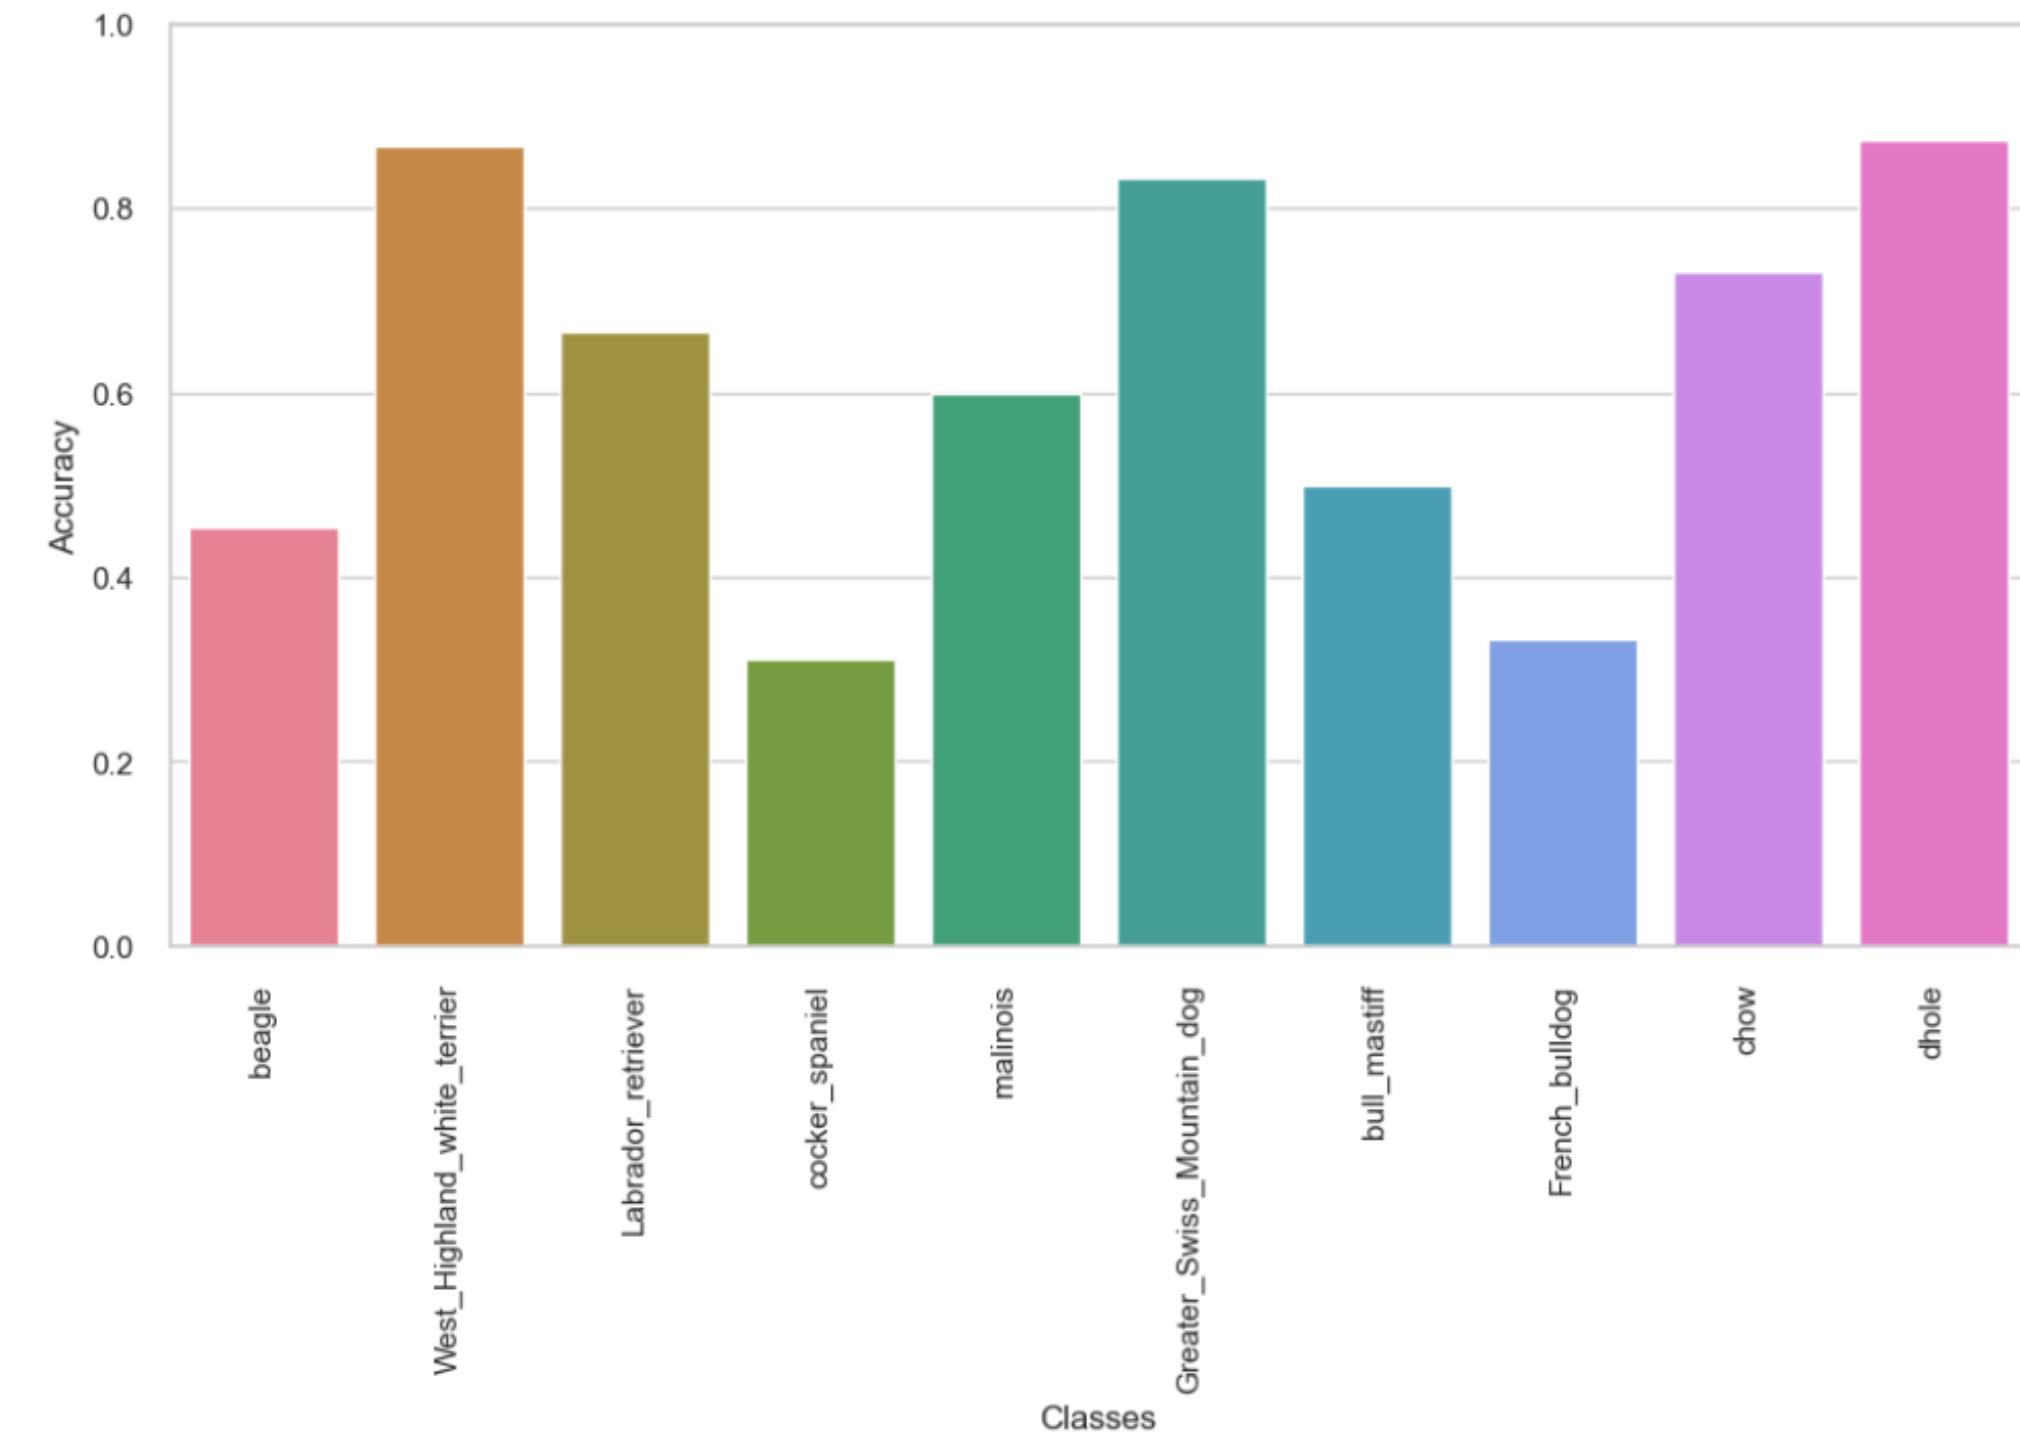
## Analyse des résultats

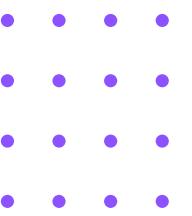




# Création d'un CNN "from scratch"

## Analyse des résultats





# Modèles avec Transfer Learning

## Xception (Google - 2017)



Test accuracy: 0.9712643623352051

## InceptionV3 (Google - 2015)

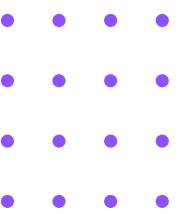


Test accuracy: 0.9942528605461121

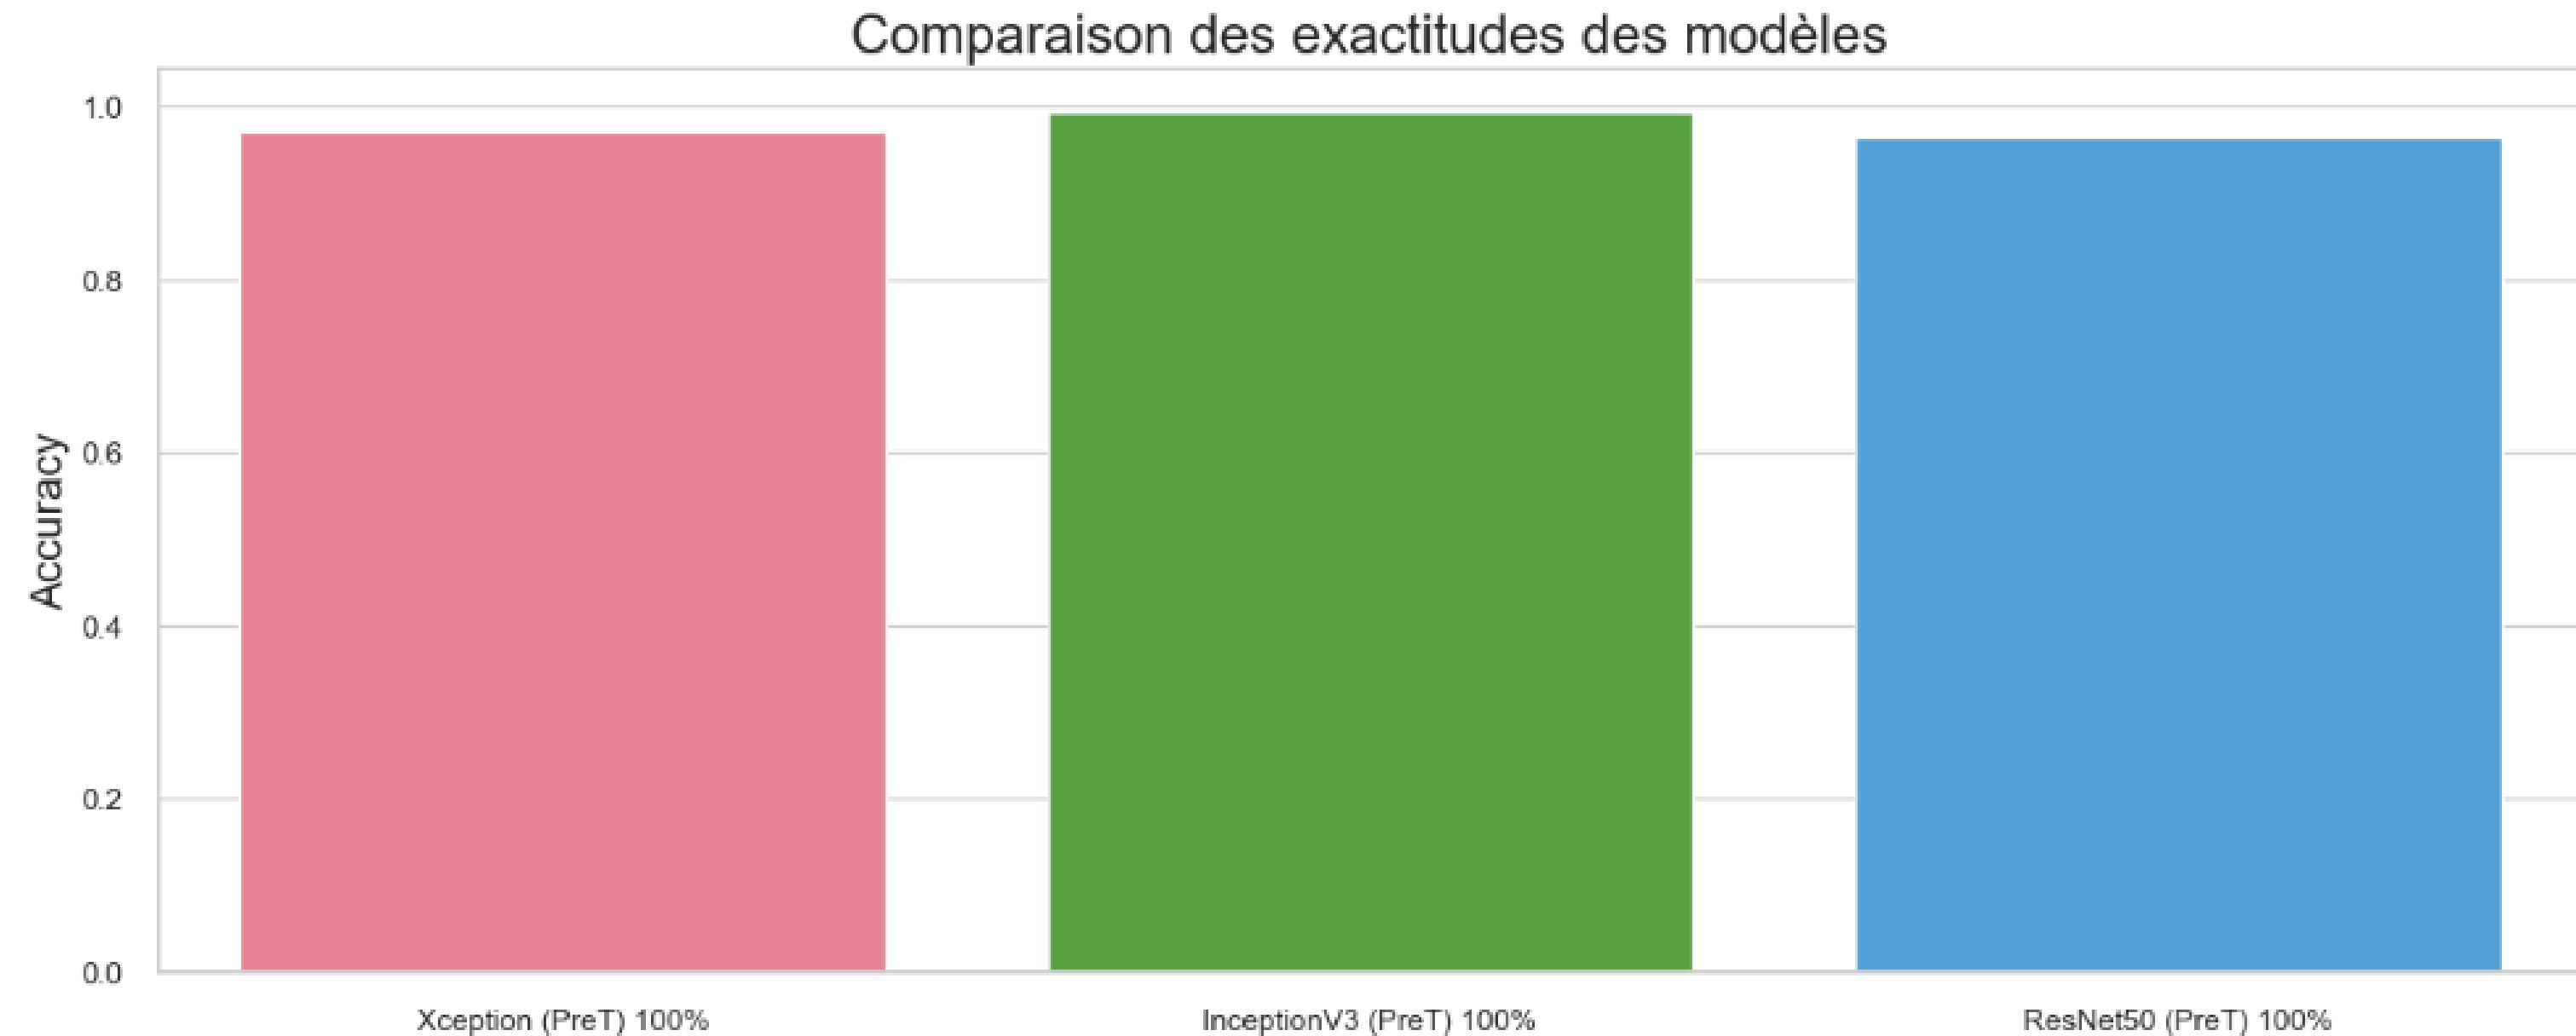
## ResNet50 (Microsoft - 2015)

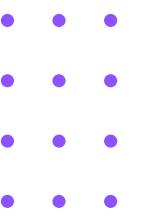


Test accuracy: 0.9655172228813171



# Modèles avec Transfer Learning





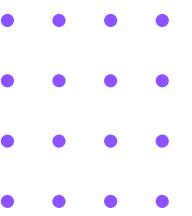
# Modèles avec Transfer Learning

## InceptionV3 sur 10 races

		Predicted label									
		beagle	West_Highland_white_terrier	Labrador_retriever	cocker_spaniel	malinois	Greater_Swiss_Mountain_dog	bull_mastiff	French_bulldog	chow	dhole
True label	beagle	18	0	0	0	0	0	0	0	0	0
	West_Highland_white_terrier	0	22	0	0	1	0	0	0	0	0
	Labrador_retriever	0	0	17	0	0	0	0	0	0	0
	cocker_spaniel	0	0	0	22	0	0	0	0	0	0
	malinois	1	0	0	0	22	0	0	0	0	0
	Greater_Swiss_Mountain_dog	0	0	0	0	0	15	0	0	0	0
	bull_mastiff	0	0	0	0	0	0	9	0	0	0
	French_bulldog	0	0	0	0	0	0	0	9	0	0
	chow	0	0	0	0	0	0	0	0	22	0
	dhole	0	0	0	0	0	0	0	0	0	15

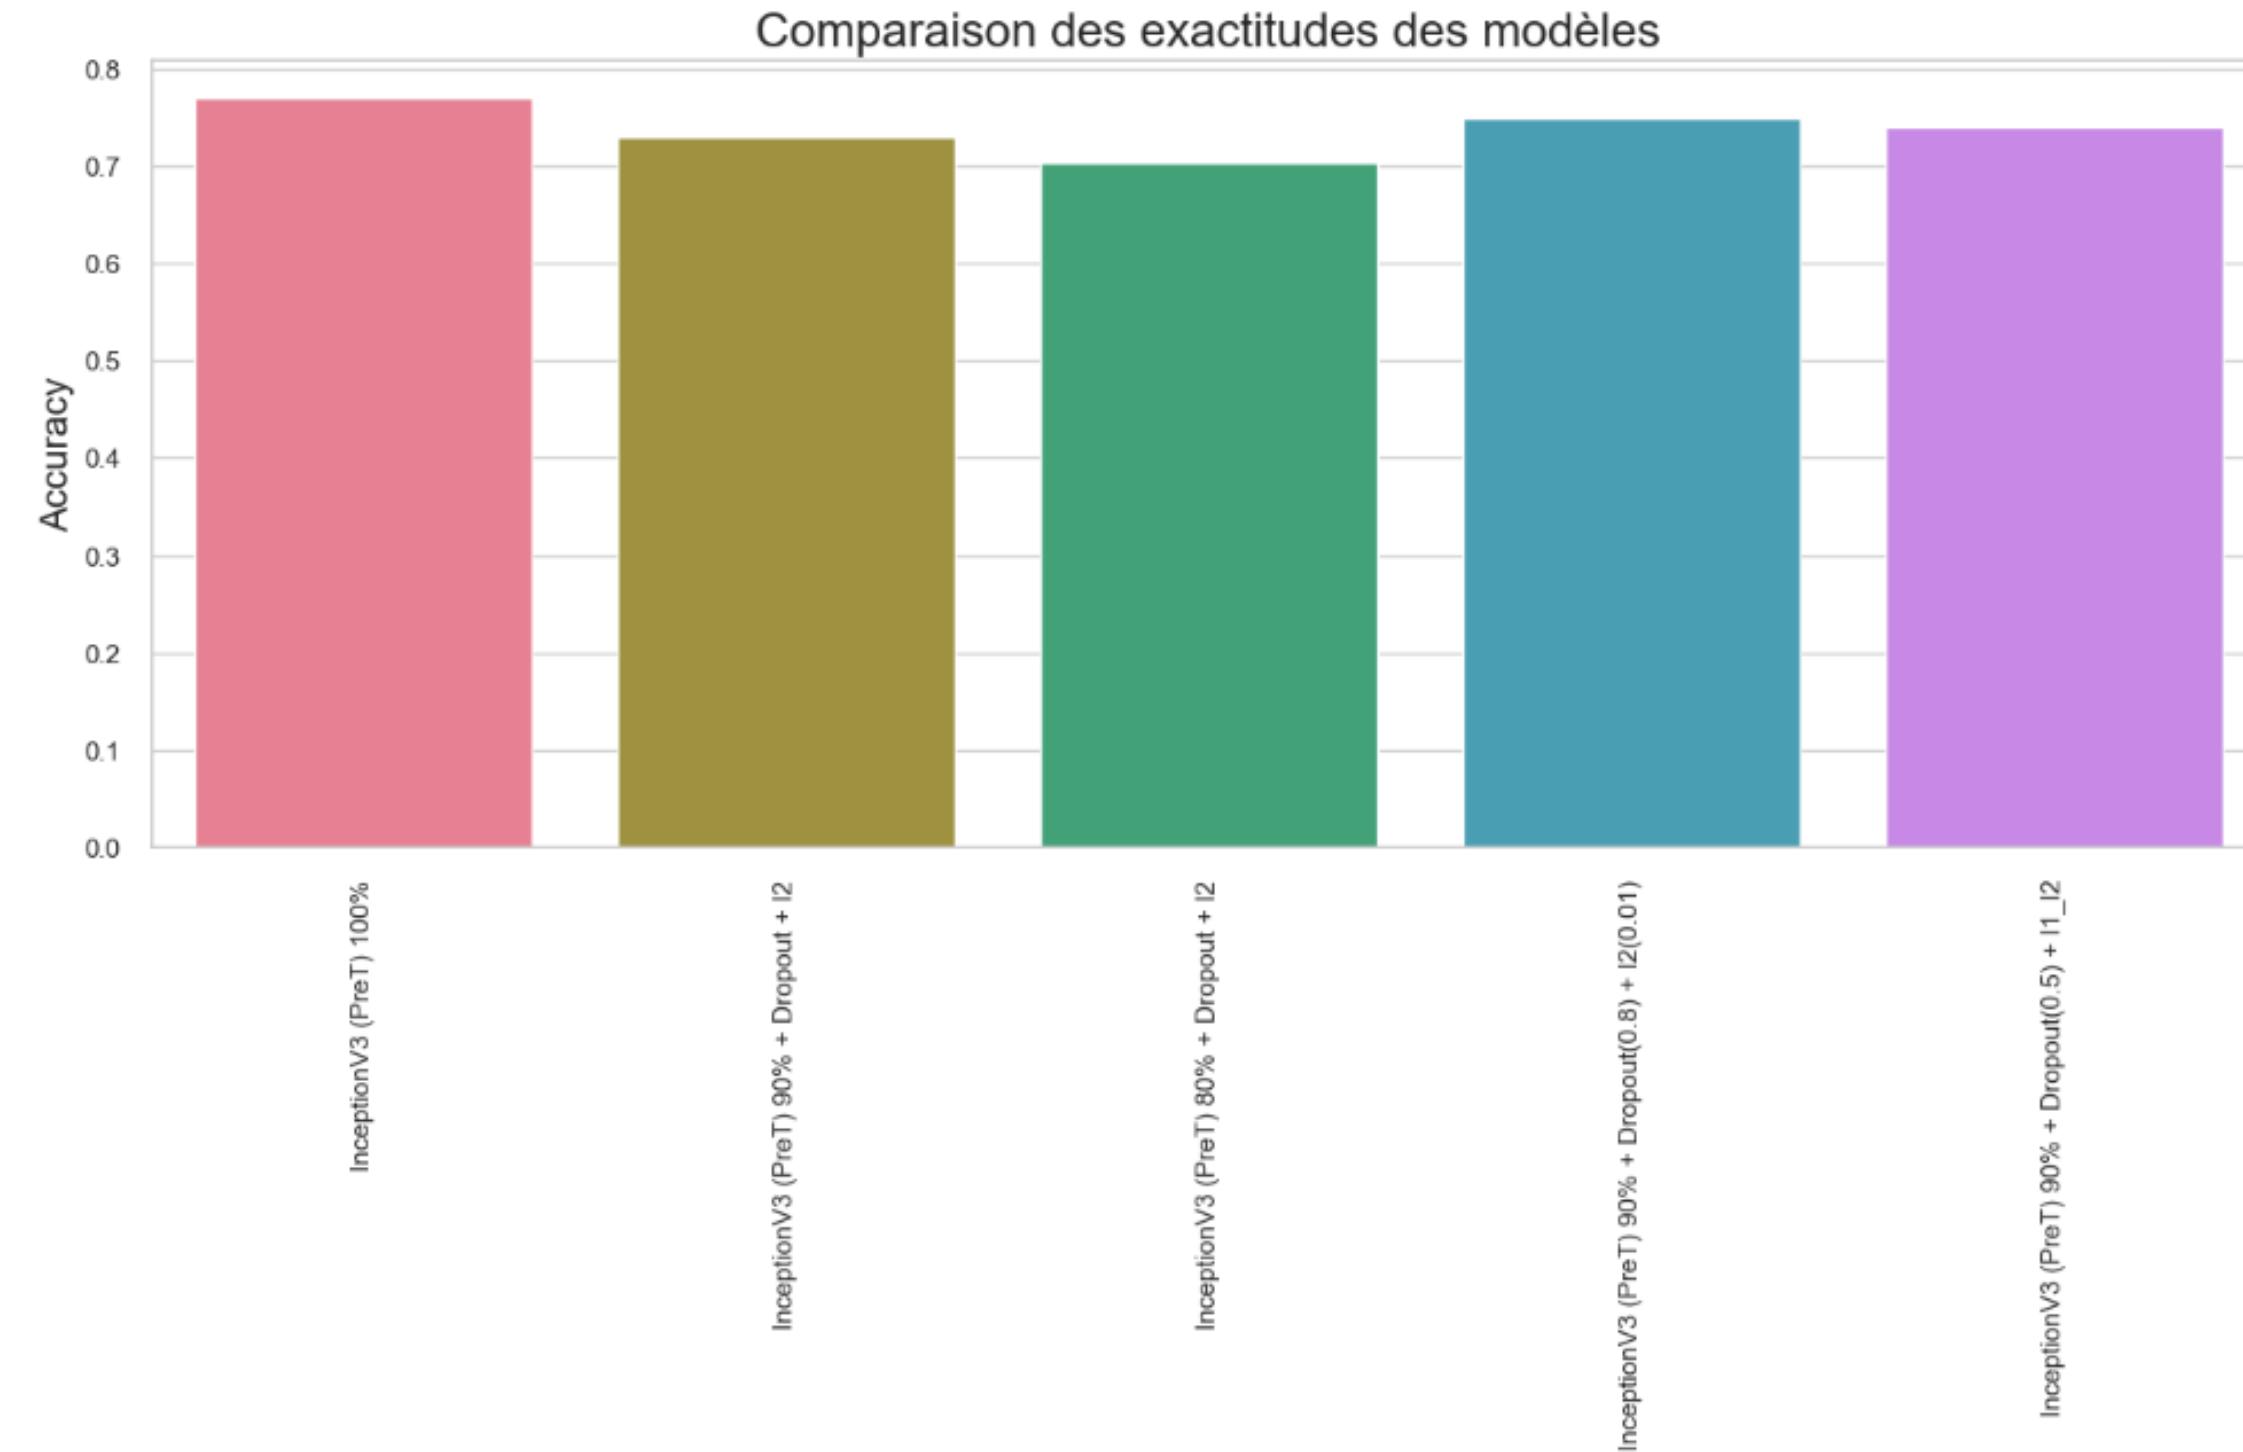
		Predicted label									
		beagle	West_Highland_white_terrier	Labrador_retriever	cocker_spaniel	malinois	Greater_Swiss_Mountain_dog	bull_mastiff	French_bulldog	chow	dhole
True label	beagle	0.95	0.00	0.05	0.00	0.00	0.00	0.00	0.00	0.00	0.00
	West_Highland_white_terrier	0.00	0.96	0.00	0.04	0.00	0.00	0.00	0.00	0.00	0.00
	Labrador_retriever	0.00	0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
	cocker_spaniel	0.00	0.00	0.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00
	malinois	0.04	0.00	0.00	0.00	0.96	0.00	0.00	0.00	0.00	0.00
	Greater_Swiss_Mountain_dog	0.00	0.00	0.00	0.00	0.00	0.00	1.00	0.00	0.00	0.00
	bull_mastiff	0.00	0.00	0.00	0.00	0.00	0.00	0.00	1.00	0.00	0.00
	French_bulldog	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	1.00	0.00
	chow	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	1.00	0.00
	dhole	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	1.00

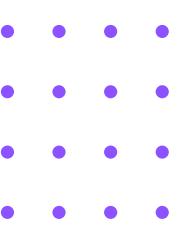




# Modèles avec Transfer Learning

## InceptionV3 sur 120 races





# Modèles avec Transfer Learning

## InceptionV3 sur 120 races

```
base_model = InceptionV3(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
for layer in base_model.layers:
    layer.trainable = True

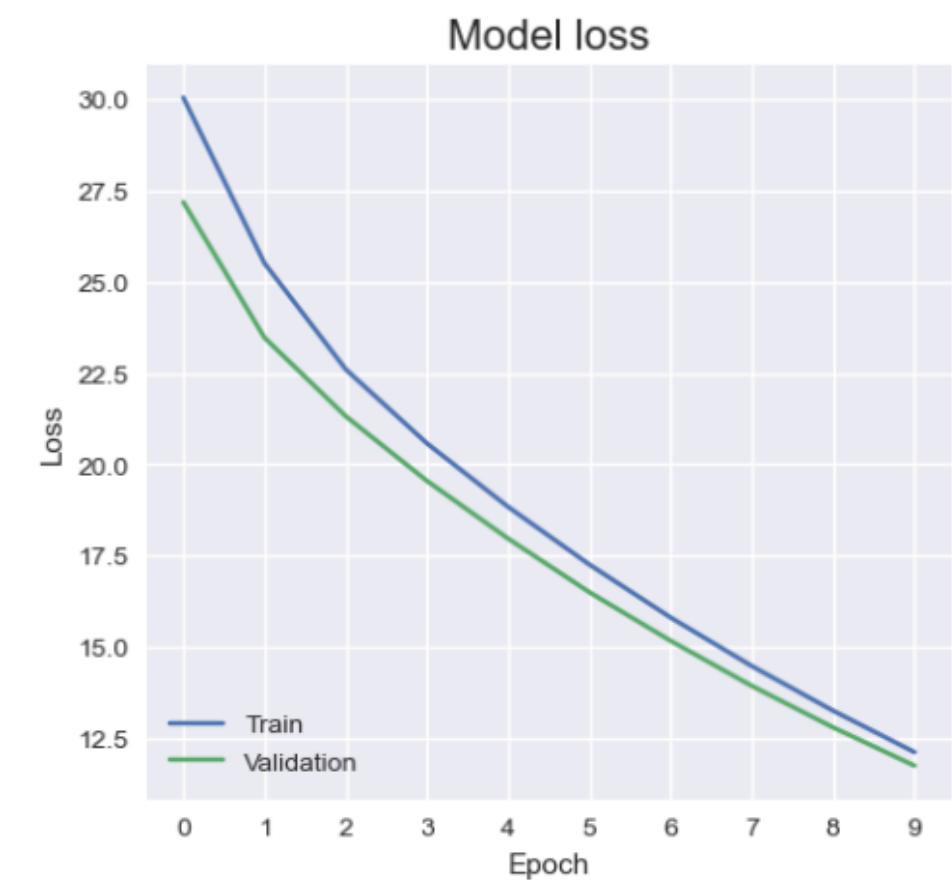
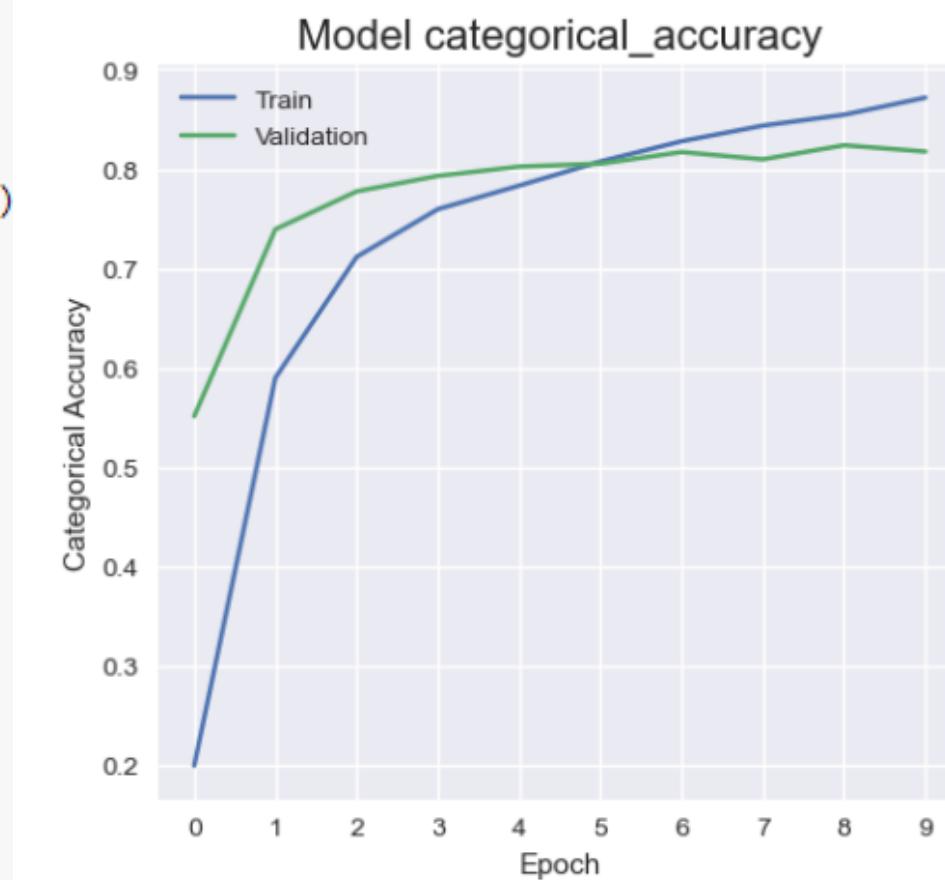
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(2048, activation='relu', kernel_regularizer=l2(0.01))(x)
x = Dropout(0.1)(x)
x = Dense(1024, activation='relu', kernel_regularizer=l2(0.005))(x)
x = Dropout(0.05)(x)

predictions = Dense(num_classes120, activation='softmax')(x)

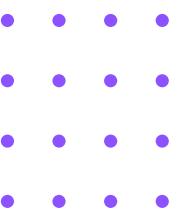
model_120 = Model(inputs=base_model.inputs, outputs=predictions)

model_120.compile(
    optimizer=Adam(lr=0.000008),
    loss='categorical_crossentropy',
    metrics=['categorical_accuracy']
)

early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)
```

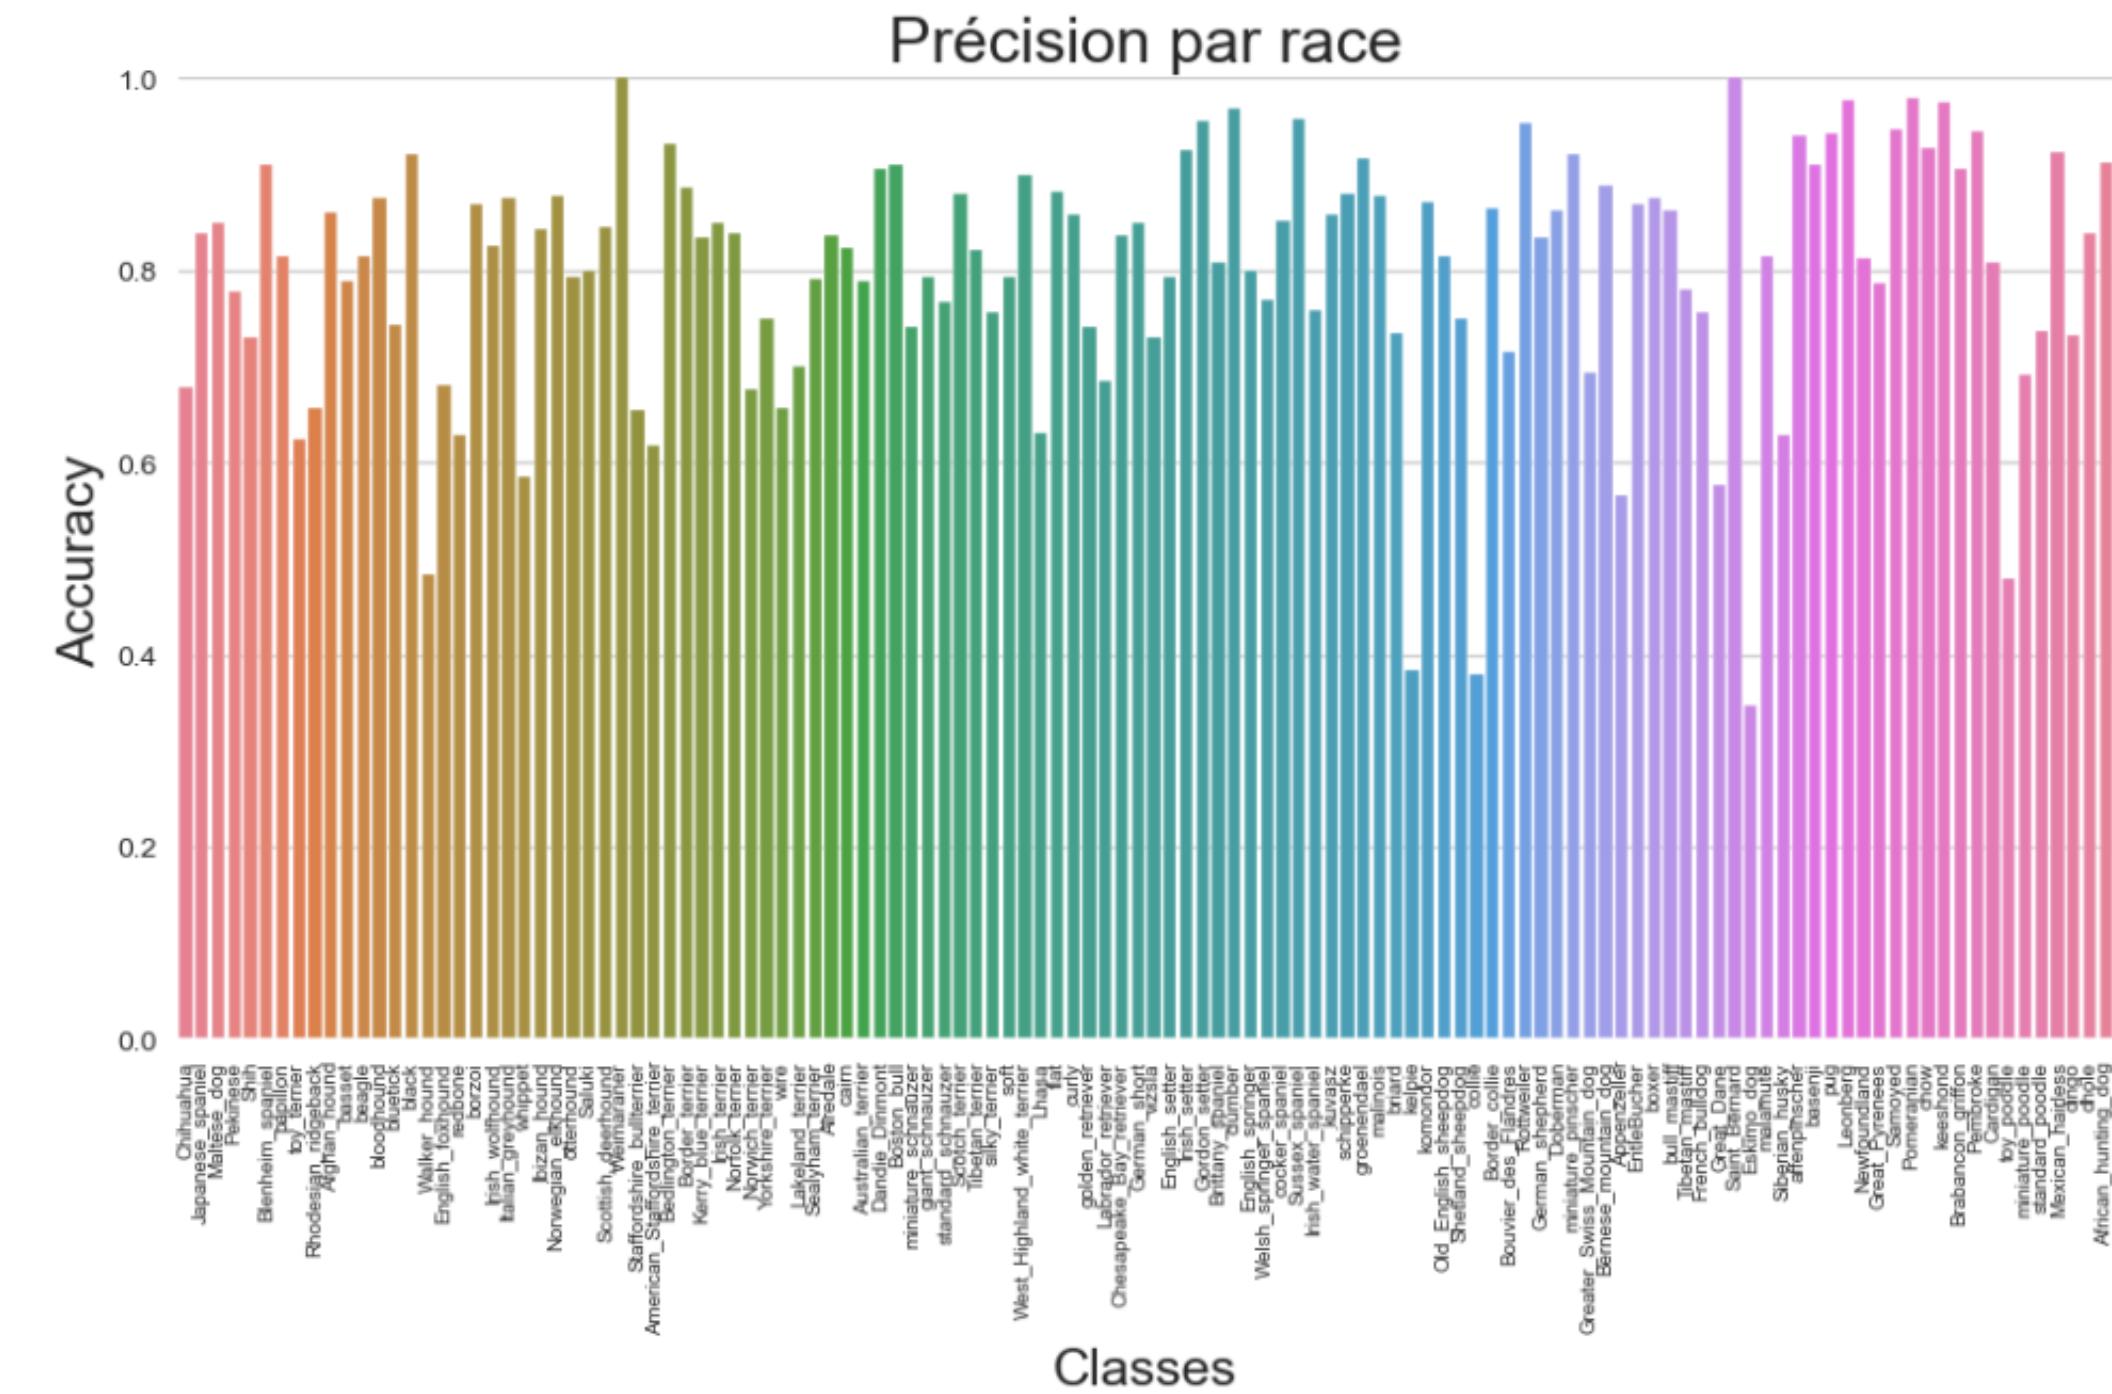


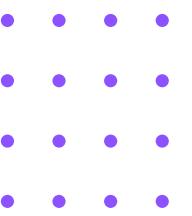
Test accuracy: 0.813829779624939



# Modèles avec Transfer Learning

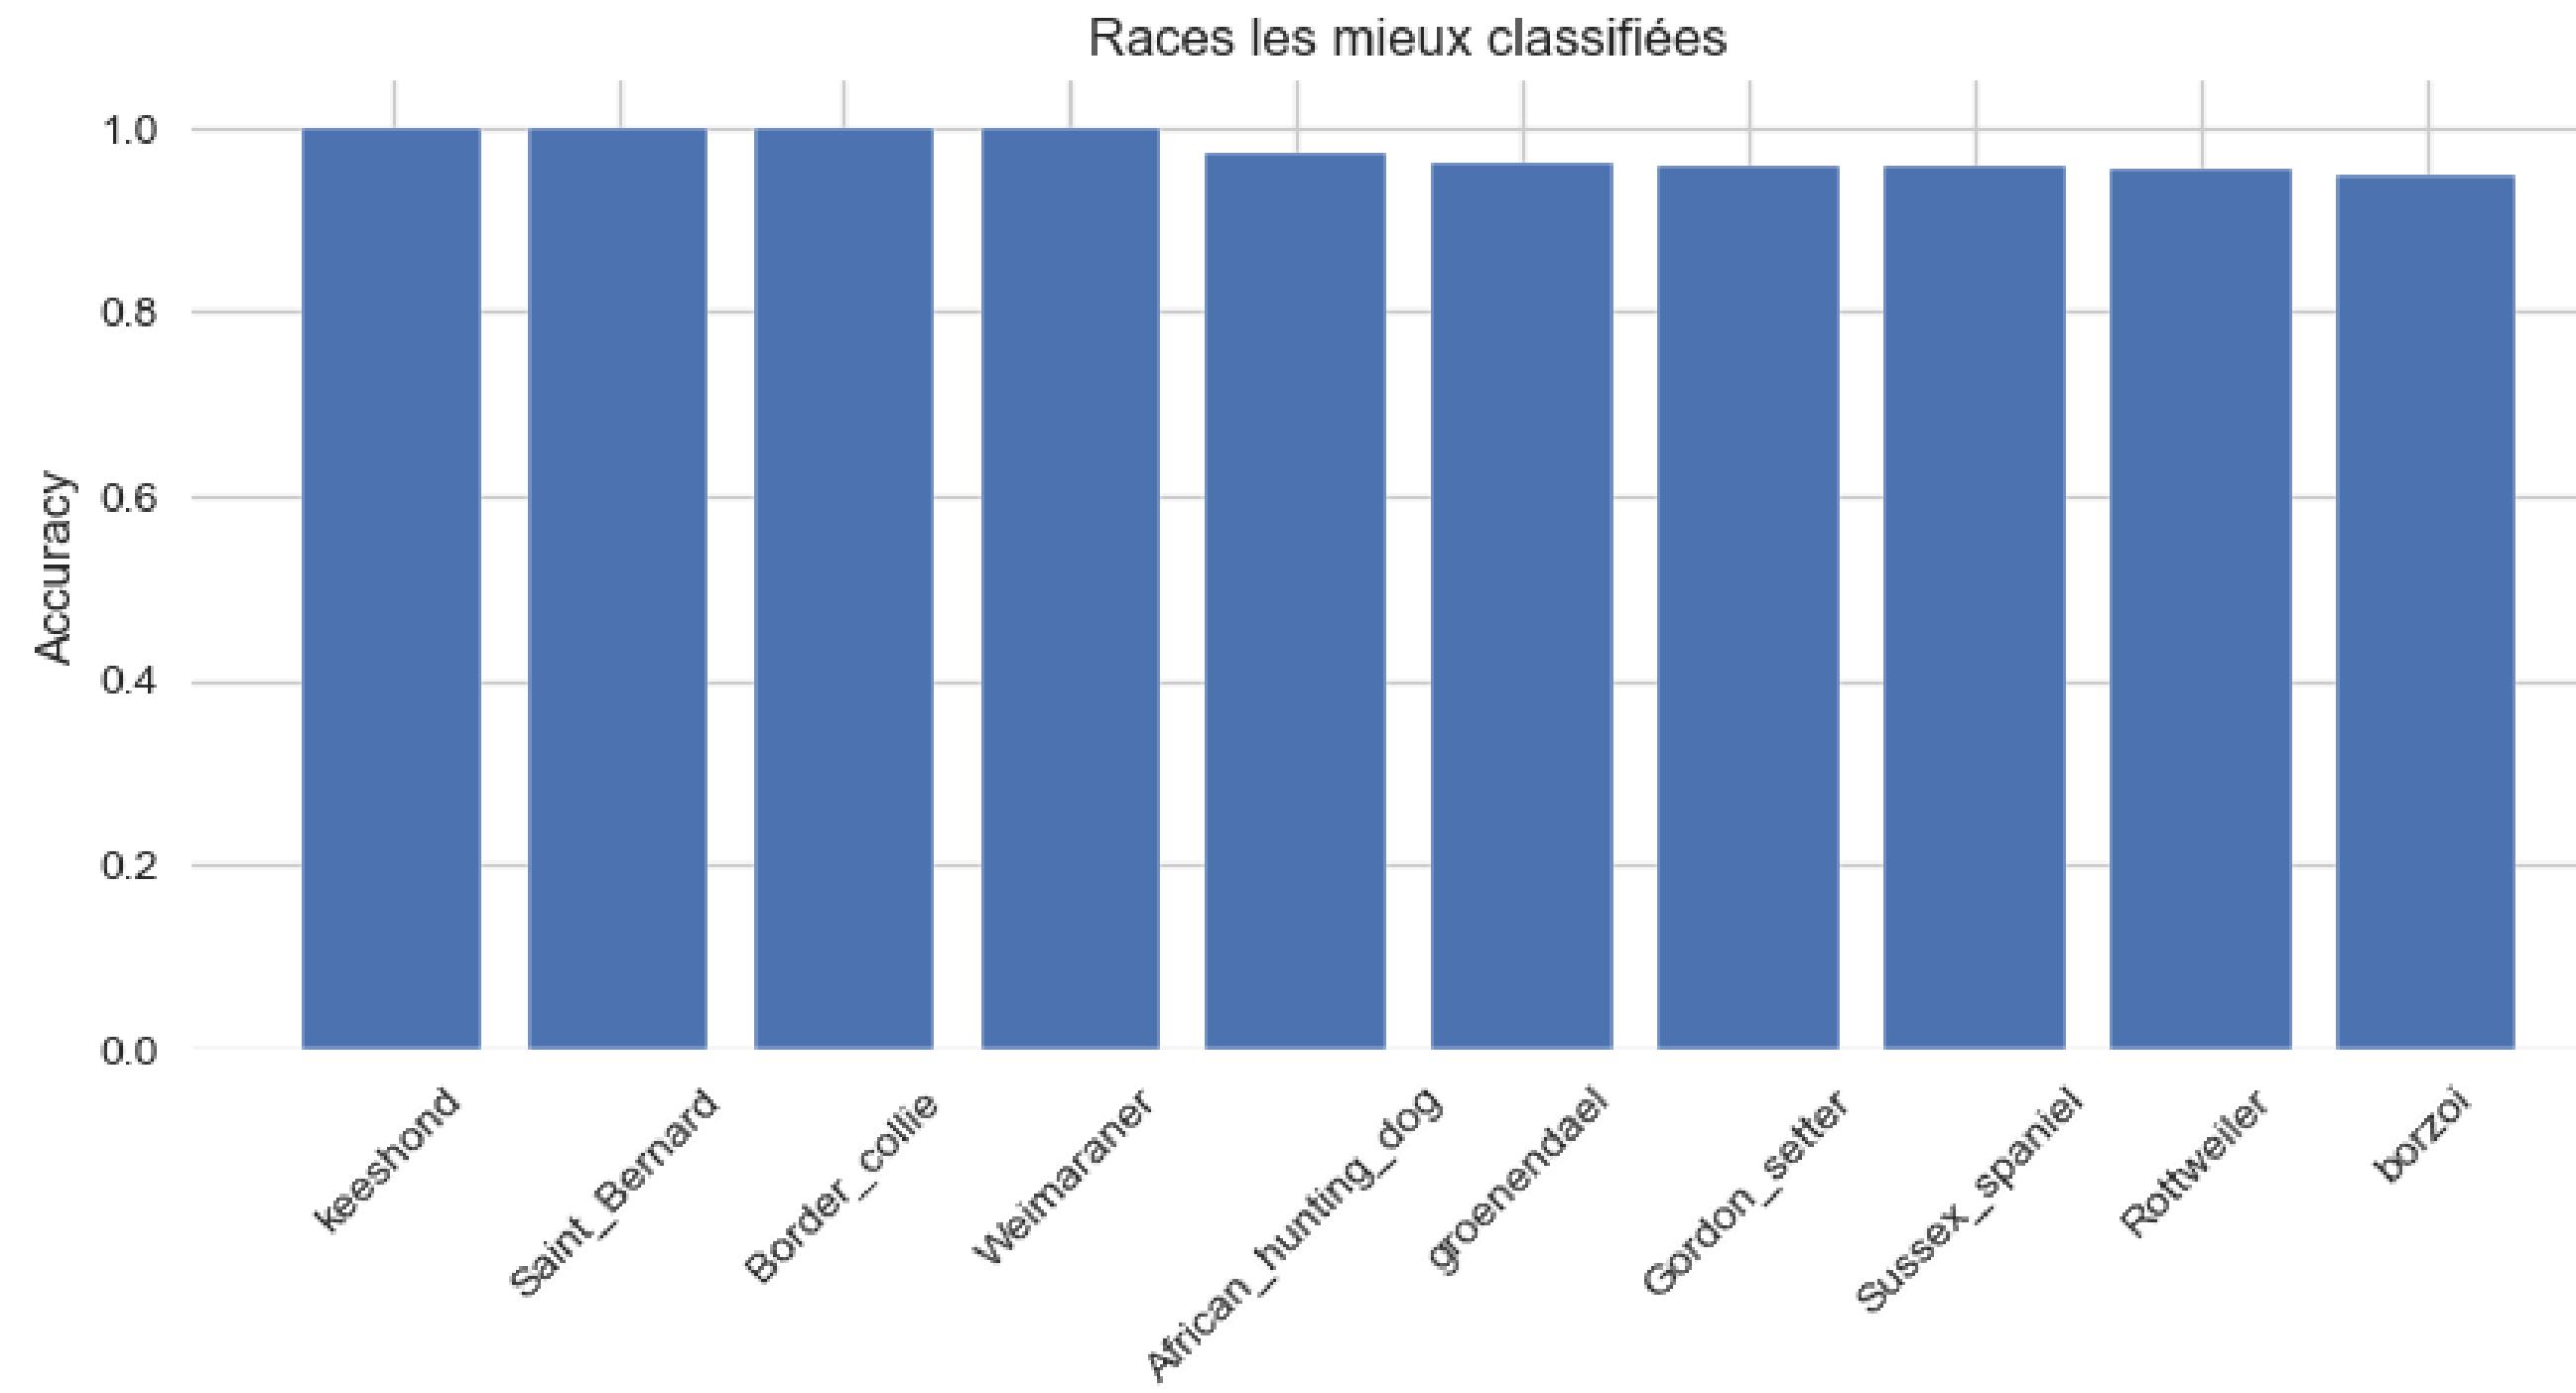
## InceptionV3 sur 120 races

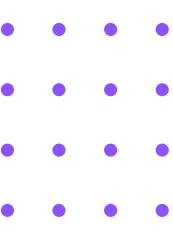




# Modèles avec Transfer Learning

## InceptionV3 sur 120 races

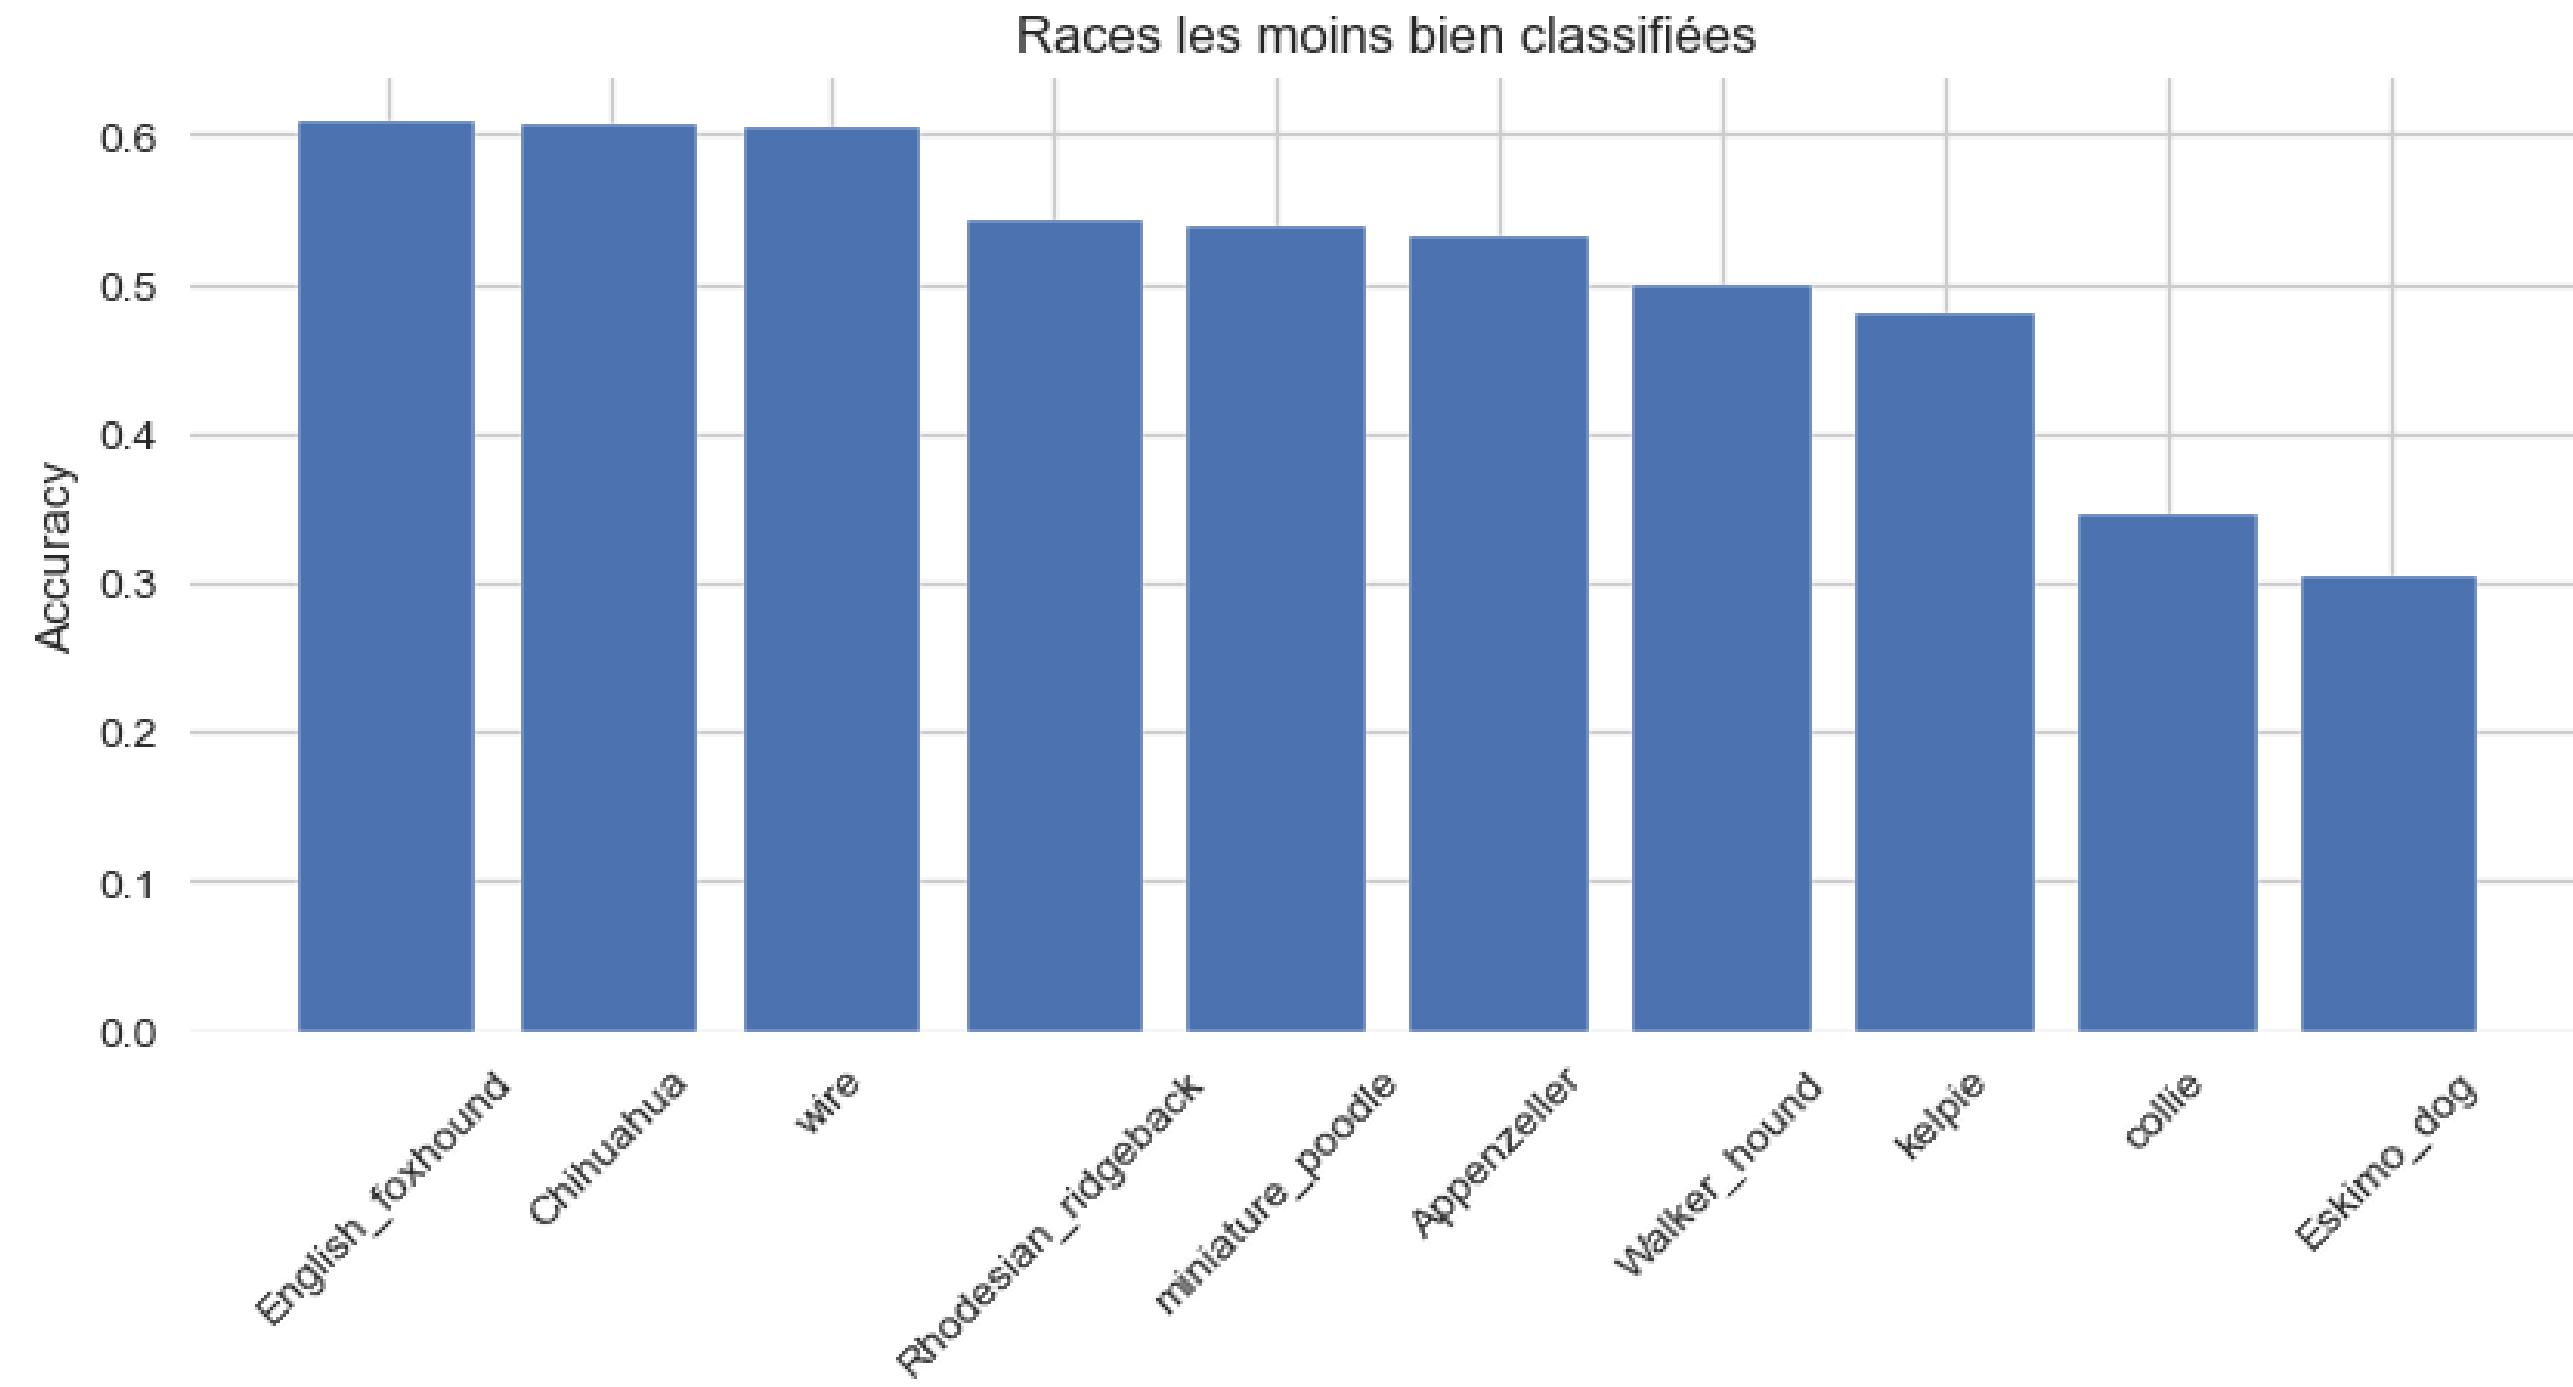


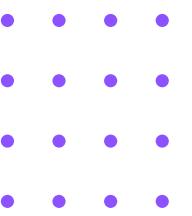


# Modèles avec Transfer Learning

InceptionV3 sur 120 races

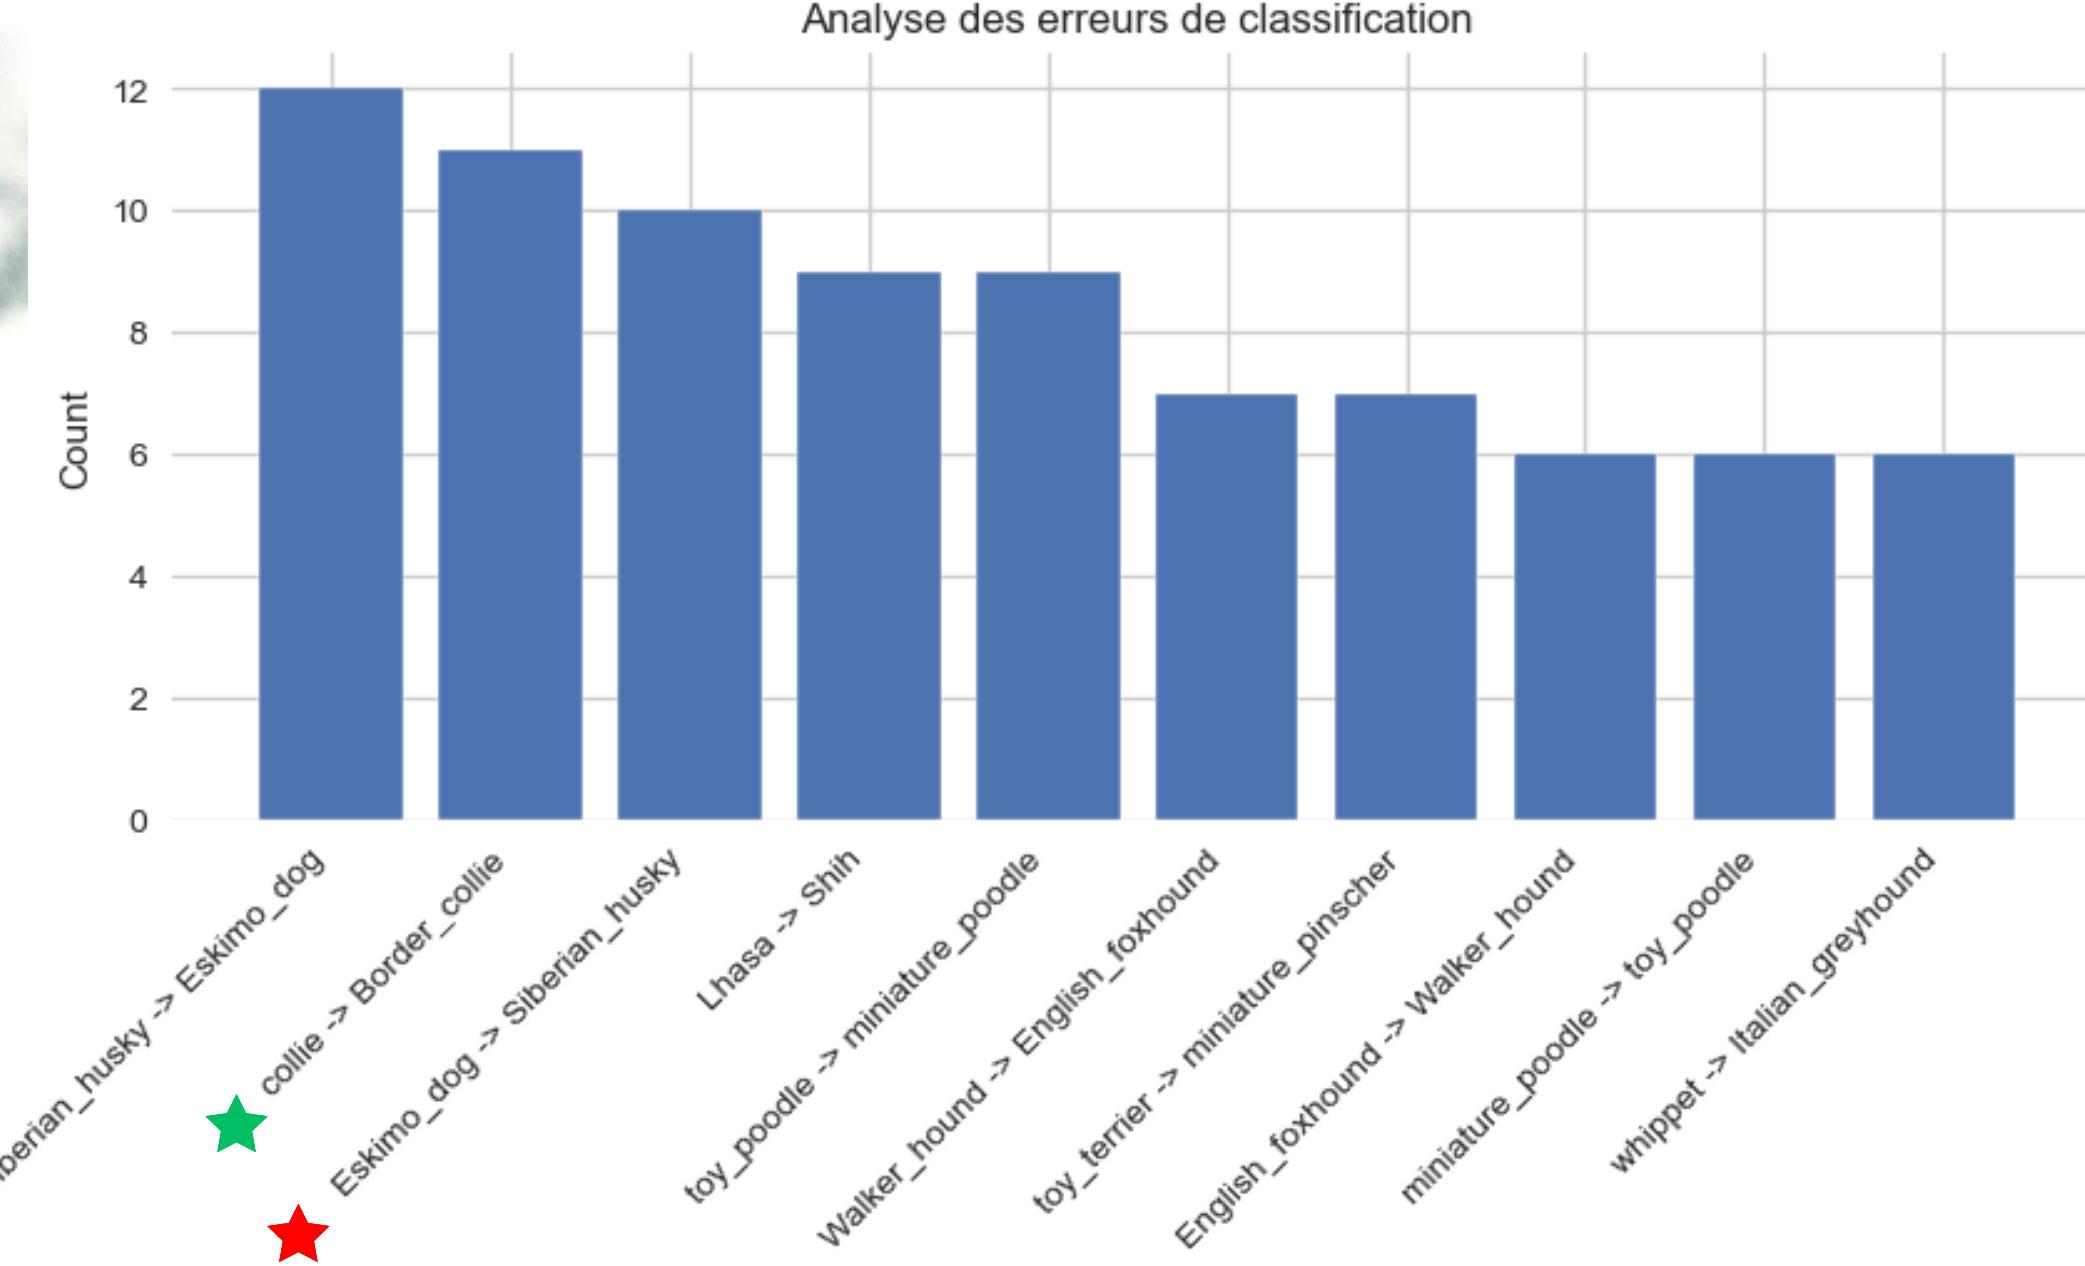
Misclassification error: 19.46%

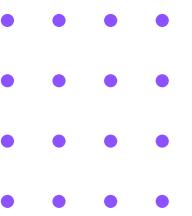




# Modèles avec Transfer Learning

## InceptionV3 sur 120 races





# Démonstration de l'API

<https://sdmwd-p6-iml-detecteur-de-race-de-chien.streamlit.app>

## Prédicteur de race de chien

Importez votre image

Drag and drop file here  
Limit 200MB per file • JPG

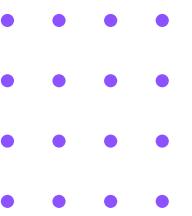
[Browse files](#)

 Sans titre.jpg 1.1MB X



Race prédictive : Labrador\_retriever

Précision de la prédition : 83.6%



# Rappel de l'objectif

- Réaliser un algorithme de détection de la race de chiens

# Conclusion

- Modèle à partir d'InceptionV3 (Transfer Learning + Fine Tuning)
- Modèle mis en application en ligne via Streamlit
- Erreur de classification très faible